

# CoreWeave Inc. (CRWV)

Consistent execution across AI market opportunity needed prior to seeing multiple expansion: Initiate Neutral (\$54 PT)

CRWV

12m Price Target: \$54.00

Price: \$35.42

Upside: 52.5%

We initiate on CoreWeave with a Neutral rating and 12-month price target of \$54 using a 13.9x FY26 EV/EBIT multiple. This looks to balance the uniqueness of the asset with the absence of a wide cohort of public pure-play Gen-AI peers and weak market reception to leverage or Gen-AI (both of which are central to the success of CoreWeave's business model). Our growth estimates underwrite the path CoreWeave can take to continue to scale (supported by over \$26bn in RPO and a ratable revenue model). However, we understand that given the company's relative short operating history and current macro uncertainty, CoreWeave will need to deliver consistent execution across the following for the stock to break out of its current range: **1) Durability of topline.** Net new logo wins and scaling GTM can ease strong customer concentration (with MSFT and NVDA accounting for 77% of 2024 revenue), **2) Sustainability of CapEx investments.** CoreWeave can improve its return on net assets by driving deployment efficiencies and expanding up the stack, and **3) Management of a scaling debt load.** While we estimate debt to exceed \$36bn by 2027, up 4.5x from current levels, CoreWeave will need to showcase that its revenue-backed financing philosophy can drive cost of capital towards the hyperscalers' low-single digit levels (vs. CoreWeave's 11.8% as of December). Given its relatively undemanding valuation relative to other companies in our coverage, we see upside to the stock if CRWV can land more sizeable customer wins, drive more RPO coverage against 2027 revenue targets and demonstrate incrementally more efficient conversion of deployed capital into revenues and EBIT.

## Kash Rangan

+1(415)249-7318 | kash.rangan@gs.com  
Goldman Sachs & Co. LLC

## Gili Naftalovich

+1(212)343-4286 | gili.naftalovich@gs.com  
Goldman Sachs & Co. LLC

## Henry Dane

+1(212)357-6378 | henry.dane@gs.com  
Goldman Sachs & Co. LLC

## Matthew Martino

+1(212)902-0695 | matt.martino@gs.com  
Goldman Sachs & Co. LLC

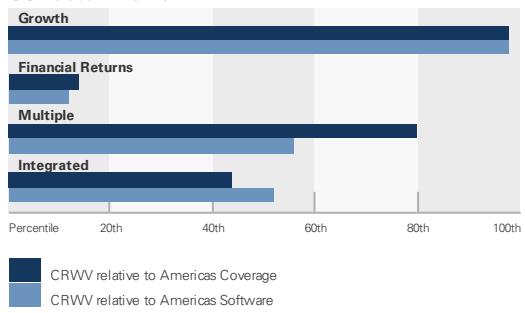
## Key Data

Market cap: \$20.3bn  
Enterprise value: \$35.2bn  
3m ADTV: NA  
United States  
Americas Software  
M&A Rank: 3

## GS Forecast

	12/24	12/25E	12/26E	12/27E
Revenue (\$ mn)	1,915.4	4,607.7	10,833.7	15,470.3
EBITDA (\$ mn)	1,219.3	3,149.9	8,164.6	12,377.9
EBIT (\$ mn)	355.8	813.7	2,554.8	4,289.5
EPS (\$)	–	(0.34)	(0.01)	2.38
P/E (X)	–	NM	NM	14.9
EV/EBITDA (X)	–	9.0	4.5	3.0
FCF yield (%)	NM	(116.5)	(64.6)	(21.0)
Dividend yield (%)	–	–	–	–
Net debt/EBITDA (X)	2.6	4.7	2.6	1.8
	12/24	3/25E	6/25E	9/25E
EPS (\$)	–	(0.05)	(0.00)	(0.03)

## GS Factor Profile



Source: Company data, Goldman Sachs Research estimates.  
See disclosures for details.

Neutral

**CoreWeave Inc. (CRWV)**

Rating since Apr 21, 2025

**Ratios & Valuation**

	<b>12/24</b>	<b>12/25E</b>	<b>12/26E</b>	<b>12/27E</b>
P/E (X)	—	NM	NM	14.9
EV/EBITDA (X)	—	9.0	4.5	3.0
EV/sales (X)	—	6.1	3.4	2.4
FCF yield (%)	NM	(116.5)	(64.6)	(21.0)
EV/DACF (X)	NM	10.1	7.3	4.1
CROCI (%)	11.0	13.5	13.5	18.3
ROE (%)	NM	(14.2)	(0.1)	39.8
Net debt/EBITDA (X)	2.6	4.7	2.6	1.8
Net debt/equity (%)	244.3	526.7	758.2	524.6
Interest cover (X)	1.0	0.8	1.1	1.7
Inventory days	NM	NM	NM	NM
Receivable days	55.4	71.6	84.6	92.4
Days payable outstanding	491.0	367.5	186.8	110.9

**Growth & Margins (%)**

	<b>12/24</b>	<b>12/25E</b>	<b>12/26E</b>	<b>12/27E</b>
Total revenue growth	736.6	140.6	135.1	42.8
EBITDA growth	173,336.6	158.3	159.2	51.6
EPS growth	—	NM	98.3	39,824.7
DPS growth	NM	NM	NM	NM
Gross margin	74.3	74.6	75.8	77.5
EBIT margin	18.6	17.7	23.6	27.7

**Price Performance**

Source: FactSet. Price as of 21 Apr 2025 close.

**Income Statement (\$ mn)**

	<b>12/24</b>	<b>12/25E</b>	<b>12/26E</b>	<b>12/27E</b>
Total revenue	1,915.4	4,607.7	10,833.7	15,470.3
Cost of goods sold	(492.0)	(1,168.8)	(2,621.0)	(3,484.6)
SG&A	(1,067.5)	(2,625.2)	(5,657.8)	(7,696.2)
R&D	—	—	—	—
Other operating inc./exp.)	—	—	—	—
<b>EBITDA</b>	<b>1,219.3</b>	<b>3,149.9</b>	<b>8,164.6</b>	<b>12,377.9</b>
Depreciation & amortization	(863.4)	(2,336.2)	(5,609.8)	(8,098.4)
<b>EBIT</b>	<b>355.8</b>	<b>813.7</b>	<b>2,554.8</b>	<b>4,289.5</b>
Net interest inc./exp.)	(360.8)	(1,053.2)	(2,263.5)	(2,512.8)
Income/(loss) from associates	—	—	—	—
<b>Pre-tax profit</b>	<b>54.3</b>	<b>(166.5)</b>	<b>387.6</b>	<b>1,879.7</b>
Provision for taxes	(119.2)	(5.4)	(391.1)	(469.9)
Minority interest	—	—	—	—
Preferred dividends	—	—	—	—
<b>Net inc. (pre-exceptionals)</b>	<b>(64.9)</b>	<b>(171.8)</b>	<b>(3.5)</b>	<b>1,409.7</b>
<b>Net inc. (post-exceptionals)</b>	<b>(863.4)</b>	<b>(482.4)</b>	<b>(431.3)</b>	<b>992.7</b>
<b>EPS (basic, pre-except) (\$)</b>	<b>--</b>	<b>(0.46)</b>	<b>(0.01)</b>	<b>3.22</b>
<b>EPS (diluted, pre-except) (\$)</b>	<b>--</b>	<b>(0.34)</b>	<b>(0.01)</b>	<b>2.38</b>
<b>EPS (ex-ESO exp., dil.) (\$)</b>	<b>--</b>	<b>--</b>	<b>--</b>	<b>--</b>
DPS (\$)	—	—	—	—
Div. payout ratio (%)	NM	0.0	0.0	0.0
Wtd avg shares out. (basic) (mn)	0.0	377.1	426.8	437.2
Wtd avg shares out. (diluted) (mn)	0.0	500.4	577.5	591.4

**Balance Sheet (\$ mn)**

	<b>12/24</b>	<b>12/25E</b>	<b>12/26E</b>	<b>12/27E</b>
Cash & cash equivalents	1,398.5	2,164.2	2,099.0	2,152.3
Accounts receivable	416.5	1,390.1	3,630.8	4,204.1
Inventory	—	—	—	—
Other current assets	101.2	239.3	350.6	421.7
<b>Total current assets</b>	<b>1,916.2</b>	<b>3,793.6</b>	<b>6,080.4</b>	<b>6,778.1</b>
Net PP&E	14,504.3	33,158.8	46,199.1	55,030.9
Net intangibles	24.5	24.5	24.5	24.5
Total investments	0.0	0.0	0.0	0.0
Other long-term assets	1,387.6	1,512.9	1,456.0	1,375.0
<b>Total assets</b>	<b>17,832.6</b>	<b>38,489.7</b>	<b>53,759.9</b>	<b>63,208.5</b>
Accounts payable	868.3	1,485.7	1,196.6	921.2
Short-term debt	2,468.4	5,311.8	4,673.0	3,419.8
Current lease liabilities	270.9	446.7	700.2	911.7
Other current liabilities	1,355.0	1,245.0	1,509.1	1,767.5
<b>Total current liabilities</b>	<b>4,962.6</b>	<b>8,489.2</b>	<b>8,078.8</b>	<b>7,020.2</b>
Long-term debt	5,457.9	17,692.2	28,141.5	32,785.5
Non-current lease liabilities	2,423.0	4,774.1	7,838.6	10,206.5
Other long-term liabilities	3,680.6	4,695.4	6,865.6	8,951.2
<b>Total long-term liabilities</b>	<b>11,561.5</b>	<b>27,161.6</b>	<b>42,845.7</b>	<b>51,943.2</b>
<b>Total liabilities</b>	<b>16,524.1</b>	<b>35,650.8</b>	<b>50,924.5</b>	<b>58,963.3</b>
Preferred shares	1,722.1	—	—	—
<b>Total common equity</b>	<b>(413.6)</b>	<b>2,838.9</b>	<b>2,835.4</b>	<b>4,245.2</b>
Minority interest	—	—	—	—
<b>Total liabilities &amp; equity</b>	<b>17,832.6</b>	<b>38,489.7</b>	<b>53,759.9</b>	<b>63,208.5</b>
BVPS (\$)	—	5.67	4.91	7.18

**Cash Flow (\$ mn)**

	<b>12/24</b>	<b>12/25E</b>	<b>12/26E</b>	<b>12/27E</b>
Net income	(863.4)	(482.4)	(431.3)	992.7
D&A add-back	863.4	2,336.2	5,609.8	8,088.4
Minority interest add-back	—	—	—	—
Net (inc)/dec working capital	1,800.6	652.5	(607.5)	1,057.1
Others	948.6	410.1	1,012.5	959.2
<b>Cash flow from operations</b>	<b>2,749.2</b>	<b>2,916.4</b>	<b>5,583.5</b>	<b>11,097.4</b>
Capital expenditures	(8,702.1)	(18,419.3)	(15,314.9)	(14,353.7)
Acquisitions	—	—	—	—
Divestitures	—	—	—	—
Others	44.0	—	—	—
<b>Cash flow from investing</b>	<b>(8,658.1)</b>	<b>(18,419.3)</b>	<b>(15,314.9)</b>	<b>(14,353.7)</b>
Dividends paid	(57.7)	(29.3)	—	—
Share issuance/(repurchase)	1.4	1,388.6	—	—
Inc/(dec) in debt	6,430.0	14,996.0	9,712.1	3,297.6
Others	673.7	(28.6)	(12.0)	12.0
<b>Cash flow from financing</b>	<b>7,047.3</b>	<b>16,268.6</b>	<b>9,666.2</b>	<b>3,309.6</b>
<b>Total cash flow</b>	<b>1,138.4</b>	<b>765.7</b>	<b>(65.2)</b>	<b>53.3</b>
Free cash flow	(5,952.9)	(15,502.9)	(9,731.3)	(3,256.3)
Free cash flow per share (basic) (\$)	NM	(41.27)	(22.88)	(7.45)

Source: Company data, Goldman Sachs Research estimates.

## Table of Contents

Executive Summary	4
Top Line Durability: Agility of strategic differentiation to determine LT positioning	7
Top Line Durability: AI Market Outlook	13
Top Line Durability: Customer diversification likely to take time	18
Top Line Durability: Contract construction drives visibility	21
Top Line Durability: Revenue as a function of improving ROA	28
CapEx Buildout: Investment philosophy based on customer commitments	31
CapEx Buildout: Depreciation dynamics key to model assumptions	37
CapEx Buildout: Strategic navigation of supply constraints	42
Debt Burden: Revenue-backed financing lowers risks associated with scaling debt needs	50
Debt Burden: Path to lower cost of capital	59
Valuation	64
Key Risks	67
Disclosure Appendix	70

# Executive Summary

Our Neutral thesis on CoreWeave is predicated on our view that 1) CoreWeave is a leading provider of AI compute at scale, sitting at the cusp of a massive paradigm shift, and 2) the bar for success is extremely high, with CoreWeave facing large hyperscaler incumbents, a relatively concentrated customer base, and a high debt burden. Further, CoreWeave is indexed towards the current overhangs of the broader market, including tariffs, AI skepticism, and caution in underwriting debt-heavy growth companies.

While our growth estimates underwrite the path CoreWeave can take to continue to scale (supported by over \$26bn in RPO and a ratable revenue model), we understand that given the company's relative short operating history and current macro uncertainty, we expect it will need to deliver consistent execution to sustainably break out of its current range. Below, we outline the dynamics driving our investment thesis.

## Top Line Durability

- **Ability of strategic differentiation to determine LT positioning:** While being a best-in-class provider of AI compute, we expect its competitive differentiation to evolve as we move through the S curve of Gen-AI adoption. We expect CoreWeave to go from leaning on the speed in which it brings next generation chips to market (short-term differentiator), to repeated success in deploying large scale, curated deployments (medium-term), and expansion of AI expertise up the stack (long-term).
- **AI market outlook:** Gen-AI is at the center of CoreWeave's market opportunity as the company provides access to a modernized tech stack configured specifically for the unique technical requirements of AI workloads. This positions CoreWeave to benefit from the proliferation of AI model training and inferencing as more use cases come to market. We note, however, that should future use cases require less performant, bespoke infrastructure, growth in inferencing may pose a headwind to CoreWeave.
- **Customer diversification likely to take time:** CoreWeave's ability to diversify its customer concentration will likely be the main near-term indicator investors look to gain comfort around the company's long-term revenue durability. While we view the company's GTM shift augmenting the market's expansion to support the longer-term broadening out of end-users, the large contracts CoreWeave has today will likely leave these dynamics largely unchanged in the near-term ([Exhibit 1](#)). Still, within these customers, CoreWeave holds a relatively small share of their CapEx spend, leaving room for expansion.
- **Contract construction drives visibility:** We see CoreWeave's visibility into future revenue as an underappreciated aspect of the company's model. With 95% committed contracts (a contrast to the more consumption-oriented hyperscalers), we conservatively estimate RPO coverage of 95%/59%/35% of our FY25/FY26/FY27 revenue estimates. The remaining revenue can come from a mix of customer expansions, new logo wins, and on-demand revenue ([Exhibit 2](#)).
- **Revenue as a function of improving ROA:** We see CoreWeave's revenue growth

benefiting from both RPO conversion and improvement on the return on net assets (ROA). We expect ROA to expand from 8.7% (on average) in 2024 to 10.6% by 2027 as on-demand revenue scales, CoreWeave drives further infrastructure efficiency, and the company potentially moves up the stack to monetize software solutions.

## CapEx Buildout

- **Investment philosophy based on customer commitments:** As CoreWeave does not purchase infrastructure without a signed contract, CapEx spend can be used as a leading indicator of revenue growth. This also reduces the potential risk of an overbuild cycle, a common concern investors hold for hyperscalers. We expect CapEx to nearly double in FY25 (with the rollout of Nvidia Blackwell chips and build out of OpenAI's +\$11bn contract) and moderate in the outer years; investors could see these estimates increase if CoreWeave signs more contracts than currently modeled (Exhibit 3).
- **Depreciation dynamics key to model assumptions:** The replacement cost of GPUs is a significant portion of CoreWeave's business. The 5-year Blackwell contract that OpenAI signed (that includes two one-year extensions) and extended useful lives of early GPU generations provide support in favor of CoreWeave's baseline 6-year useful life estimate. However, the high degree of sensitivity of CoreWeave's operating margins to these assumptions given the scaling number of GPUs results in a wide range of potential outcomes to our future margin estimates (Analysis below).
- **Strategic navigation of supply constraints:** Our Data Center, Utilities, Sustainability analysts, have started to forecast an alleviation of constraints in data center capacity, although still forecast strong power demand through 2030. Given time to market has been a key differentiator for CoreWeave, a more balanced market will likely accelerate the need for CoreWeave to lean in and solidify the other components of its value proposition (GPU-focus, improved reliability, dedicated infrastructure, software, etc).

## Debt Burden

- **Revenue-backed financing lowers risks associated with scaling debt needs:** We see CoreWeave's current valuation reflecting an outsized focus on the magnitude of the company's debt load while under-appreciating its favorable contract dynamics (with a 2.5-year average cash payback a ~4-year contract average). While debt will be the primary financing instrument to sustain CoreWeave's capital-intensive business, financing needs are largely an output of contracted deals, with debt vehicles amortized against the cash flows of customer contracts. This lowers the risk of competing cash flow needs and leads us to see the debt ramp as largely manageable.
- **Path to lower cost of capital:** Achieving a lower cost of capital is critical for CoreWeave to compete with competitors, who hold a weighted average cost of debt of 3.5%. We expect CoreWeave's cost of debt to come down from 11.8% in FY24 to 7.8% by the end of 2027 as the company benefits from scale, a longer operating history, the transparency and liquidity of the public markets, and customer

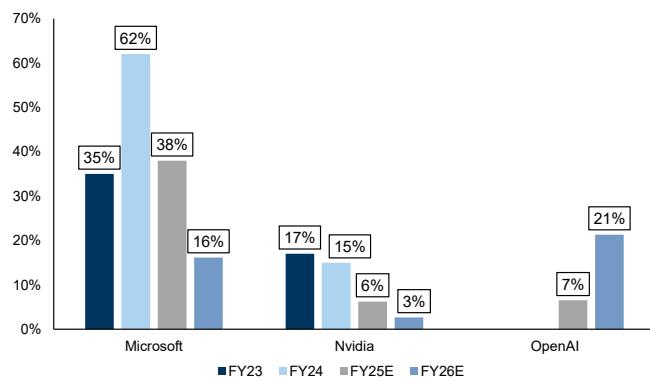
diversification. We estimate that any 100bps change in SOFR or cost of debt can drive a ~\$0.40-\$0.50 impact to 2027 EPS.

## Valuation

- We value CoreWeave at a second 12-month price target of \$54. We leverage an EV/EBIT methodology and apply a 13.9x multiple, based on a peer group of general-purpose cloud providers, high-growth hardware, and value-add hardware redistributors.

### Exhibit 1: Expect revenue concentration among a few large customers in short- to medium- term

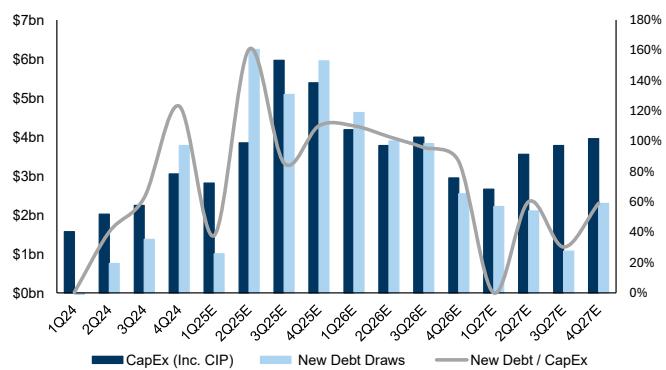
Large customers as a % of GSe CoreWeave Revenue



FY25/FY26 Microsoft estimates assume revenue is at midpoint of the \$1.2bn spent in FY24 and OpenAI's expected revenue contribution of ~\$2.3bn. FY25/FY26 Nvidia estimates assumes revenue \$ spend is flat relative to FY24. FY25/FY26 OpenAI revenue is GSe.

Source: Company data, Goldman Sachs Global Investment Research

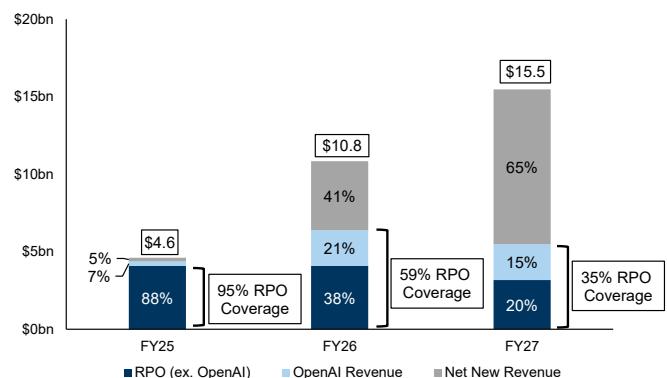
### Exhibit 3: Expect CapEx spend to fluctuate based on timing of deal wins and lean less on debt as ROA improves



Source: Company data, Goldman Sachs Global Investment Research

### Exhibit 2: We estimate CoreWeave has RPO coverage of 95% in FY25, 59% in FY26, and 35% in FY27

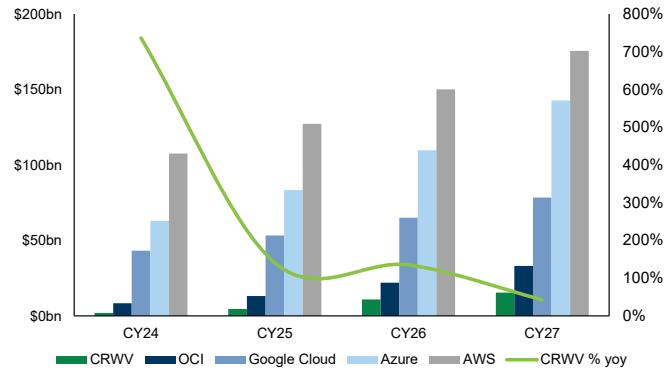
This includes OpenAI's +\$11bn contract, closed post Dec 2024



Source: Company data, Goldman Sachs Global Investment Research

### Exhibit 4: CoreWeave scaling disproportionately fast vs. hyperscaler peers

GS Cloud Revenue Estimates, Calendarized



Source: Company data, Goldman Sachs Global Investment Research

## Top Line Durability: Agility of strategic differentiation to determine LT positioning

While being a best-in-class provider of AI compute, we expect CoreWeave's competitive differentiation to evolve as we move through the S-curve of Gen-AI adoption.

CoreWeave provides its customers with access to large clusters of GPUs at a faster pace and with better performance-adjusted pricing than many of their nearest competitors. For companies that are looking to drive the course of the industry, the speed, scale, performance, and pricing of AI compute are essential considerations to maintain an edge over their competition - especially as it powers the solutions they bring to market.

While we appreciate the various uncertainties hovering over the durability of AI, we lay out the ways in which CoreWeave has positioned itself to take on this opportunity. As we move through the S curve of adoption and we progress toward a more steady state of supply/demand, we see CoreWeave needing to lean into different areas of its competitive differentiation to remain strategically positioned.

**We expect this to result in the company's strategy evolving from leaning on the speed in which it brings next generation chips to market (short-term differentiator), to repeated success in deploying large scale, curated deployments (medium-term), and expansion of AI expertise up the stack (long-term).**

CoreWeave competes with the following two groups to bring GPU capacity to market: 1) Cloud hyperscalers (such as Amazon, Microsoft, Google, Oracle), and 2) Neoclouds (such as Nebius and Lambda) which are pure-play GPU cloud providers. Since the onset of Gen-AI, CoreWeave has been able to bring the latest generation chips to market at an accelerated pace and with ~20% better performance once online vs peers. This has largely been attributed to an organization-wide focus on speed to market, a strategic navigation of supply constraints (with a notable tie to Nvidia), and the company's Gen-AI centrality, which has brought to market GPU-centric solutions.

Over time, we expect CoreWeave's time-to-market advantage (particularly vs. hyperscalers) to wane as we move toward supply/demand equilibrium. However, CoreWeave's willingness to commit to large, customizable deployments while maintaining its speed to market and reliability will likely drive large deals. Its success over the medium term in winning and deploying next-gen chip architectures, at scale, will be central to building the company's brand awareness, which can lay the foundation around the durability of its market position. In the longer term, CoreWeave may look to extend its AI expertise up the stack. This could help build the peripheral developer ecosystem needed to gain wallet share and further entrench its strategic positioning. We expand on each of these differentiators below.

## Exhibit 5: CoreWeave demonstrates clear differentiation vs. Hyperscaler and Neocloud peers

	CoreWeave	Hyperscalers	Neoclouds
<b>Offering</b>	Deploys highly performant, large-scale, next-gen GPU clusters for specific customer use cases	Largest global footprint of cloud services. Offers AI, cloud, legacy and managed services	Flexible, cost-optimized GPU compute for AI developers and businesses (including SMBs). Offers expanded full-stack AI solutions
<b>Time to Market new GPUs</b>	Procurement, configuration, and deployment of components can occur within weeks - up to 5 months. No owned data centers	GPU time to market likely to improve in 2026 as added capacity comes online. Lean toward full TCO of data centers	Time to market likely to improve in 2026 as added capacity comes online. Have full TCO of data centers
<b>Reliability / Performance</b>	High performance (bare metal, low overhead) driving 20% better utilization rates vs benchmark	Industry-leading reliability for Cloud workloads, offer global redundancy; Build out of GPU infrastructure ongoing	Good performance, less mature, fewer global regions
<b>Size of Clusters</b>	Can scale to hundreds of thousands of GPU clusters, large-scale AI training possible	Largest global compute scale. Constrained GPU availability in near-term	Can scale to tens of thousands of GPUs
<b>Customizability</b>	High for GPU/AI workloads (bare metal, Kubernetes, custom networking). Can take up to 3-6 months from signing to go-live	Broad range of deployments across both direct sales and self-serve. Limited all-around customization of GPU needs.	Limited. Self-serve onboarding leverages pre-built Gen-AI infrastructure, driving ease of onboarding
<b>Target Users</b>	AI/ML researchers, enterprises needing rapid, large-scale GPU access	Enterprises needing global reach, with various cluster needs and compliance, with a broad range of IT workloads	In-house enterprise AI teams, startups, AI labs, cost-sensitive users focused on AI/ML

Source: Goldman Sachs Global Investment Research

- **Time-to-Market:** We expect CoreWeave to be well positioned to continue to be first to market with future generations of top-of-the-line GPUs. For customers who operate at the bleeding edge of AI, and particularly for model trainers, the first weeks of access to AI compute are essential to maintain a competitive edge. This is likely one of the reasons CoreWeave has been able to win large customers like OpenAI. This consistent execution was notable to customers we spoke to during our diligence who expressed that other vendors faced rather consistent difficulties around on-time procurement, configuration, and deployment of GPUs (expanded upon later in this report).

CoreWeave was the first cloud provider to deliver large-scale clusters of Nvidia H100, H200, GH200, and now GB200. This has largely been attributed to: an organization-wide focus on speed, its strategic navigation of supply constraints (with a notable tie to Nvidia) and Gen-AI centrality. CoreWeave holds a sole focus on AI compute vs. hyperscalers who need to balance priorities between traditional and AI workloads.

- **Relationship with Nvidia:** CoreWeave has a symbiotic relationship with Nvidia whereby both companies are customers of the other (CoreWeave exclusively purchases Nvidia GPUs and Nvidia leverages CoreWeave's

capacity). The strategic rationale behind this, which allows Nvidia to have a large, at scale environment to test new GPU capabilities, has also given way to CoreWeave getting priority allocation of in-demand GPUs.

- **Deployment of Large Clusters:** Few cloud providers have deployed the same number of massive clusters as CoreWeave, giving the company unique visibility into the challenges of operating next-gen chips at that scale. More specifically, CoreWeave is unique among the emerging Neocloud peers in its deployment of large scale GPUs for a single customer (up to +100K GPUs). CoreWeave invests in a strong engineering base and proprietary solutions (such observability solutions) that have helped build its expertise in this area. Keeping large clusters of GPUs online for days at a time is a difficult technical challenge, especially as the first deployments of next gen GPU clusters typically see higher failure rates. The failure of just one GPU can disrupt an entire training workload, causing hours of downtime. Similarly, for an inferencing workload, a GPU failure can increase latency significantly. These errors are compounded at scale. Our customer diligence conversations emphasized CoreWeave's strong reliability vs. the downtime experienced with competitors, where they saw daily - if not hourly - failure rates. CoreWeave is able to minimize this level of disruption via the granular visibility it has into some of the largest deployments of next-generation GPUs.

Smaller scale peers do not have the technical expertise to run such complex workloads reliably. This helps CoreWeave perform well on industry benchmarks, with CoreWeave reporting 29x faster H100 MLPerf performance in 2023 than its nearest competitor and a 20% improvement in MFU vs. the benchmark.

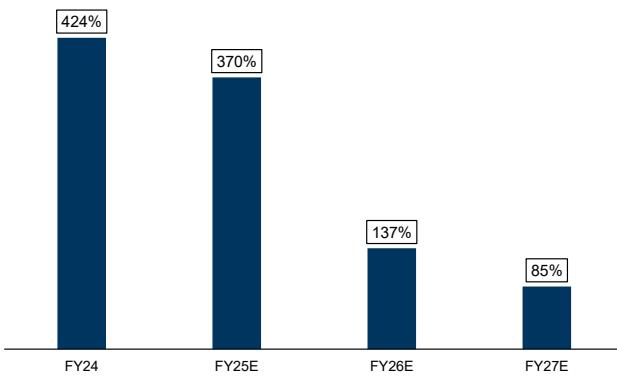
- **Observability tools support higher relative performance vs peers:** CoreWeave's GPU-focused observability tools give CoreWeave a better understanding of how to run large, volatile clusters. This has resulted in a 20% improvement in CoreWeave's MFU utilization vs. comparable benchmarks. An MFU (model FLOPS utilization) is a metric that compares how efficiently a model is using the theoretical maximum compute of given hardware, indicating CoreWeave is able to drive better performance out of the same architecture (also a result of its hardware-level innovations). CoreWeave's access on both system-level and node-level metrics to quickly identify failures gives it the visibility to evaluate potential failure points, improve uptime and drive efficiencies. As failures for large-scale training workloads are often a common occurrence, CoreWeave's automatic detection and remediation of errors are important differentiators. Further, its engineering team is in close contact with customers to proactively flag and remediate any potential issues. To minimize interruptions and failures, CoreWeave runs a Fleet LifeCycle Controller (FLCC) and a Node LifeCycle Controller (NLCC) on the back end. FLCC automates the provisioning, testing, and monitoring of each node while NLCC ensures problematic nodes are replaced with healthy nodes before failure.
- **Better Performance-Adjusted Pricing:** Every component of CoreWeave's architecture is designed to maximize efficiency and minimize latency, including the datacenter layout and components. This includes optimizing rack density,

implementing the latest liquid cooling technology, and securing access to large quantities of power that can support the demands of large-scale training and inferencing workloads. Driving more efficient utilization of its hardware allows the company to charge a similar price as peers, while delivering better value, therefore delivering better performance-adjusted pricing.

□ **CoreWeave's CapEx intensity ahead of peers; Expect further improvements:**

We see CoreWeave's scaling revenue and continued improvement in return on their infrastructure to drive improvements in the company's capital intensity (CapEx/Revenue). CoreWeave has seen a lower capital intensity on its Gen-AI investments ([Exhibit 6](#)) relative to its hyperscaler peers ([Exhibit 7](#)) given its CapEx spend has been more directly tied to revenue (with a spread of ~3-6 months). The time gap for hyperscalers has been much longer (reaching years) as they needed to allocate a majority of their CapEx investments towards long-duration building blocks (i.e. land) in order to expand their Gen-AI-centric, GPU-based architectures. We expect CoreWeave to continue to outpace that of peers as they drive improvements via scaling infrastructure, increased efficiencies and expansion of software-related revenue (expanded upon further in the report).

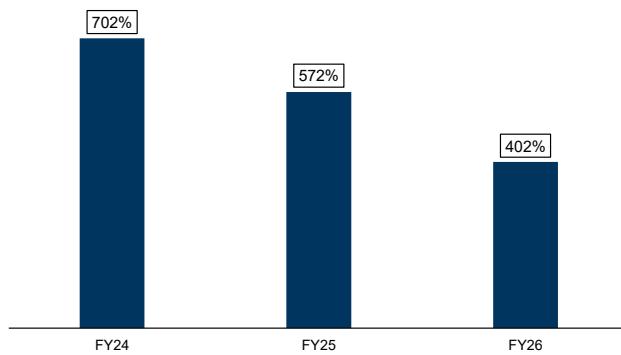
**Exhibit 6: Expect CoreWeave's capital intensity to contract meaningfully as revenue scales**



Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 7: Microsoft's AI CapEx intensity**

Microsoft's Azure AI CapEx Divided by AI Revenue (GSe)



Source: Company data, Goldman Sachs Global Investment Research

■ **Committed contract dynamics give way to customized deployments, predictable ROI:**

While many of CoreWeave's hyperscaler peers are building broad, generalizable AI architectures that can be used for a variety of use cases, CoreWeave is more flexible for its largest customers and willing to build customer-specific hardware. For example, other AI clouds may prioritize the flexibility of their architectures, and are therefore reluctant to take on a large single-customer commitment with specialized configurations. CoreWeave, however, embraces such opportunities - as seen with the signing of the +110k GPU commitment signed with OpenAI.

This is enabled by the unique structure of CoreWeave's contracts. Rather than purchasing and building the infrastructure ahead of demand signals, CoreWeave waits until a contract is signed before purchasing their infrastructure. With a ratable revenue model, healthy 2.5 year cash payback period over their average 4-year

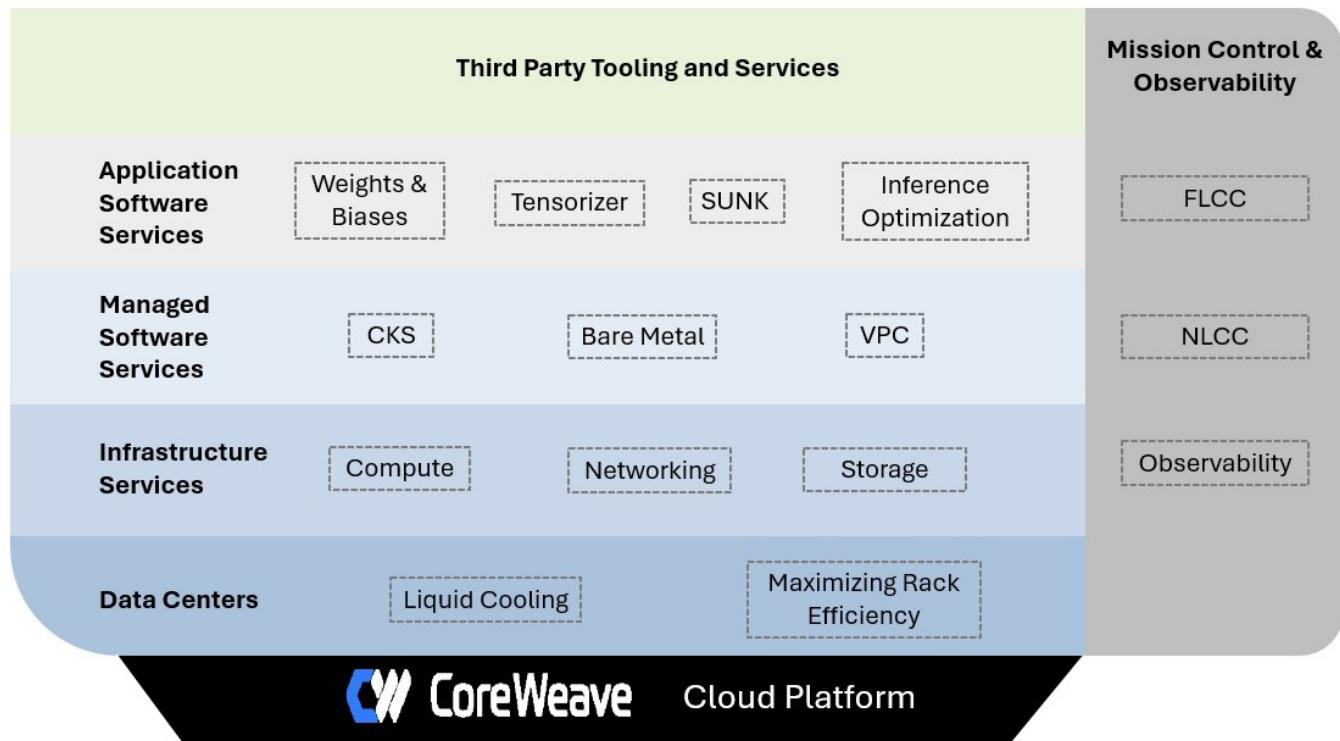
contract length, CoreWeave has been able to drive strong visibility and an attractive ROI profile from such contracts. Hyperscalers, in contrast, generally build customer agnostic architectures that are monetized by their more consumption-oriented pricing models.

■ **AI specialization can be used to flesh out Gen-AI-centric solutions up the stack:**

Mirroring prior tech paradigm shifts, we believe the technological shift to Gen-AI will bring on new tech stack functionalities. RAG, Vector Search, Vector Databases, and ML ops are examples of these emerging needs. Similarly, CoreWeave has brought to market the below solutions to meet the specific needs of AI workloads. If CoreWeave can successfully identify the opportunities on its platform that can draw buy-in and engagement, this may drive longer-term market leadership. Its build out of a GPU-optimized observability solution, as well as the below offerings, suggest CoreWeave may take on this goal. We acknowledge the nascence of where CoreWeave is today and look for consistent execution to be displayed against a long horizon before we determine whether CoreWeave was able to capture this long-term differentiation.

- 1) **Slurm on Kubernetes** gets customers the ability to deploy Slurm (a job scheduler for high-performance parallel workloads, co-locating different AI workloads on the same cluster) on top of CoreWeave's Kubernetes (an open-source tool to deploy and manage applications);
- 2) **CoreWeave Tensorizer** allows customers to quickly load models into GPU memory. This has accelerated average model load times by 1.7x and 1.4x faster than the two alternatives against which CoreWeave has benchmarked itself. This is achieved by compartmentalizing job tasks so the infrastructure can run through numerous smaller tasks simultaneously instead of overloading the capacity by aiming to load one large data set at once;
- 3) **Inference Optimization Services** which will help customer right-size and more accurately provision their workloads; and
- 4) **The recent acquisition of the AI developer platform Weights & Biases**, which expands CoreWeave's application layer functionality, allowing customers to use their infrastructure more effectively. Weights & Biases helps organization build, deploy, and monitor AI applications, with visibility into how models are performing and what kind of traffic they are seeing.

## Exhibit 8: CoreWeave's tech stack



Source: Company data, Goldman Sachs Global Investment Research

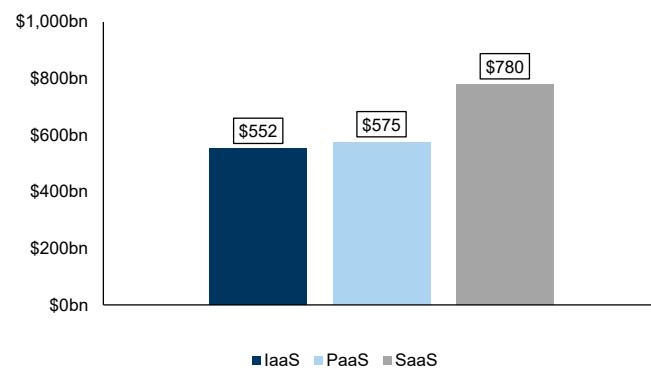
## Top Line Durability: AI Market Outlook

Gen-AI is at the center of CoreWeave's market opportunity as the company provides access to a modernized tech stack configured specifically for the unique technical requirements of AI workloads. This positions CoreWeave to benefit from the proliferation of AI model training and inferencing as more use cases come to market. We estimate that Generative AI could catalyze software revenues (across IaaS, PaaS and SaaS) to \$2 trillion by 2030 (Exhibit 9), with Gen-AI potentially garnering 10-15% of that spend. This presents an opportunity for cloud providers such as CoreWeave to power the infrastructure that will support these applications. CoreWeave estimates its TAM will reach \$399bn by 2028, underscoring the potential strategic nature of providing the underlying platform to Gen-AI. This can be broken down as \$330bn from AI training infrastructure (i.e. servers, storage, training costs), \$49bn from inferencing infrastructure, and \$20bn from workload monitoring.

While we acknowledge that the industry has evolved rapidly since we published our TAM estimate in August 2024, the uncertainty in the macro environment limits the visibility of upside at this point in time. While companies are likely to prioritize AI investments and innovations, in the near term, the magnitude may be softer than what might otherwise have been deployed in a more stable environment (which has historically been conducive to higher spending).

**Exhibit 9: GS expects AI TAM to reach nearly \$2tn by 2030; Gen-AI could garner 10-15% of spend**

2030 AI TAM Across IaaS, PaaS, and SaaS



Source: Goldman Sachs Global Investment Research

### **Proliferation of AI to support infrastructure needs**

#### **AI Infrastructure demand durable as AI set to percolate to enterprise over time.**

Across the market, we continue to see evidence of AI moving from early AI product trials to larger-scale deployments. This can drive robust and sustainable demand for AI infrastructure. Early AI adoption is likely to be most evident at the consumer and prosumer layer before eventually percolating to the enterprise (see: *GS Private Company Conference: Software, Internet*). As CoreWeave is an AI-centric infrastructure provider, the company stands to benefit from enterprises expanding the scope and complexity of AI deployments over time. This is still in early innings, with demand outpacing supply.

Adoption in the application side is also still nascent, with Microsoft seeing Copilot deployment start off small before gradual expansion across organizations. Still, utilization is growing, with a 60% increase in usage intensity reported last quarter.

- **Current state of the industry sufficient to propel Gen-AI.** As we discussed in our 2025 Software Outlook, we do not see AGI as necessary for Gen-AI progression and permeation across the enterprise, therefore leading us to expect the current state of the industry being sufficient to drive the expected growth in the number of Gen-AI apps. The momentum in which Gen-AI applications (i.e. Salesforce's Agentforce, Microsoft Copilot, Adobe Firefly, etc.) are improving with current models and can leverage newer models underwrite our view.

**CoreWeave use cases agnostic to training vs inferencing needs.** While CoreWeave rents its infrastructure agnostic to the ultimate use case, the company generally has visibility into how customers are using its GPUs. CoreWeave estimates the mix between inferencing/training is split relatively evenly today. This is determined by evaluating the utilization of the hardware, in which CoreWeave has visibility into. Training, for instance, generally displays consistent power demand across any given period. Inferencing, in contrast, generally shows more variable power demand as usage as it is based on product usage.

- **Need for continuous training supportive of CoreWeave's offerings.** Training has been an instrumental part of Gen-AI to date as it underpins a foundational model's ability to understand context and drive accurate model results. We expect this process to continue to hold technological and market relevance - though across a smaller number of companies - given: 1) the exponential growth of data (across the web and within specific organizations), 2) the ongoing innovation in the industry (that is adding modalities and parameters to future models) and 3) the continuous nature of learning and reinforcement needed for these models. The compute-intensive nature of this exercise is likely to support demand for CoreWeave's value proportion as the company can provide the large scale GPU clusters needed for this type of process. The range of configuration needs across different training approaches (i.e. more densely packed data centers, liquid cooling, etc) also feed into CoreWeave's strategy as the company is one of a handful of providers willing to custom build data centers.

- **Mix-shift to inferencing presents more uncertainty, may drive capacity needs.** Greater proliferation of AI is likely to drive exponential growth in the engagement with and - therefore - the usage of AI infrastructure. This presents an opportunity for cloud providers to build out the infrastructure that can be used to power these applications, leading the hyperscalers into an elevated CapEx cycle (Gen-AI Deep Dive, Data Center Supply/Demand, MSFT CapEx Deep Dive). For CoreWeave this may increase customers' need for infrastructure, as they start supporting more data pulls. As this consumption can scale rapidly, customers may choose to expand their spend with CoreWeave as the company provides better performance-adjusted pricing (See prior section). This market dynamic can then drive additional investments in models as demand for such solutions grows. A shared infrastructure between training/inferencing processes may also result in improved feedback loops between

the two workflows. On the other hand, if inferencing needs are largely standardized across providers (leading to lower differentiation) and customers look to diversify across multiple clouds, the growth in inferencing may pose a headwind to CoreWeave.

### **Market composition highly concentrated today, likely to bifurcate over time**

As an emerging technology, the AI market has a high concentration across a relatively small number of scaled players. This holds true as it pertains to the number of AI-native customers in the market (OpenAI, Anthropic, etc) as well as those providing the necessary technology to enable AI-native solutions (i.e. hyperscalers). We expect the market to evolve over time in a similar manner to what occurred during the Cloud build out cycle.

- On the customer side, the market is likely to expand over a long period of time with new companies emerging on the back of broader adoption. This can broaden the scope of the market and present an opportunity for CoreWeave to diversify its customer base.
- As it pertains to the enablers of this market, we expect the prominence of a small group of providers to be a lasting dynamic. The infrastructure for Gen-AI has similar characteristics to the cloud in the fact that it requires high capital intensity. This creates a high barrier to entry which we do not expect many companies to be able to reach. This also adds to the uniqueness of CoreWeave's success to date, and suggests it can remain a key player over time.

### **Progressing on path to supply/demand equilibrium**

#### **Supply constraints of key data center components has been the main limiting factor of industry growth.**

The rapid pace of innovation and desired adoption of Gen-AI have outpaced the constrained production capabilities of GPUs, bringing power online and building out data centers (the main components of AI-based data centers). This was the industry's biggest challenge to meeting customer needs and driving additional growth. CoreWeave was able to benefit as a relatively new player as it had access to these needs. Since CoreWeave acquires the needed GPUs on a just-in-time basis based on customer contracts, the company's close supply chain relationships have helped CoreWeave insulate itself from the full extent of these constraints.

- With 1.3GW of power contracted and operating across 32 data centers with over 250K GPUs, CoreWeave currently has enough capacity to meet existing commitments (with subsequent agreements for 70 MW from a single landlord and with Galaxy Digital for 133 MW, per Galaxy's 4Q release). However, given three suppliers accounted for 46%, 16% and 14% of purchases in FY24, any strains of any these relationships could have a material effect on CoreWeave's operations. Should bottlenecks or competition increase, there me added risk to revenue conversion of CoreWeave's backlog and/or its ability to drive future growth.

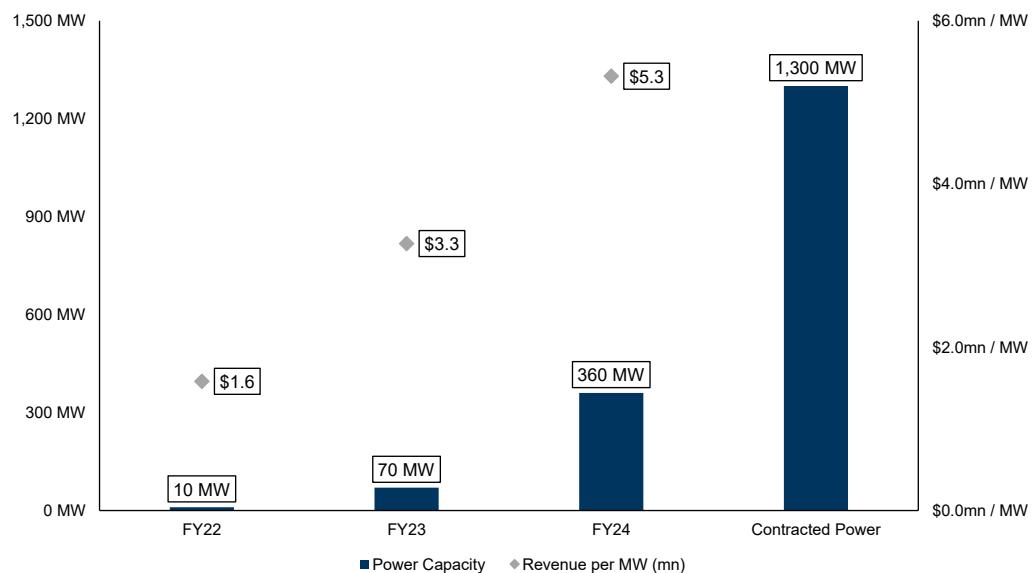
#### **Improving visibility around alleviation of some constraints; CoreWeave likely to lean into key differentiators, efficiencies to drive share gains.**

Our teams (across

Data Center, Utilities, Sustainability) have started to forecast an alleviation of constraints in data center capacity though forecast an unrelenting need for power demand through 2030. GS' Data Center model (run by Jim Schneider) expects loosening constraints may occur earlier than expected, with peak datacenter occupancy likely to occur in 2025 (vs expectations 2H26 prior). Given that time to market has been a key differentiator for CoreWeave, a more balanced market will likely accelerate the need for the company to lean in and solidify the other components of its value proposition (GPU-focus, improved reliability, dedicated infrastructure, software, etc).

Meanwhile, our Utilities and Sustainability teams (led by Carly Davenport and Brian Singer, respectively) expect a ~160% increase in global datacenter power demand by 2030 vs. 2023 levels to be met with moderating power intensity. CoreWeave will need to demonstrate continued ability to source power and drive internal efficiencies ([Exhibit 10](#)) to meet industry demand.

**Exhibit 10: CoreWeave has a demonstrated history of scaling power capacity and driving higher price per MW**



As of December 31, 2024

Source: Company data, Goldman Sachs Global Investment Research

**Long-term evolution likely to show resemblance to Cloud build out**

We believe the evolution of the AI ecosystem is likely to resemble the buildup of the cloud ecosystem, with value first accruing to the Infrastructure layer, then the Platform layer, and finally Applications (outlined in our [Gen-AI Deep Dive](#)). However, we see this cycle scaling faster and more efficiently at the Infrastructure layer (with CapEx efficiency already scaling ahead of the cloud buildup). Mirroring prior tech paradigm shifts, we believe the shift will require new ways of thinking about the tech stack and evolution of platform solutions to unlock unexpected applications. Gen-AI is already showing this potential, as foundational models, RAG, Vector Search, Vector Databases, and ML ops are emerging.

We believe the opportunity for Infrastructure-level players to partake in building out the

AI ecosystem is large. We refer back to AWS at the beginning of the cloud cycle and how its tech-forward platform was able to drive customers to choose and scale on its cloud platform. We are observing several Neocloud providers taking different approaches with the goal of yielding similar results. If CoreWeave can successfully identify the opportunities on its platform that can draw a compelling level of buy-in and engagement, this could drive longer-term market leadership. Its successful start building a GPU-optimized observability solution offers very early proofpoints that CoreWeave may take on this strategy. The acquisition of Weights & Biases is another early step that may help with software monetization strategies. We acknowledge the nascence of where CoreWeave is today and look for consistent execution to be displayed against a long horizon before determining the leading ecosystem for AI companies.

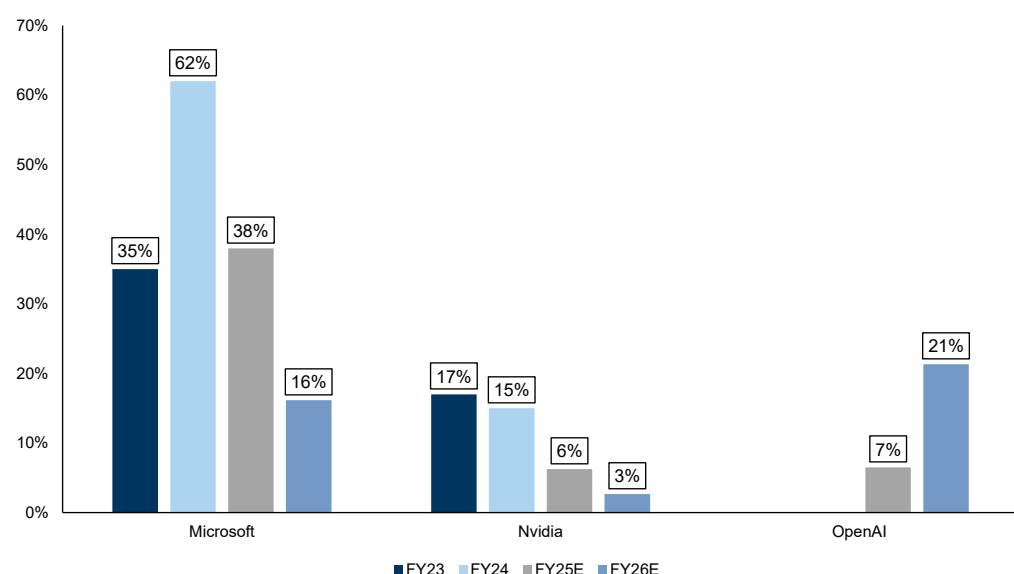
## Top Line Durability: Customer diversification likely to take time

### **Customer concentration, contract composition contrasts that of broader software.**

CoreWeave has significant customer concentration, with a substantial portion of revenue driven by a limited number of customers. Most prominently, OpenAI is positioned to be the company's largest customer as the +\$11bn deal announced on March 10th is the largest in CoreWeave history. Prior to that (and as of Dec 2024), Microsoft was the largest customer, comprising 62% of revenue in 2024 (vs. 35% 2023), with Nvidia the company's second largest customer (comprising ~15% of revenue). The company's top two customers accounted for 77% of revenue in FY24, 56% in FY23 and 29% in FY22.

The asset-heavy nature of CoreWeave's business, that requires net new contracts to be signed and new capital to be deployed whenever a customer comes off of an existing contract or needs additional capacity, adds to the risk investors may attribute to CoreWeave's model. These dynamics heavily contrast to software peers, whose solutions can be scaled rapidly at an incremental margin and whose contracts can be easily renewed, with expansion clauses typically embedded.

**Exhibit 11: Expect revenue concentration among a few large customers in short to medium term**  
Large Customers as a % of GSe CoreWeave Revenue



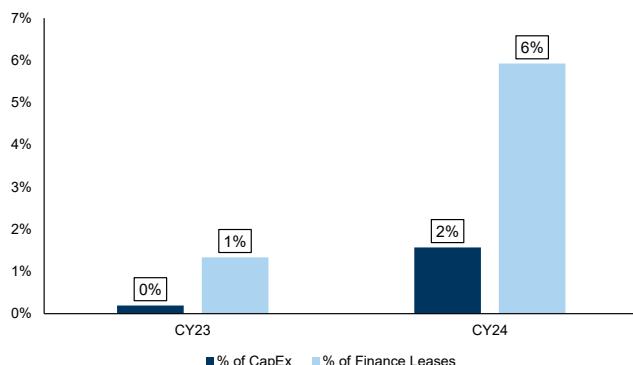
FY25/FY26 Microsoft estimates assume revenue grows from \$1.2bn in '24 but are below OpenAI's expected revenue contribution. FY25/FY26 Nvidia estimates assumes revenue \$ spend is flat. FY25/FY26 OpenAI revenue is GSe.

Source: Company data, Goldman Sachs Global Investment Research

- **CoreWeave wallet share still below 10% within largest customers.** While CoreWeave's customers comprise a large portion of its revenue, CoreWeave still holds relatively nascent wallet share. For example, even within the company's largest customers, CoreWeave represents less than 10% of either Microsoft's ([Exhibit 12](#)) or Nvidia's CapEx investments([Exhibit 13](#)). This leads us to believe the company still has room to grow its presence within existing customers.

**Exhibit 12: CoreWeave represents <10% of Microsoft's CY23/24 CapEx**

CoreWeave's Revenue from Microsoft as a % of Microsoft's CapEx and Finance Lease Spend

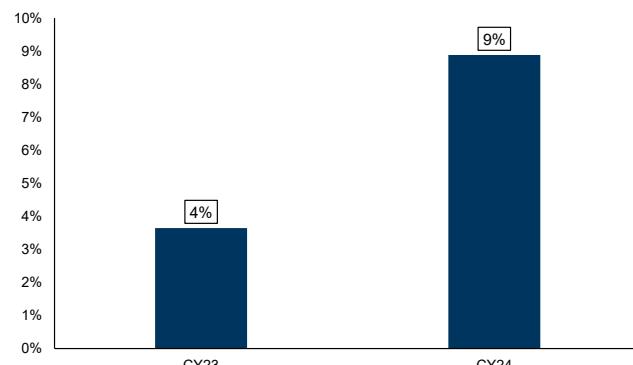


MSFT CapEx/Finance Leases are calendarized. Disclaimer: If MSFT spend on CoreWeave were part of the company's joint research agreement with OpenAI, it would not fall under either CapEx category.

Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 13: CoreWeave represents <10% of Nvidia's total CapEx spend**

CoreWeave's Revenue from Nvidia as a % of Nvidia's CapEx Spend



Source: Company data, Goldman Sachs Global Investment Research

### GTM to shift to direct sales and product-led growth from light-touch approach

**Engineering focus helped build strategic customer relationships.** CoreWeave has historically prioritized AI labs (Cohere, Mistral) and AI-focused enterprises (Microsoft, Nvidia, Meta, IBM) in terms of GTM. The focus was on securing the workflows of the handful of companies that are driving the innovation and enablement of Gen-AI. This would lead to scalable opportunities that can drive increased wallet share. CoreWeave largely leveraged an engineering-heavy strategy to close deals, with a strong supply/demand imbalance present in the market that was conducive to a light-touch sales motion (non-GAAP S&M intensity <1% in 2024). This allowed the company to establish deep technical partnerships with their customers that can drive longer relationships. While building its operating history, this also yielded consistent expansion activity with its customers.

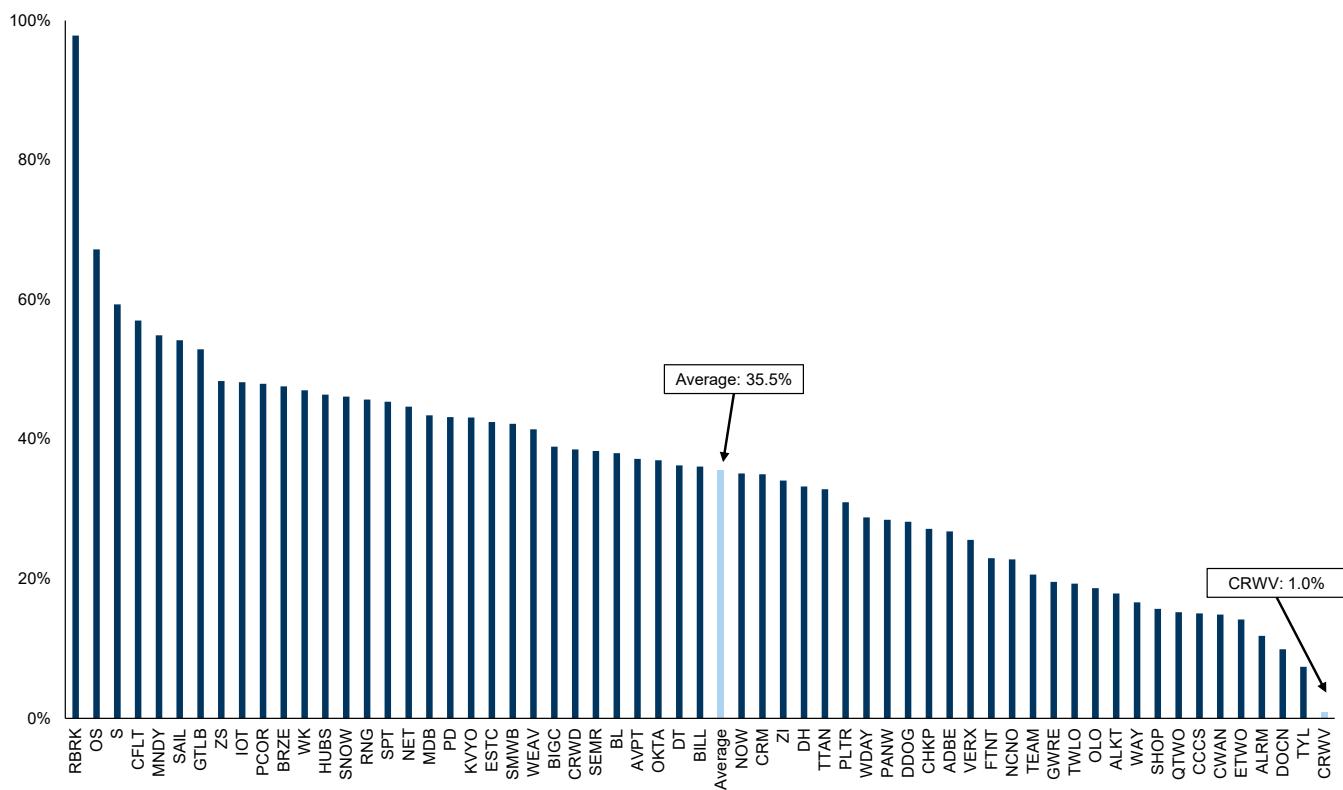
**GTM likely to shift towards outbound selling motion as market expands.** Looking ahead, we see the expansion of CoreWeave's GTM to a broader range of customers as a key pillar of its growth strategy. As Gen-AI use cases scale, more Gen-AI businesses may emerge, with a preference toward Gen-AI-forward platform solutions. Additionally, customers who may have an indirect usage of CoreWeave may begin to look to build a more direct relationship with the company to reduce costs and/or increase the control and reliability of their applications.

Rolling out dedicated solutions and partnerships is expected to help attract these customers to CoreWeave more directly. The company is also verticalizing its sales teams, establishing regional expansion teams, building partnerships and actively scaling its sales leadership. We see this blending a direct sales strategy with product-led growth, vs the (almost exclusively) light touch and engineering-led approach today. CoreWeave plans to invest in scaling these initiatives over time, driving Sales and Marketing (non-GAAP) spend to increase to \$232mn in FY27 (GSe) from \$15mn in FY24.

- CoreWeave has been able to benefit from rapid revenue growth and large deals that didn't require a large sales motion. With Sales and Marketing spend expected to account for less than 1% of revenue in revenue (vs. 36% for our broader coverage, [Exhibit 14](#)), despite 8x in revenue, its evident how the large demand pool and supply constrained environment drove customer spend toward CoreWeave.

**Exhibit 14: CoreWeave S&M intensity is lowest across our software group**

TTM S&M Expense (GAAP) as a % of TTM Revenue



CRWV non-GAAP S&M intensity was 0.8% in 2024

Source: Company data

**Weights & Biases acquisition can expand buy-in from broader scope of personas.**

The acquisition of Weights & Biases builds on the effort to diversify CoreWeave's customer base. Weights & Biases serves over 1,400 organizations, across AI labs and enterprises, as well as over 1mn AI researchers. Though much smaller in revenue scale, CoreWeave has an opportunity to cross sell across their respective customer bases. We note that the current user persona of Weights & Biases is not typically that of the infrastructure buyer. As an AI developer platform to build agents, applications and models, however, its users have a stake in which infrastructure is used. Should CoreWeave integrate the technologies and user experience well, it can build buy into CoreWeave's cloud platform.

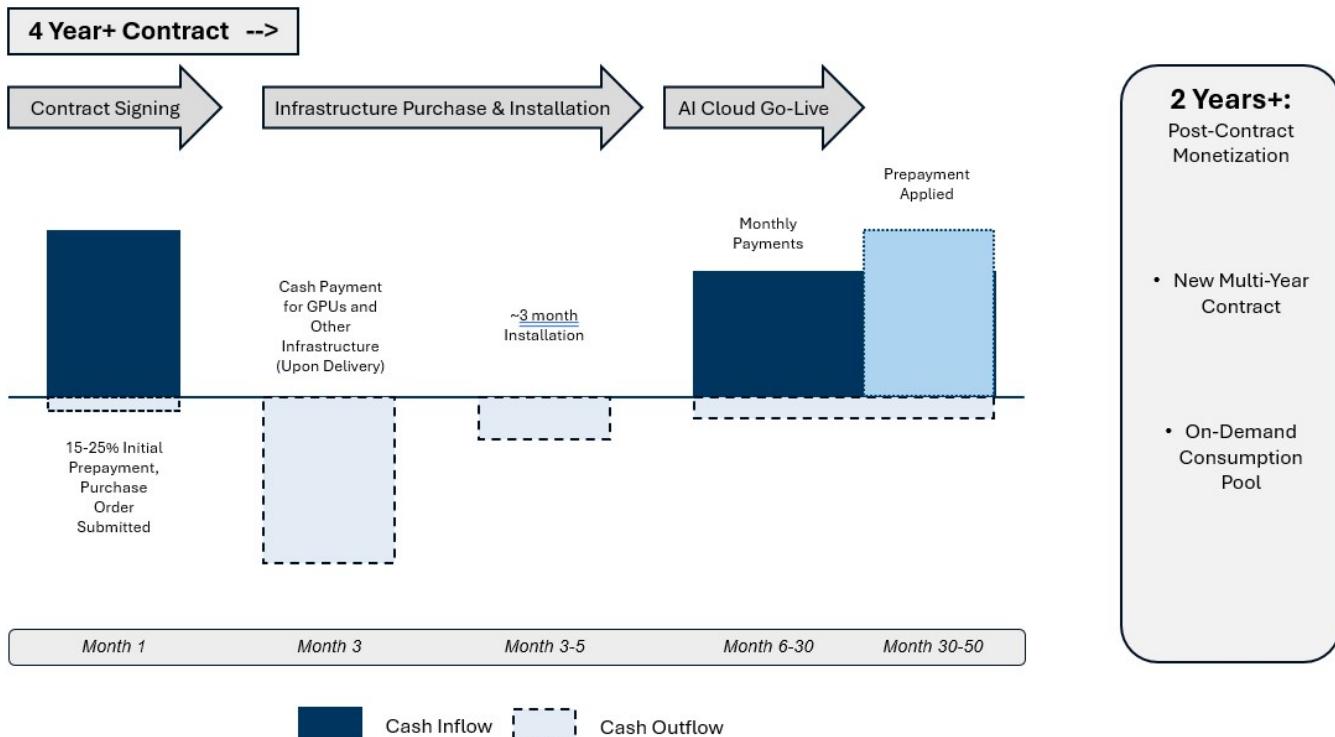
## Top Line Durability: Contract construction drives visibility

CoreWeave has strong revenue visibility given the committed nature of its multi-year contracts and the associated ratable revenue model – a dynamic we do not expect to change even as its customer base diversifies over time. Unlike general-purpose clouds, CoreWeave is able to garner this commitment given the access it provides to dedicated hardware configured specifically for the unique technical requirements of that customer. Other cloud providers typically build out their infrastructure in a customer-agnostic way that creates a pool of capacity to be utilized across different customers or applications. This approach fundamentally offers more flexibility, giving room to manage costs as needed.

Expected utilization of CoreWeave's provided infrastructure is agreed upon prior to the deal signing, which determines the implied cost per token. Though ultimate utilization is up to the customer once the infrastructure goes live and is theirs to use for the full duration of the contract. This allows CoreWeave to generate over 95% of its revenue from committed contracts, vs consumption-dependent comps.

We outline the dynamics of CoreWeave's contracts in [Exhibit 15](#) and discuss each phase of the process in each of the paragraphs below.

- **At contract signing.** CoreWeave typically receives 15-25% of the Total Contract Value (TCV) up front with the remaining recognized ratably over the duration of the contract. CoreWeave leverages the upfront payment to purchase GPUs and other infrastructure components. The pre-payment sits within deferred revenue, while the company's Construction in Progress (CIP) CapEx metric signifies the costs related to GPUs before it goes live. The installation period typically takes 3-6 months.
- **Once the infrastructure is live.** CoreWeave starts recognizing revenue on a ratable basis (excluding the upfront payment). This typically occurs at a 1-2 quarter lag from signing as it takes time to procure and deploy the GPU cluster. CoreWeave typically has a 2.5 year cash payback period over their weighted-average 4-year contract duration. Contract lengths range from 2-5 years. The customers' pre-payment amount is also applied to the tail end of the contract.
- **After the contract expires.** CoreWeave monetizes the infrastructure by either signing a shorter-duration contract (~1 year) or selling the compute on an on-demand basis. This gives access to customers who want to maintain their workloads on the same GPUs or need to quickly scale up a short burst of compute. Given the GPU's economic value may deteriorate after four years, CoreWeave does not want to commit to providing that infrastructure for longer periods of time. On-demand revenue is billed monthly in arrears on a pay-as-you-go basis for hourly usage of the platform.

**Exhibit 15: Example of a typical contract lifecycle with CoreWeave**


This exhibit is meant to serve as a general representation of the company's contracted business, but may not be applicable to every contract. Not drawn to scale.

Source: Company data, Goldman Sachs Global Investment Research

**End-of-contract or infrastructure needs typically trigger new contract wins.**

CoreWeave's contract renewals also do not follow that of the typical SaaS company.

Given the supply/demand imbalance currently in the market and the asset-heavy nature of the business, each additional need for a customer triggers a new contract to be signed. This contrasts the asset-light nature of software that can be deployed at scale in a rapid manner, with minimal incremental costs and follows a 4Q-weighted renewal cycle.

CoreWeave contracts are based on dedicated infrastructure provided, therefore relying on the useful life of the physical components underlying the contract. As the economic value of these GPUs (currently 6 years) is more unpredictable toward the end of its useful life, CoreWeave typically leans on an average contract duration of 4 years.

Customers whose committed contracts are expiring either: 1) sign a new contract and move to a new infrastructure, or 2) leverage the existing infrastructure on a pay-as-you-go, on-demand basis. The on-demand basis relieves CoreWeave from its obligations to provide continued uptime while continuing to monetize the underlying assets.

- **Customers often opt for more capacity, even amidst ongoing contracts.** As companies continue to deliver on their product roadmap, additional capacity is often procured, often leading to several commitments associated with one customer. Expansion in contracted commitments typically happens through three different vectors:

- a. A customer bought compute capacity, it worked well, and they now need more as they expand and scale the deployment of that use case.
- b. A customer bought capacity in one geography and now needs capacity in an additional geography.
- c. A customer bought capacity for one generation of silicon and now needs capacity for the next generation of silicon.

For these reasons contract expansions do not necessarily happen at a certain stage in a given contract, but are rather done on an as-needed basis. Notably, three of CoreWeave's top five committed contract customers had signed agreements for additional capacity within 12 months of their initial purchase dates, growing their spend roughly 4x vs. their initial contract value (an increase of \$7.8bn). As laid out in the prior section, CoreWeave still holds relatively nascent wallet share. For example, even within Microsoft and Nvidia (two of CoreWeave's largest customers), CoreWeave spend represents less than 10% of CapEx spending.

**On demand revenue likely to grow as more cohorts move through expiry cycle.** As the majority of CoreWeave's revenue was added in FY24, the company has yet to go through a meaningful end-of-contract cycle. For this reason, we expect the base of on-demand revenue to still scale from current levels. As larger cohorts of customers move past their committed contracts, there is likely to be a layering effect to the company's on-demand revenue base. Customers are likely to use their infrastructure until the GPUs fail, which typically occurs after six years. CoreWeave has noted that it has customers that still leverage Nvidia Ampere-based infrastructure, which was released in 2020. Still, we do not expect the mix of revenue to materially shift (with On-Demand accounting for ~5% of revenue today) as CoreWeave is likely to grow committed contracts at a more rapid pace.

## Exhibit 16: Mechanics of CoreWeave's committed contracts and on-demand

	Committed Contracts >95% of Revenue	On-Demand <5% of Revenue
Pricing & Contract Mechanics	<ul style="list-style-type: none"> <li>Primarily take or pay contracts</li> <li>Price on \$/GPU/hr based on type of infrastructure and contract duration</li> <li>Weighted avg. contract duration ~4 years</li> <li>\$27bn in remaining performance obligations (inclusive of OpenAI)</li> </ul>	<ul style="list-style-type: none"> <li>On-demand compute capacity consumed on an hourly basis</li> <li>Priced on \$/hr/GPU basis based on type of infrastructure</li> </ul>
Payment Terms	<ul style="list-style-type: none"> <li>Weighted avg. prepayment typically 15-25% of contract</li> <li>Generally upfront prepayment received post-signing offsets payments due at the end of the contract</li> <li>Prepayment may be applied to the first part of the contract</li> </ul>	<ul style="list-style-type: none"> <li>Payments due at end of billing period</li> </ul>
Billing	<ul style="list-style-type: none"> <li>Generally billed monthly in arrears</li> <li>Interest earned on prepayments recognized as interest income on the I/S</li> </ul>	<ul style="list-style-type: none"> <li>Generally billed in arrears based on usage</li> </ul>
Revenue Recognition	<ul style="list-style-type: none"> <li>Recognized ratably on a monthly basis throughout the contract duration</li> </ul>	<ul style="list-style-type: none"> <li>Recognized upon usage</li> </ul>

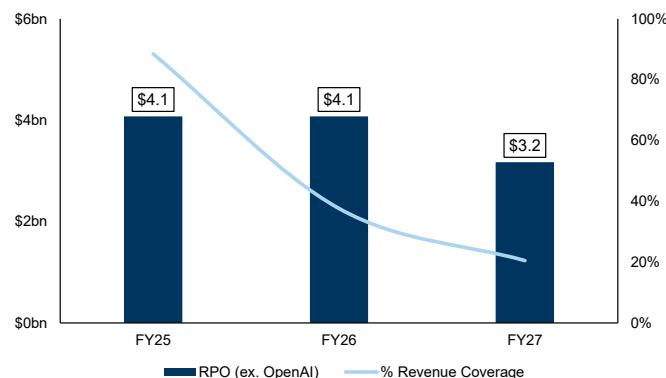
Source: Company data, Goldman Sachs Global Investment Research

**Assessing RPO coverage given ratable rev rec provides high degree of visibility:**

Based on the contract dynamics laid out above, we assess CoreWeave's committed RPO in relation to our revenue projections. This evaluation helps us understand how much of future revenue is secured and how much remains to be signed. To do so, we break this down into two main buckets:

- **Estimate \$15.1bn in RPO (excluding OpenAI deal closed in March) to convert to revenue and account for 88%/38%/20% of GSe in FY25/26/27.** CoreWeave expects to recognize 54% of their RPO in the next 24 months. If we estimate a roughly 50/50 split, the RPO flow through to revenue would account for 88% of revenue in FY25 and 38% in FY26 revenue. We note that this may be overestimating the contribution in FY25 and underestimating FY26 coverage as a portion of the RPO base may be related to Blackwell chips, which are only expected to come online closer to the back half of FY25, and therefore driving a more meaningful percentage of revenue in FY26. CoreWeave also expects to recognize 42% of their RPO in months 25-48 (FY27 and FY28). We apply the same conservative estimate that the 42% of RPO will be recognized equally across FY27 and FY28, leading us to estimate at least 20% revenue coverage in FY27 from existing RPO.

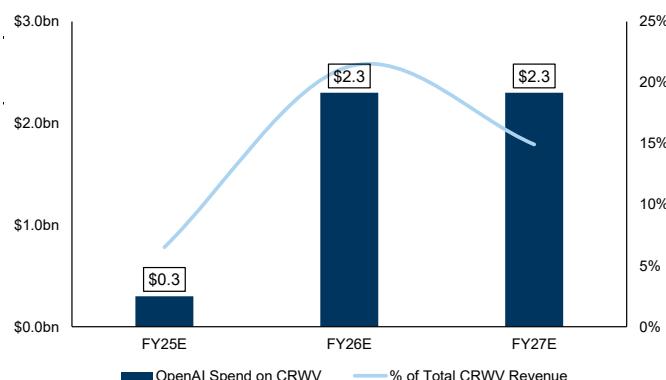
**Exhibit 17: We estimate CoreWeave's RPO coverage (ex. OpenAI) 88% in FY25, 38% in FY26 and 20% in FY27**



Source: Company data, Goldman Sachs Global Investment Research

- **OpenAI's \$11.9bn commitment, announced March 10, 2025, likely to drive \$2.3bn/year. We also set out to determine the contribution of CoreWeave's largest deal to date.** Announced in March (1Q25), this was not included in the last reported RPO figure CoreWeave provided us (\$15.1bn). We also do not expect material revenue flow through from this deal in 2025. Given the scale of the deployment, we expect the procurement, installation and go-live of the infrastructure to span several timelines over six months, with full go-live only likely at some point in 4Q. Based on this, we only account for ~\$300mn of revenue flow through in 2025. It is also important to note that CoreWeave has highlighted that \$11.55bn of the \$11.9bn will be recognized in RPO. The remaining \$350mn is recognized as an equity investment in OpenAI it is making in CRWV. This also suggests no upfront prepayment made by OpenAI (which is typically recognized in deferred revenue). Dividing this \$11.55bn by the contract duration of 5 years suggests an annual runrate of \$2.31bn to be recognized in each year until expiration. Assuming 2026 will be the first full year will recognize this revenue, we assume this deal will account for 21% and 15% of 2026/2027 our revenue expectations ([Exhibit 18](#)).

**Exhibit 18: We estimate OpenAI spend to be 7% of CoreWeave's Revenue in FY25, 21% in FY26, and 15% in FY27**

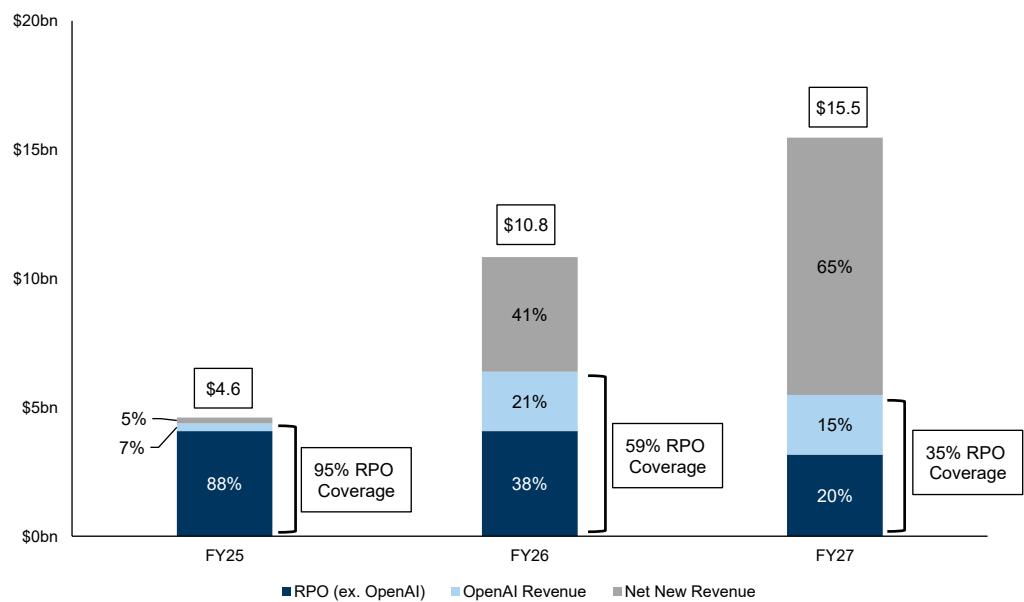


Source: Company data, Goldman Sachs Global Investment Research

**Collectively, committed deals expected to account for 95%/59%/35% of GSe in 2025/26/27; Remainder can be driven by mix of existing customers and net new wins.** In aggregate (Exhibit 19), we assume that CoreWeave's RPO commitments account for 95%/59%/35% of our revenue expectations for 2025/2026/2027. While a portion of the remainder will likely be a function of adding net new customers, we also outline the various avenues in which existing customers can drive a higher portion of revenue than accounted for in the above analysis. This can occur via:

- 1. Customers continuing to leverage existing infrastructure on a pay-as-you-go, on-demand basis.** As noted above, customers are likely to use their infrastructure until the GPUs fail, which typically occurs after six years. As CoreWeave's average contract lifecycle is four years, CoreWeave has noted that they have customers still leveraging their Ampere-based infrastructure, which was an Nvidia GPU released in 2020. As larger cohorts of customers move past their committed contracts, there is likely to be a layering effect to the company's on-demand revenue base. While the economic value of the asset continues, the alleviation of the commitment to support this infrastructure relieves CoreWeave from its obligations to provide continued uptime should the chip fail to run. This has not yet been a meaningful driver for CoreWeave as the company drove exponential growth over the last two years and has yet to move through a large expiry base of contracts.
- 2. Customers need additional capacity.** Expansion in contracted commitments typically happens through three different vectors: 1) A customer bought compute capacity, it worked well, and they now need more capacity, 2) A customer bought capacity in one geography and now needs capacity in an additional geography, and 3) as customers need to move off of last-gen chips, they may need to move the workloads to the next generation of silicon. This is fairly common, with three of CoreWeave's top five committed contract customers having signed agreements for additional capacity within 12 months of their initial purchase dates, growing their spend roughly 4x vs. their initial contract value (an increase of \$7.8bn).

**Exhibit 19: Including OpenAI, we estimate RPO coverage of 95% in FY25, 59% in FY26, and 35% in FY27**



Source: Company data, Goldman Sachs Global Investment Research

## Top Line Durability: Revenue as a function of improving ROA

While revenue can be evaluated as a function of RPO, we look to get ahead of deal signings by evaluating the trajectory in which CoreWeave can improve its return on its net assets. Given revenue today is a byproduct of how CoreWeave can monetize its fleet of AI infrastructure, we discern the dynamics that can yield an improvement in the return profile of the company's spend. Net Assets is the metric we see best representing the portion of CoreWeave's infrastructure that is online and being utilized, net of depreciation. The average of the company's Net Assets is based on the average over the past two quarters, to account for timing of capacity becoming available. Net Assets are calculated by subtracting PP&E Depreciation from Gross PP&E (on the Balance Sheet). This takes Gross Assets, which represents the total monetizable assets (without accounting for depreciation) and subtracting Construction in Progress – the portion of PPE that has not yet come online, and is, therefore, not yet contributing to revenue.

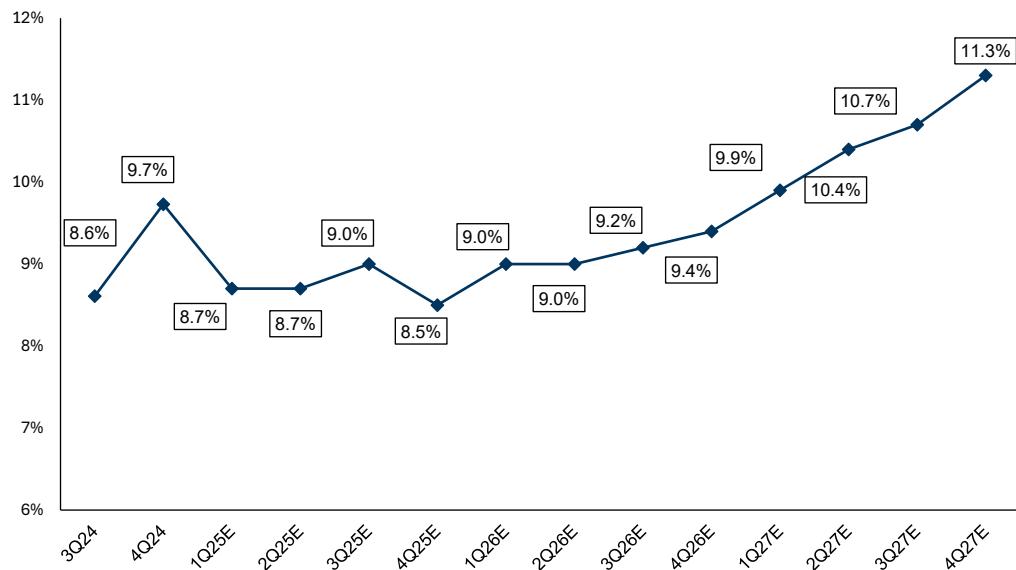
We expect CoreWeave to show gradual improvement against the return on Average Net Assets, with the ratio expanding to 8.7%/9.1%/10.6% (on average) in 2025/2026/2027, from 9.7% at the end of 2024 ([Exhibit 20](#)). The key drivers of expansion to the company's average net asset ratio are:

- **Growing expiry base:** When customers reach the end of their committed contract with CoreWeave, they may opt to continue to utilize the underlying infrastructure on a pay-as-you-go basis. Given CoreWeave can't guarantee the utilization of these chips past the predetermined duration, this spend is not committed. We therefore do not have visibility into the utilization and contribution of these customers beyond the ROI backlog. On demand revenue comprises only 5% of CoreWeave's revenue today, though we would note that given CoreWeave's exponential growth over the last two years, the company has not yet seen a large customer base expire. As we move through 2026 and 2027 and more contracts reach the end of their duration, should customer utilization go beyond the committed contract duration, greater revenue may be generated on a steady infrastructure base, leading to expansion of ROA.
- **Improved Infrastructure Efficiency:** CoreWeave is singularly focused on driving efficiencies in GPU-centric architectures (detailed in 'Agility of strategic differentiation to determine LT positioning' section of this report). Should the company drive ways to optimize CapEx investments and costs without compromising available capacity, this can reduce the dollar base of spend and improve ROA over time.
- **Software to be layered into model in outer years:** CoreWeave is investing in providing additional software-based solutions to compliment its infrastructure offerings. Currently, they have integrated software to improve the performance of their configured nodes and data center systems. Management has expressed intentions of moving up the stack to expand the scope of users engaging with CoreWeave and to further standardize and entrench organizations on the platform.

Current solutions are rolled out at no additional cost to drive the value add that CoreWeave provides. As CoreWeave establishes itself in the market, we expect future innovations to be monetized more directly. This can contribute to growth at limited incremental R&D. Regardless, these costs will sit outside of CapEx, and can support expansion of the revenue/net asset ratio. The company's recent Weights & Biases acquisition is the first step the company has taken to monetize software solutions.

**Exhibit 20: See return on average net assets expanding over time**

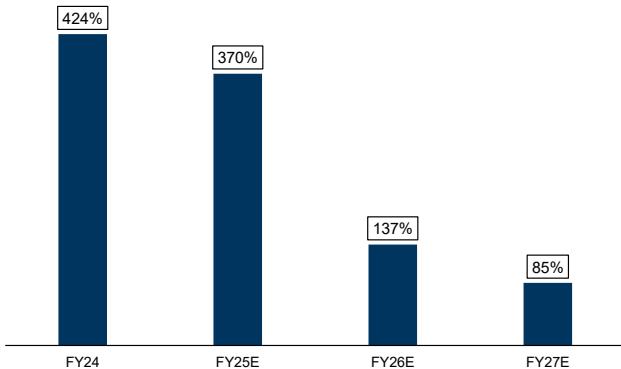
Revenue as a % of average net assets



Source: Company data, Goldman Sachs Global Investment Research

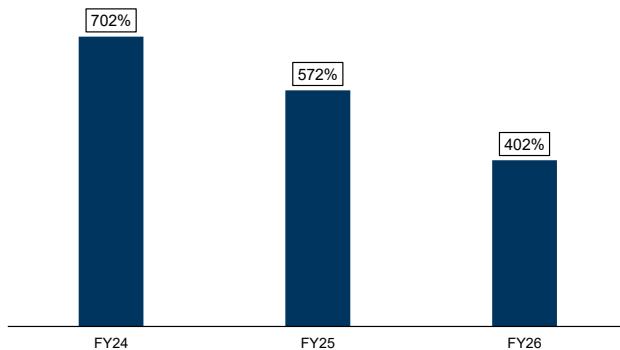
**Capital intensity related to AI infrastructure ahead of cloud peers; Expect further improvements.** The improvement in ROA will likely drive improvements in CoreWeave's capital intensity (CapEx/Revenue). CoreWeave has seen a lower CapEx intensity on its Gen-AI infrastructure vs. its hyperscaler peers ([Exhibit 21](#), [Exhibit 22](#)) given its CapEx spend is more directly tied to revenue. The spread between CapEx spend and CoreWeave's revenue recognition is typically 3-6 months. This has not been the case with other cloud providers who have needed to allocate a majority of their CapEx investments towards long-duration building blocks (i.e. land) in order to expand their Gen-AI-centric, GPU-based architectures to address future demand. As these projects start to come online, companies such as Microsoft have laid out the expectation to see a major mix-shift of their spend to short-duration, component-CapEx (Source: MSFT NDR). As GPU-related spend becomes a majority of spend over time, peers' capital intensity is likely to be more comparable to CoreWeave's. However, we expect CoreWeave to continue to outpace that of peers as they drive improvements via scaling infrastructure, increased efficiencies and expansion of software-related revenue. We lay out the progression of Microsoft's CapEx intensity since its Azure build out ([Exhibit 23](#)) as a reference to the long-term trajectory that may be evident through the evolution of Gen-AI.

**Exhibit 21: Expect CoreWeave's CapEx intensity to be reduced to <100% by FY27**  
Revenue divided by CapEx



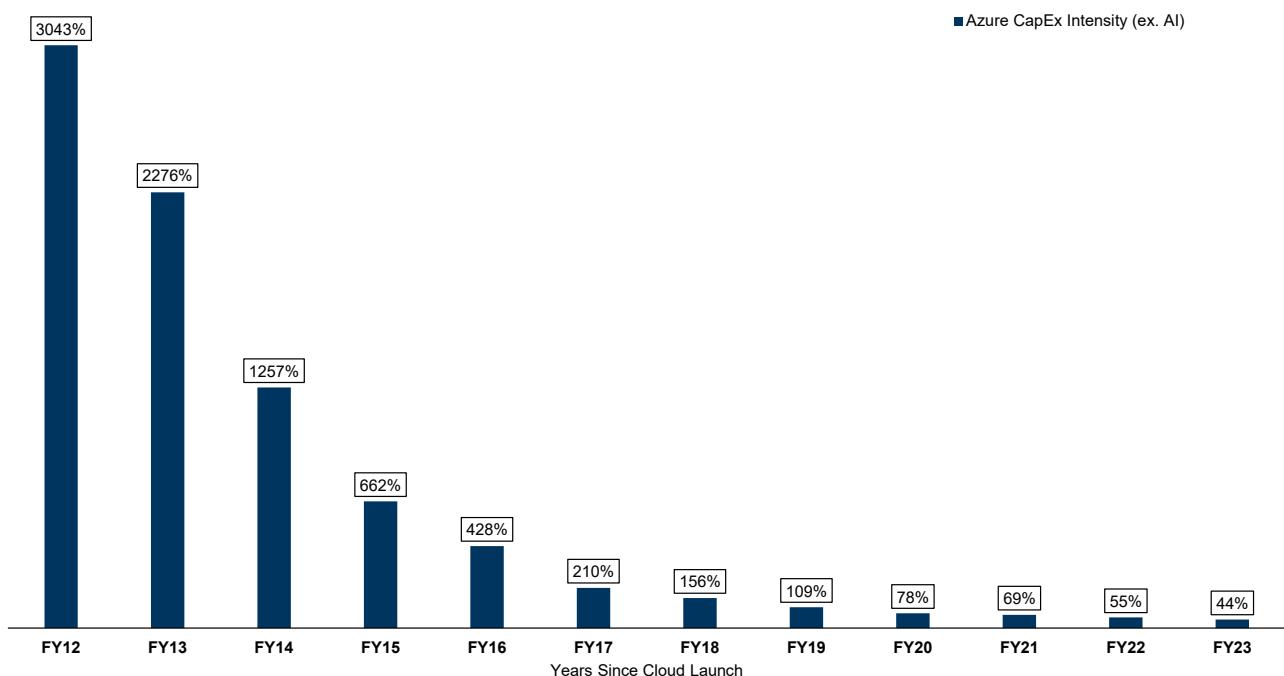
Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 22: See Azure AI CapEx intensity as a comparable benchmark for CoreWeave**  
Azure AI revenue divided by Azure AI CapEx



Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 23: Azure CapEx intensity fell dramatically as revenue scaled**  
Azure Revenue (pre-Resegementing) Divided by Total CapEx



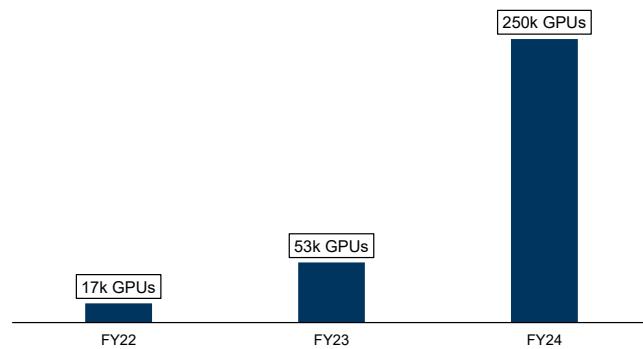
Source: Company data, Goldman Sachs Global Investment Research

## CapEx Buildout: Investment philosophy based on customer commitments

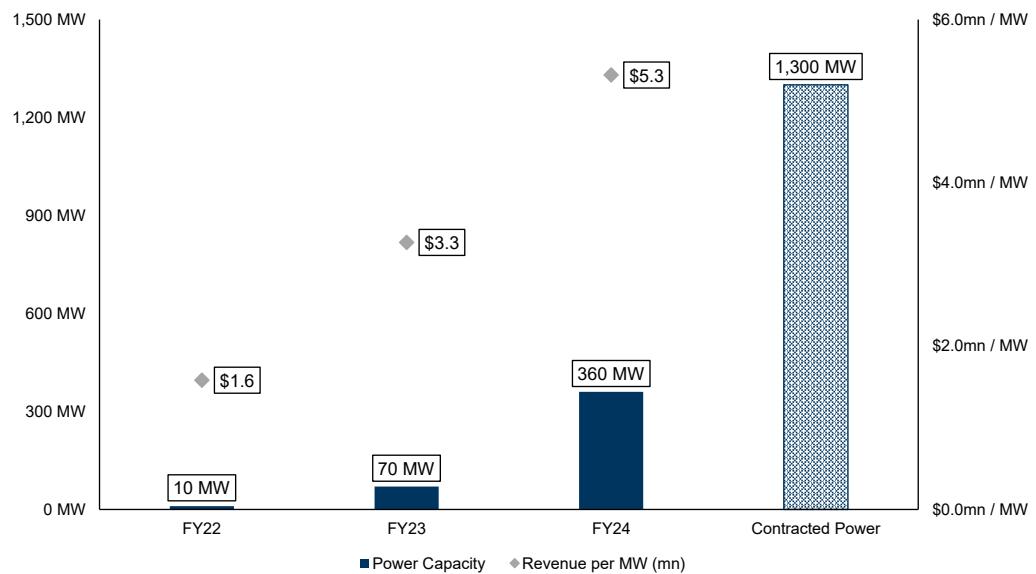
Total CapEx is an important metric for CoreWeave. The company's offerings are directly related to the reserved capacity it can deliver to its customers and CapEx is a direct barometer of that spend. For that reason, CapEx can serve as a leading indicator for revenue generation.

As of the end of 2024, CoreWeave operated more than 250k GPUs, comprising H100, H200, Hopper and others generations ([Exhibit 24](#)). As we look to 2025, we expect the mix to increasingly include Nvidia Blackwell chips, with more next-gen infrastructure as we look further ahead. The evolution of silicon will determine the mix of spend CoreWeave needs for resources such as data center spend, cooling, power and other configurations. For example, Nvidia's GB200's computing power is anywhere from six to thirty times greater than the H100, while reducing energy consumption by 25 times. While CoreWeave has 1.3GW of power contracted, the rate of utilization and durability of the manufactured infrastructure will be key to CoreWeave's ability to deliver its customer commitments ([Exhibit 25](#)).

**Exhibit 24: GPU capacity expanded over 4x from FY23 to FY24**



Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 25: CoreWeave has demonstrated an ability to rapidly scale power capacity**


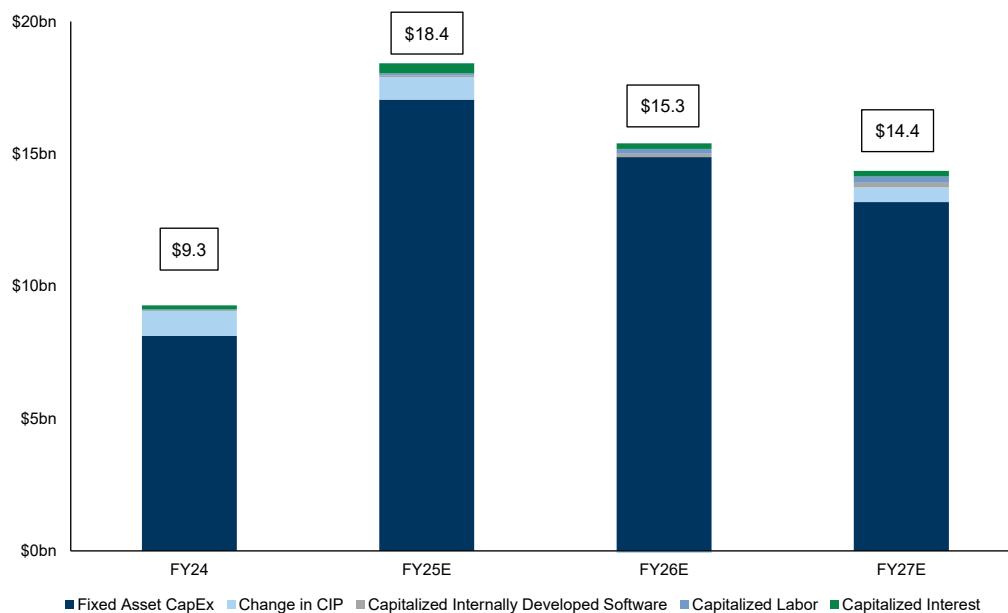
1,300 MW does not include additional contracted power of 70 MW from a single landlord and 133 MW from Galaxy

Source: Company data, Galaxy Digital Holdings Ltd., Goldman Sachs Global Investment Research

The biggest component of CapEx is CoreWeave's GPUs. We discuss the different buckets of CapEx (Fixed Asset vs Construction in Progress, or CIP) as CIP signals capacity (and therefore revenue) expected to come online in 1-2 quarters. We estimate Total CapEx of \$18.4bn in FY25, \$15.3bn in FY26, and \$14.4bn in FY27 ([Exhibit 26](#)).

- **Construction in Progress (CIP):** Construction in Progress is spent capital that has yet to be put in service (such as pre-revenue server equipment). As a preferred partner for Nvidia, CoreWeave procures their GPUs on a just-in-time basis. The period in which CapEx is procurement but not yet configured as a live deployment leads CapEx to be recognized in CIP. Further, Since this equipment is not yet generating revenue, we exclude it from the calculation of ROA (return on net assets) but see it as an important leading indicator to top-line. CIP can be the first initial signal of new deal signings.
- **Fixed Asset CapEx:** Fixed Asset CapEx represents all of CoreWeave's spending on computing equipment (GPUs, CPUs, etc.) and data center equipment, such as procurement and lease needs. Smaller expenses on software, furniture/other assets, and leasehold improvements are also included but are minimal. Fixed asset expense is largely a function of signed contracts for which CoreWeave is building out capacity. Quarterly variability in Fixed Asset spending is driven by timing of data center capacity coming online and new generations of GPUs being placed into service. Currently CapEx is moderating over time in absolute dollar terms by function of limited visibility into additional capacity needs further out (given typical 1-2 quarter lead time once contract is signed). If CoreWeave signs more contracts than currently modeled, investors would see this upside as increased Fixed Asset spending.

**Exhibit 26: Expect CapEx to nearly double in FY25 and to moderate in the outer years**  
**CapEx Components**

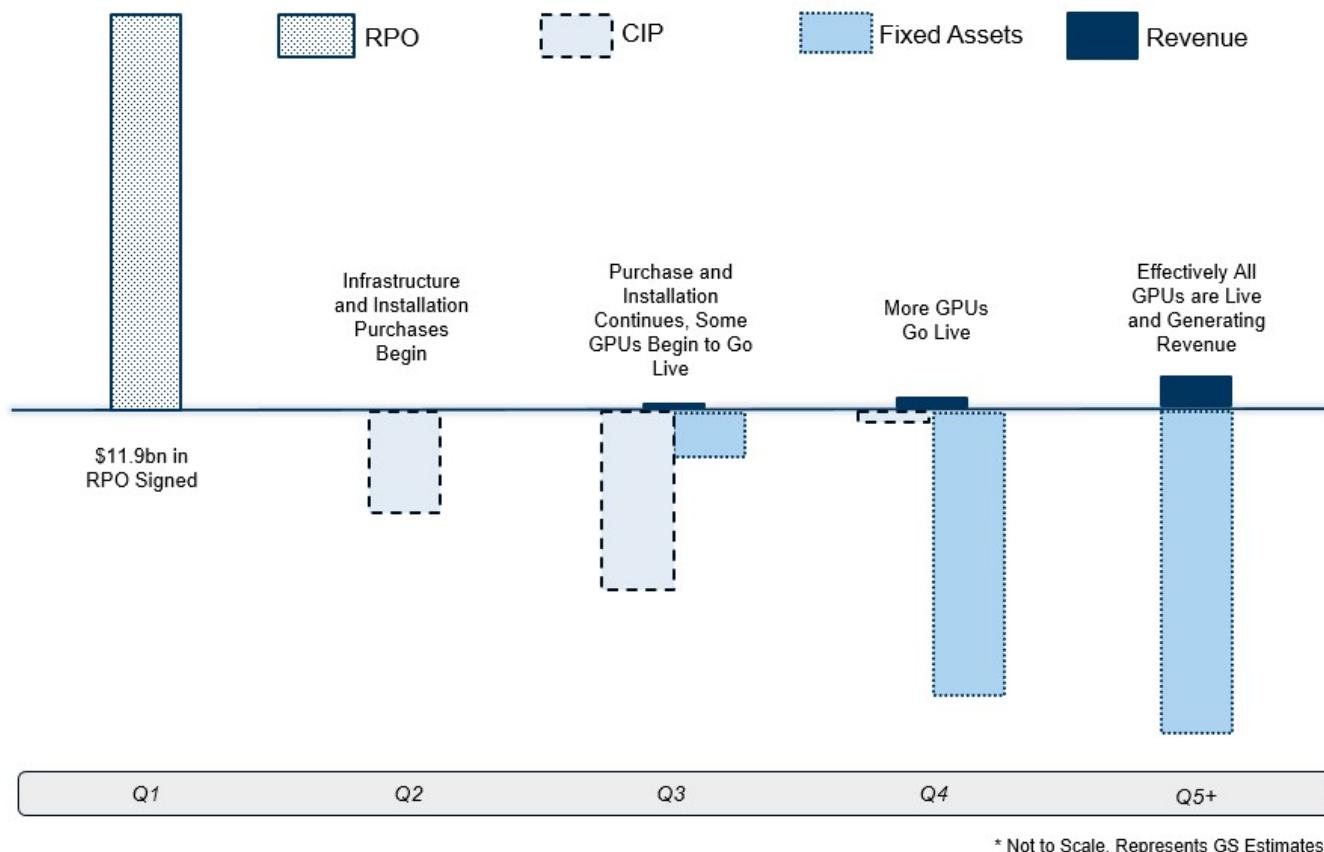


Source: Company data, Goldman Sachs Global Investment Research

**Case Study: Flow through of a contract.** We use CoreWeave's March 2025 deal with OpenAI as a case study of how a contract can flow through CoreWeave's model once signed ([Exhibit 27](#)). The scale of the +\$11bn+ OpenAI deal suggests CoreWeave is committed to supplying OpenAI with ~110K Nvidia GB200 Blackwell GPUs. Assuming an average cost of \$65K per GB200 chip, we presume the associated cost of GPU procurement for this project is ~\$7.15bn. Given the size of the contract, it is likely to take 6 months and several different deployment milestones within that timeframe. As CoreWeave executes against each delivery date, the associated CapEx spend will move from CIP to Fixed Asset CapEx. The ratable revenue associated will also start being recognized. For this reason, we model an increase in CIP of \$5.7bn between 2Q-3Q (with the remainder likely in 4Q). We also expect Fixed Asset CapEx increasing by \$7.7bn in 4Q. While note all of this will be attributed to OpenAI- with CoreWeave having prior commitments to deploy Blackwell chips in 2H25, the magnitude of the deal leads us to believe OpenAI will drive the majority of this spending.

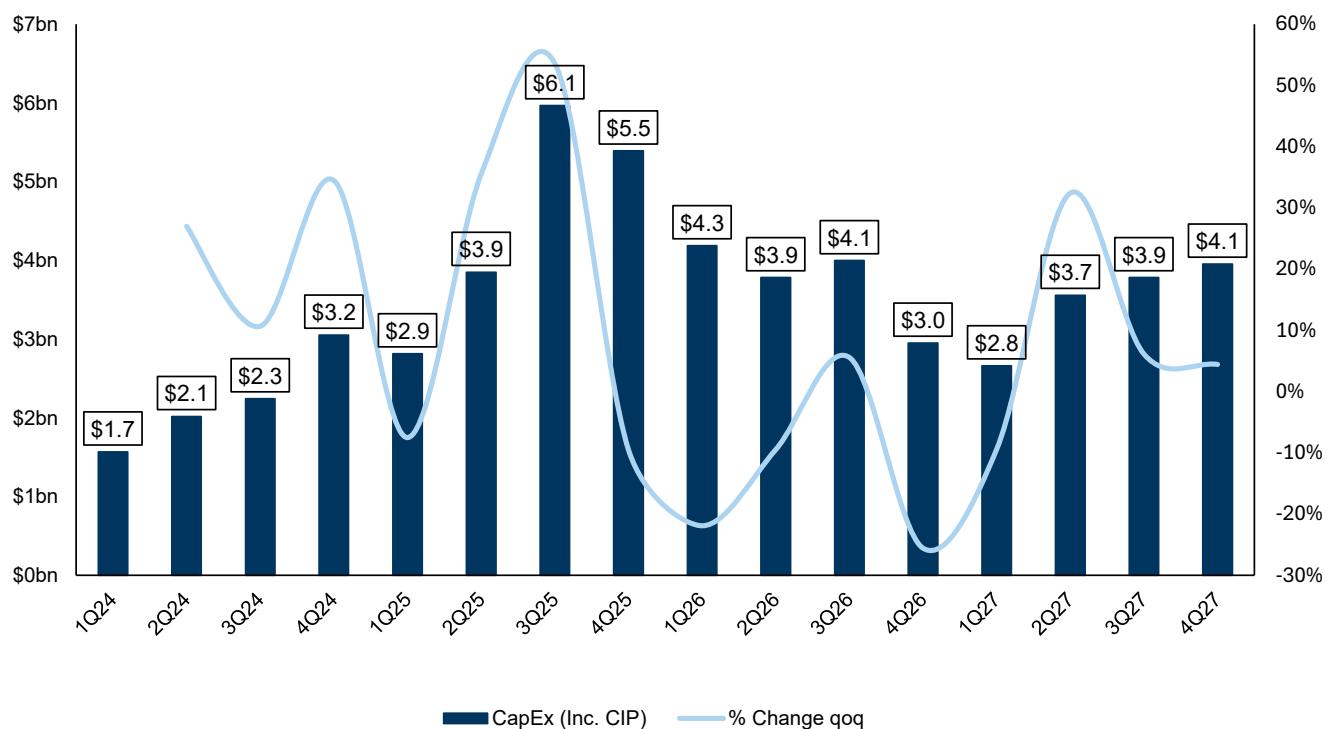
We note that RPO would be adjusted higher once the deal is signed. Should OpenAI have paid an upfront payment (often part of a CoreWeave contract) we would have likely seen 15-25% of the TCV show up in deferred revenue. We note that the front-end of a contract is typically the most cost intensive, with hardware investments occurring before any revenue is recognized. The tail-end of the contract, is much more cash flow accretive given CoreWeave sees an avg payback period of 2.5 years (vs avg contract duration of 4 years).

## Exhibit 27: LifeCycle of CoreWeave's OpenAI contract



Source: Company data, Goldman Sachs Global Investment Research

**CapEx cadence dependent on timing of contract wins.** Given the cadence of CapEx is dependent on the timing of deal signings, we expect our forward looking CapEx expectations to vary from what the more baseline and largely stable level of spend we model in 2026/2027 today. The timing of the OpenAI contract and the fact that this is already factored into our model shows the implications that any deal can have on the spending trajectory of the business. As such, 2025 sees a much higher spike to CapEx vs our estimates for 2026/2027, that are more in line with pre-OpenAI deal baseline ([Exhibit 28](#)).

**Exhibit 28: Expect CapEx spend to fluctuate with the OpenAI buildout, moderate in the outer-years**


Source: Company data, Goldman Sachs Global Investment Research

**CoreWeave's differentiated CapEx philosophy vs. cloud peers.** A key debate surrounding cloud providers pertains to the expected ROI of their CapEx build out. We expect this narrative to be somewhat different for CoreWeave driven by the majority of the company's deployed capital being allotted to component CapEx on a 'just-in-time' basis. The GPUs associated with a given contract are therefore ordered following contract signing, with the procurement and deployment of such infrastructure leveraging the prepayment from the contract and taking place over a 3-6 month period. The contract go-live often occurs in 1-2 quarters post signing, at which point CoreWeave starts recognizing ratable revenue on these contracts. This is different from peers who build a customer-agnostic capacity pool to get ahead of future demand. CoreWeave is able to operate as such given its preferred partnership agreement with Nvidia and other suppliers.

This ultimately means that CoreWeave has less "at risk" capital versus that of hyperscalers, where investors hold concerns over utilization rates and potential oversupply. The areas where CoreWeave does need to deploy 'at-risk' capital largely pertain to procurement of power and data center needs (more in the next section). Similar to broader cloud providers, however, management looks at market conditions to inform their expectations around these investments.

It is worth noting that CoreWeave does not currently build or own its own data centers today. While this may change, should management choose to build its own data centers, this would be a pivot from its current operational model. While doing so would reduce operating costs it would increase control over delivery timelines (a key pillar of

CoreWeave's value prop to customers). We do not expect CoreWeave to change its philosophy around renting until the company generates durable FCF.

## CapEx Buildout: Depreciation dynamics key to model assumptions

**Replacement costs = necessary cost of CoreWeave's business.** The continued innovation in silicone and the ongoing introduction next-gen chips will be an essential dynamic to keep in mind for CoreWeave. The company has a rapidly growing fleet of GPUs, with both costs and revenue directly correlated to the underlying infrastructure. For this reason, replacement costs of GPUs is a necessary cost of CoreWeave's business.

This leads us to conduct the below depreciation analysis as we look to discern the potential financial implications should future silicone generations have shorter economic or useful life than the six years currently expected. As noted earlier, CapEx spend primarily consists of investments in technology equipment (such as GPUs) and Data Center equipment (to build out or maintain data centers). Computing equipment is depreciated over 6 years whereas DC buildout has a useful life between 8-12 years.

**OpenAI contract as case study for depreciation assumptions' impact on operating margin.** With investors mostly apprehensive around the right useful life of next-gen GPUs, such as Blackwell, we utilize the company's latest OpenAI contract as a case study to determine how the useful life of the underlying GPUs associated with this contract may have on operating margin. For background, CoreWeave signed a 5-year contract with OpenAI for ~110K Nvidia GB200 Blackwell chips. OpenAI also has an option to extend the contract by two one-year extensions for a potential duration of seven years. If we apply the same cost assumptions as the example in previous sections, assuming an average cost of \$65K/GPU, we assume the associated cost to procure these GPUs will be ~\$7.15bn.

In the below analysis, we look to review the implications to the company's non-GAAP operating margin across a number of useful life scenarios ([Exhibit 29](#)).

- We use six years as a base case as this is the company's current depreciation period per GPU. This implies annual depreciation on this contract will be \$1.2bn - based on ~\$7bn of upfront GPU spend. With the contract not likely to be fully live until 4Q25, we look at the impact this will have to our 2026 operating margin estimates of 24%. While this already accounts for the incremental \$1.2bn in depreciation, we look to determine the compression of this margin should the realized useful life be 5 or 4 years instead.
- Adjusting our estimates for this compressed timeline, we see D&A increase by 50% and 20%, respectively. This would result in operating margin contracting 220bps and 550bps, respectively.
- On the other hand, OpenAI be able to exercise its two one-year extensions and leverage this underlying infrastructure for seven years, the inverse effect could take place, with depreciation decreasing 14% per year and 2026 operating margin increasing ~160bps.

**Exhibit 29: Sensitivity analysis suggests a -500bps impact to FY26 OpM if useful life is lowered to 4 years, +150bps if raise to 7 years**

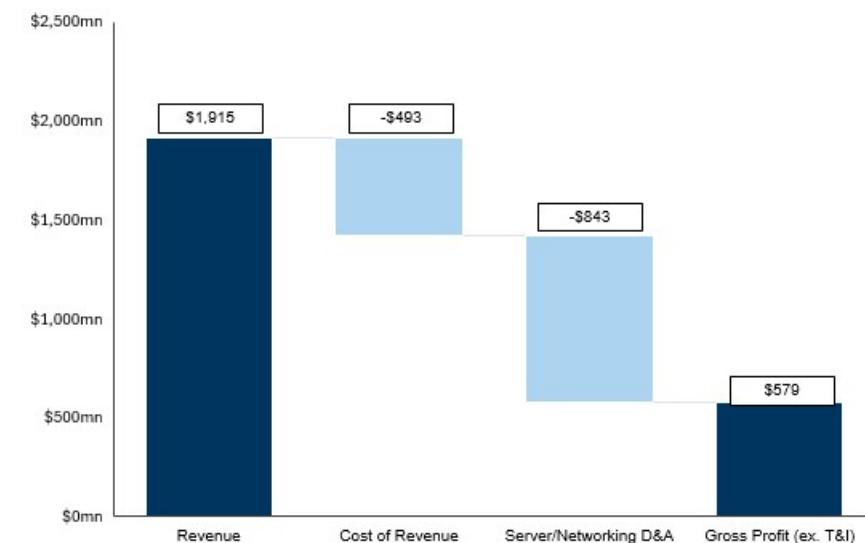
OpenAI		Adjustments to useful life accounting impact to OM			
		7 years	6 years	5 years	4 years
# of GPUs	\$110k	Depreciation (\$mn)	\$1,021	\$1,192	\$1,430
Cost per Blackwell	\$65k	<i>Delta vs base case</i>	-14%	+20%	+50%
Blackwell Investment	<b>\$7,150bn</b>	Operating Margin (2026)	25%	24%	21%
		<i>Delta vs base case</i>	157 bps	-220 bps	-550 bps

Source: Company data, Goldman Sachs Global Investment Research

**Adjustment to reported Gross Margin needed to encompass holistic cost of serving contracts.**

The above scenario highlights the magnitude of depreciation as a function of CoreWeave's business model. We take this into consideration as we evaluate the cost of servicing a CoreWeave contract. Since the company's GPU base (and therefore the D&A recognized) is a core part of the service provided, we believe investors should include D&A when evaluating their gross margins ([Exhibit 30](#)).

CoreWeave's reporting structure excludes D&A in COGS but rather buckets it into its Technology and Infrastructure (T&I) operating line. CoreWeave's T&I consists primarily of depreciation and amortization related to infrastructure costs (servers, switches, networking equipment, and internally developed software), with a small portion allotted to Research and Development (R&D) personnel costs. However, we generally consider servers a direct cost of selling their product (as opposed to a more indirect operating expense), particularly as CoreWeave does not purchase their infrastructure until a contract has been signed.

**Exhibit 30: FY24 GAAP gross profit calculated including Server and Networking D&A**

Source: Company data, Goldman Sachs Global Investment Research

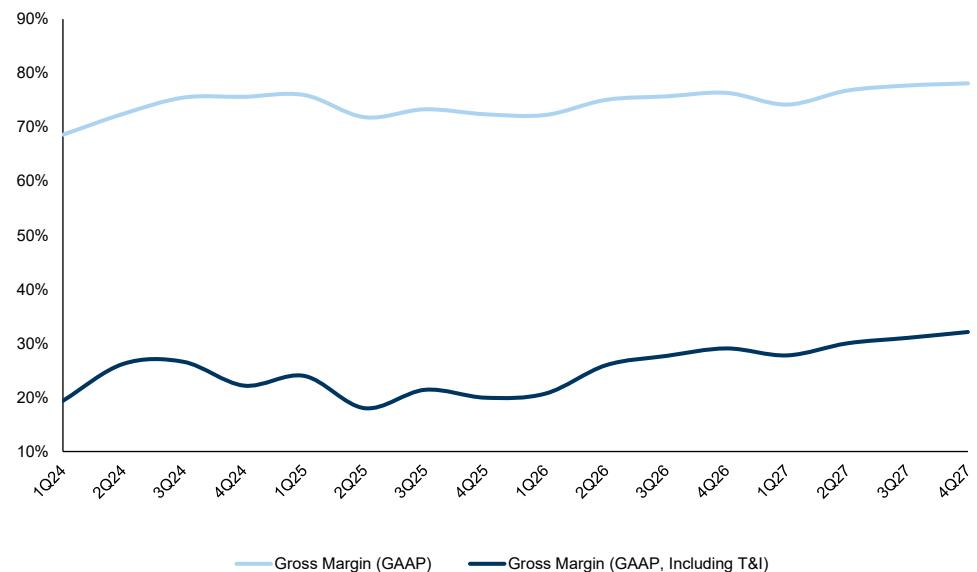
**Incorporating Depreciation and Amortization into Gross Margin evaluation.** Since CoreWeave has such a capital intensive business model and depreciation expenses are material, we choose to calculate Gross Margin in two ways (both including and excluding T&I) to have added visibility into the unit economics of the business. For reference, Research and Development costs in FY24 and FY23 increased \$56mn and \$17mn respectively compared to an increase in D&A costs of \$742mn and \$90mn

respectively. Since CoreWeave will need to go out and repurchase GPUs over time to sustain the growth of the business, incorporation of this as a real cost gives way to the magnitude and timeline of such reinvestment.

We break out the different components in this calculation in [Exhibit 30](#) and represent the par between CoreWeave's GM including and excluding these T&I costs in [Exhibit 31](#).

**Exhibit 31: Gross Margin including and excluding T&I expected to expand 600bps and 300bps respectively by FY27**

Gross Margin Estimates Including and Excluding T&I Expense



T&I is inclusive of R&D spend

Source: Company data, Goldman Sachs Global Investment Research

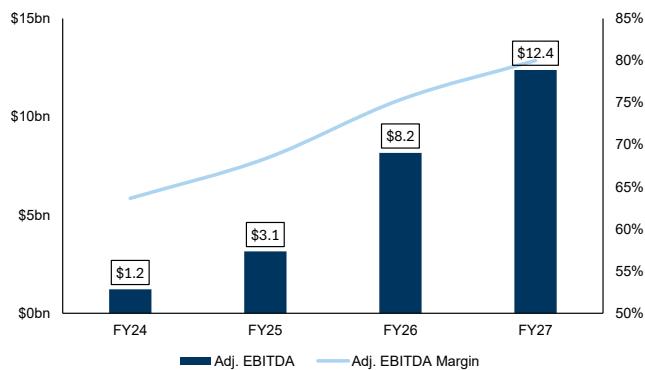
**Business needs likely to drive high rate of cash burn, push out FCF profitability.**

Given the asset-heavy nature of CoreWeave's business, the company's investment in fixed assets has a direct correlation with the future revenue growth of the business. For this reason, our estimates do not expect material improvement in the business' free cash flow through FY26. While adjusted EBITDA progress is likely given the expected improvement in return on net assets ([Exhibit 32](#)), our estimates assume FCF remains in the negative \$3-16bn range through FY27 ([Exhibit 33](#)). This is a function of both debt costs and reinvestment needs to replace aging infrastructure. As long as business prospects continue to require a high degree of CapEx investments, the achievability of FCF profitability is likely to be pushed out.

- **Investors will need to be comfortable with the potential for the business' current cash** burn rate to persist until clarity is provided around either CoreWeave's ability to reduce its dependence on net asset growth to drive revenue, or management reduces its reinvestment ratio.
- **CoreWeave's nascent software solutions, which are currently not being monetized, are a long-term lever that may offset these dynamics and drive cash generation.** Software can be scaled without incremental CapEx investments,

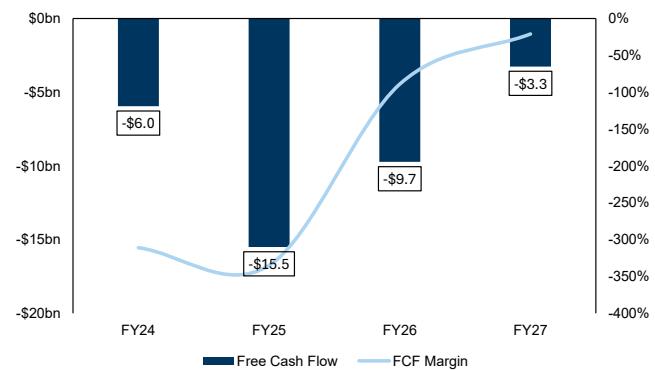
making it a growth avenue that can improve unit economics and drive FCF improvement should management flow the high-margin revenue through to FCF.

**Exhibit 32: Expect expanding EBITDA margin as return on net assets improves, revenue scales**



Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 33: Expect FCF margin to expand from -311% to -21% in FY27 as OCF scales, CapEx moderates**



Source: Company data, Goldman Sachs Global Investment Research



## CapEx Buildout: Strategic navigation of supply constraints

### **Progressing on path to supply/demand equilibrium**

**Supply constraints of key data center components was the main limiting factor of industry growth.** The rapid pace of innovation and desired adoption of Gen-AI have outpaced the constrained production capabilities of building out GPUs, bringing power online and building out data centers (the main components of AI-based data centers). This was the industry's biggest challenge to meeting customer needs and driving additional growth. CoreWeave was able to benefit as a relatively new player as it had access to these needs. Since CoreWeave acquires the needed GPUs on a just-in-time basis based on customer contracts, the company's close supply chain relationships have helped the company insulate itself from the full extent of these constraints.

- **Strategic supply chain partnerships.** With 1.3GW of power contracted and operating across 32 data centers with over 250K GPUs, CoreWeave currently has enough capacity to meet existing commitments (with an additional 70MW from a single landlord and 133MW from Galaxy, per Galaxy's 4Q earnings release). However, given three suppliers accounted for 46%, 16% and 14% of purchases in FY24, any strains of any these relationships could have a material effect on CoreWeave's operations. Should bottlenecks or competition increase, there might be added risk to revenue conversion of CoreWeave's backlog and/or its ability to drive future growth.
  - **Symbiotic relationship with Nvidia.** CoreWeave has a Symbiotic relationship with Nvidia, where both companies are customers of the other (CoreWeave exclusively purchases Nvidia GPUs and Nvidia leverages CoreWeave's capacity). The strategic rationale behind this, which allows Nvidia to have a large, at scale environment to test new GPU capabilities, has also given way to CoreWeave becoming Nvidia's first Elite Cloud Services Provider for compute within the Nvidia partner network. These inter-dependencies that gives Nvidia diversification from hyperscalers (who are also designing their own silicone) while differentiating CoreWeave given its priority allocation of in-demand GPUs. Nvidia has also invested in CoreWeave, owning ~5% of the company pre-IPO, with additional investment via its participation in the IPO offering.

**Improving visibility around alleviation of some constraints; CoreWeave likely to lean into key differentiators, efficiencies to drive share gains.** Our teams (across Data Center, Utilities, Sustainability), have started to forecast an alleviation of constraints in data center capacity though forecast strong demand for power through 2030. GS' Data Center model (run by Jim Schneider) expects loosening constraints may occur earlier than expected, with peak datacenter occupancy likely to occur in 2025 (vs expectations 2H26 prior). Given time to market has been a key differentiator for CoreWeave, a more balanced market will likely accelerate the need for CoreWeave to lean in and solidify the other components of its value proposition (GPU-focus, improved reliability, dedicated infrastructure, software, etc).

Meanwhile, our Utilities and Sustainability teams (led by Carly Davenport and Brian Singer, respectively) expect a ~160% increase in datacenter power demand by 2030 vs. 2023 levels to be met with moderating power intensity. CoreWeave will need to demonstrate continued ability to source power and drive internal efficiencies to meet industry demand.

### **Data Centers: Current state of supply/demand and path to steady state:**

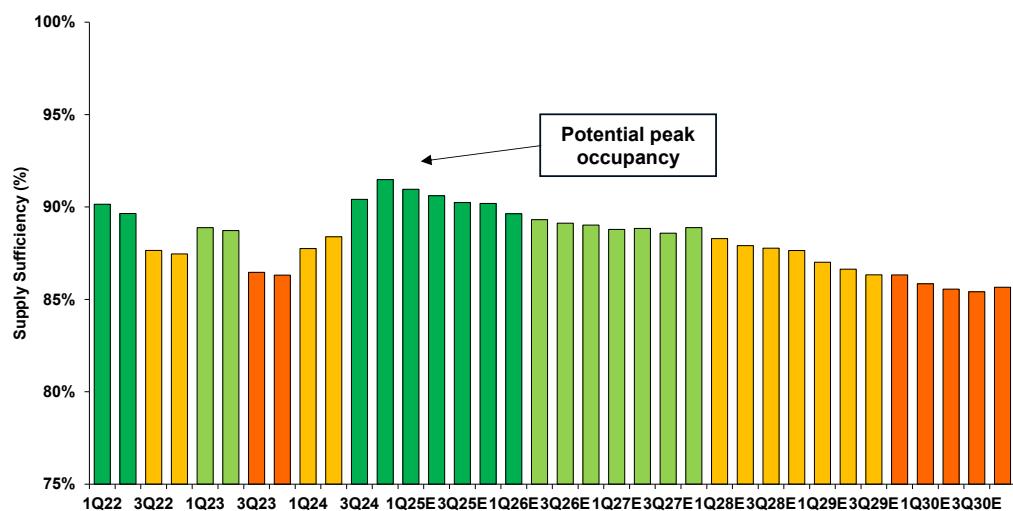
We collaborate with our Data Center team (led by Jim Schneider) to flesh out where the market capacity stands today in terms of data center capacity and how CoreWeave is positioned against that.

**Our Data Center model expects loosening constraints may occur earlier than expected, with peak datacenter occupancy likely to occur in 2025 (vs expectations 2H26 prior).** Since we introduced our first datacenter supply/demand model, the datacenter market has contended with the impact of multiple AI-related announcements, including the impact of DeepSeek on AI training (and datacenter power demand) and multiple supply announcements including Stargate's datacenter facilities. Our global technology team also cut estimates for AI training servers based on a combination of product transition issues and demand uncertainty.

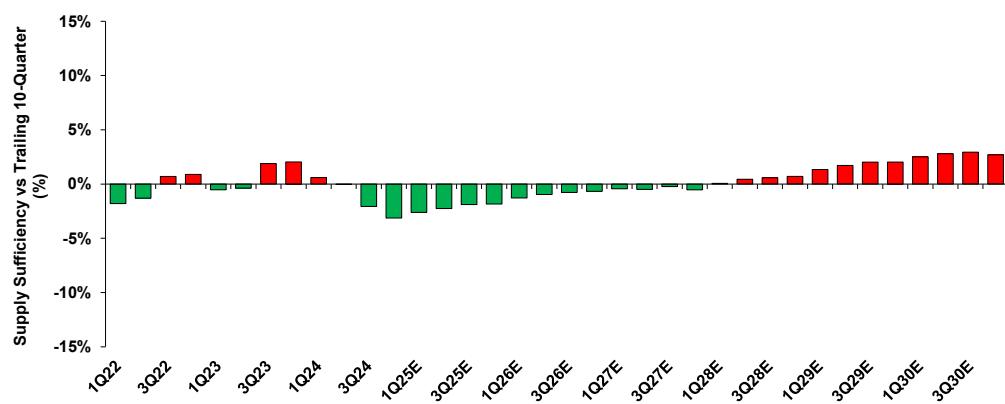
Today, our teams see relatively little change in terms of the health of short-term supply/demand balance - occupancy rates remain near record highs for third-party leased datacenter facilities across most US markets, and our own checks suggest that market lease pricing continues to move higher - albeit at a slowing rate of increase than in past years.

Looking out, however, our data center team (covered by Jim Schneider) now expects the peak of datacenter supply sufficiency is likely to be pulled forward into 2025 from 2H26 prior. We believe the datacenter market's current supply/demand tightness will begin to loosen sooner than we previously expected, although the teams' model suggests that market occupancy will stabilize around average levels seen over the past 18 months. In summary, the outlook for datacenter supply, demand, and their implied supply sufficiency remains relatively healthy for now (Exhibit 34). Jim Schneider watches for incremental datapoints that could cause a shift in expectations or changes in GPU demand, AI model efficiencies as well as announced supply additions (such as Stargate) that could significantly impact medium-term supply/demand balance.

- Although it is difficult to have visibility beyond 2027 at this stage, Jim Schneider expects supply constraints easing beyond this point, with occupancy rates gradually falling back to ~88% by 2028 and leveling out thereafter. For context, at its projected peak in 2025, occupancy rates would still exceed record levels achieved over the past 15 years. Therefore, even as supply constraints ease, he models occupancy rates remaining above the trailing 10-quarter rate for the next 18-24 months (Exhibit 34, Exhibit 35). At an aggregate level, this implies that global supply/demand is likely to remain favorable in the medium term, with peak markets likely to be more heavily constrained (with less utilized secondary and tertiary markets) as we currently see in the market.

**Exhibit 34: GS supply sufficiency estimate**

Source: 415 Research - part of S&amp;P Global Market Intelligence, Goldman Sachs Global Investment Research

**Exhibit 35: GS supply sufficiency estimate vs trailing 10-quarter supply sufficiency**

Source: 415 Research - part of S&amp;P Global Market Intelligence, Goldman Sachs Global Investment Research

**CoreWeave likely to need to lean into Gen-AI expertise, other differentiators, as capacity comes online.** Given time to market has been a key differentiator for CoreWeave, a more balanced market will likely accelerate the need for it to lean in and solidify the other components of its value proposition (GPU-focus, improved reliability, dedicated infrastructure, software, etc). With more supply coming online in 2H25 from hyperscalers' prior investments over the last two years, CoreWeave's ability to sustain its speed to market and strong reputation for on-time delivery will be imperative to increase customer confidence around choosing CoreWeave. More on CoreWeave's differentiation in the first section of this report.

Jim Schneider also highlights the following scenarios that can lead to variability vs base case. We opine on the implications each may have around CoreWeave's positioning,

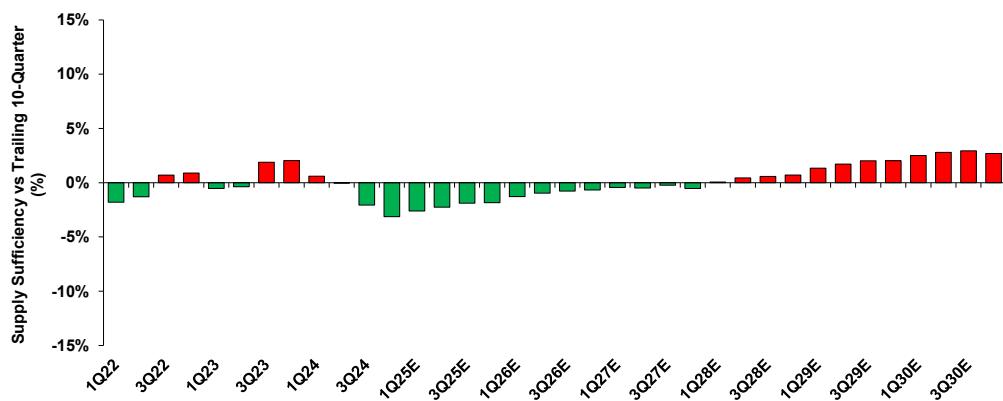
investments and strategy.

- **"AI upside" demand scenario:** This scenario assumes that the next generation of GPUs requires more power than the industry currently expects, and/or that demand for AI workloads accelerates faster than we currently estimate. We increase the AI portion of our demand model by 30%-35% in the period from 2026-30 in this scenario, resulting in a total demand 3-year CAGR from 2024 to 2027 of 19% versus our base case of 17%. In this scenario, we believe occupancy rates could exceed 95% in peak regions by 2030, and at an overall level this implies 13 points higher occupancy than in our base case scenario ([Exhibit 36](#)).

- **CoreWeave positioning to be strengthened in such a scenario.** Given CoreWeave's GPU and AI expertise, should demand for AI workload accelerate, it can elevate CoreWeave's differentiation. Customers often note faster deployments and higher uptimes, which we expect to be key considerations for customer buying behavior. The cost visibility CoreWeave provides (with committed contracts) can be another factor, especially after many companies needed to optimize spend with other cloud providers just a few years ago.

---

**Exhibit 36: "AI upside" growth scenario implied supply sufficiency**



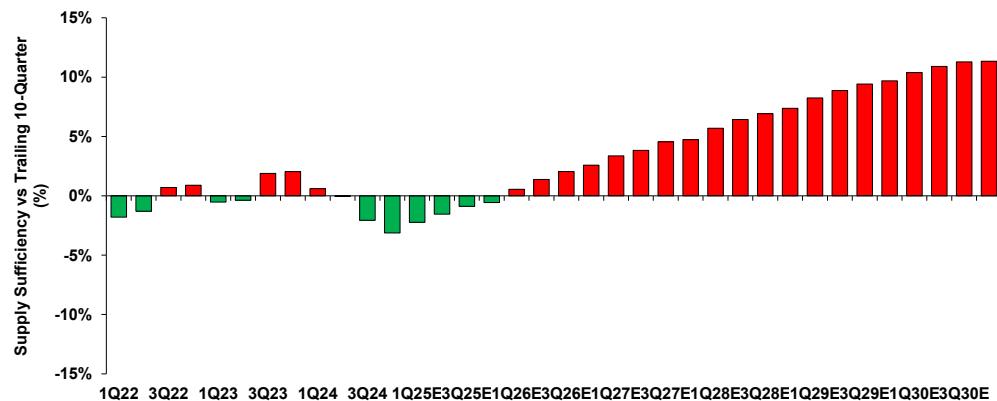
Source: 451 Research - part of S&P Global Market Intelligence, Goldman Sachs Global Investment Research

- **"Excess supply" scenario:** We also model a scenario where datacenter companies are able to bring on more incremental capacity than we expect in our base case, while assuming no change in datacenter demand. In this scenario, we assume the global datacenter market will bring on an additional 15 GW of capacity by 2030 versus our base case, which results in a difference in supply sufficiency downside of approximately 9 points by 2030 relative to our base case ([Exhibit 37](#)).

- CoreWeave does not currently build its own data centers but rather has strong relationships with Data center providers to allocate capacity as needed. Therefore, more available capacity can help with pipeline planning, although it will likely reduce time to market as a key differentiator. This can yield a similar strategic shift as the base, case, where we outlined our view that CoreWeave will need to lean into Gen-AI expertise, other differentiators, if capacity

availability eases.

#### Exhibit 37: Excess supply scenario implied supply sufficiency

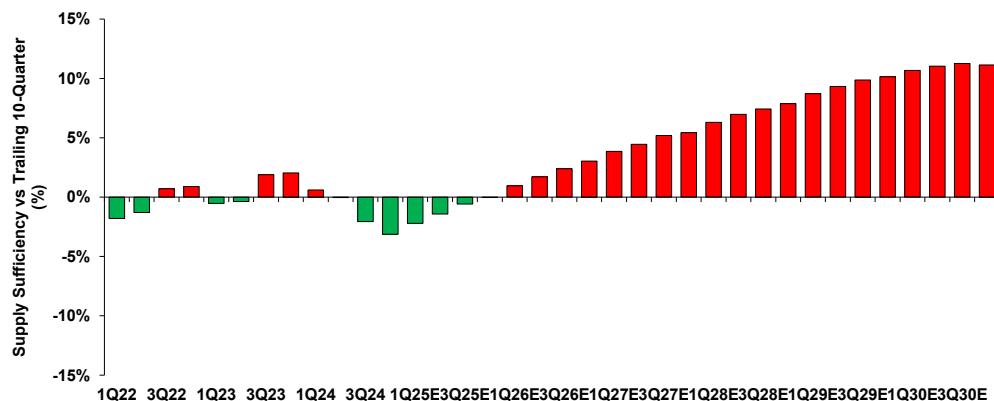


Source: 451 Research - part of S&P Global Market Intelligence, Goldman Sachs Global Investment Research

- **"AI downside" demand scenario:** In this scenario, Jim Schneider sensitivity tests a meaningfully lower growth rate for AI-focused workloads in the event that monetization of AI proves to be slower than we currently expect. Our AI forecast is already below our prior estimates due to the AI training server cuts noted previously. Depressing AI demand growth by an additional 20% in 2025-30 results in a total demand 3-year CAGR of 14% in the period 2025-27 relative to our base case growth CAGR of 17%. In this scenario, the difference in occupancy between GS' base case and "AI downside" case is approximately eight percentage points in 2030, which translates to a forecast difference of 11 GW in power demand terms ([Exhibit 38](#)).

- **Represents downside case for CoreWeave given moderating need for customized, GPU-centric architectures.** This scenario would present a downside case for CoreWeave as AI/GPU-based workloads are deprioritized, leading to moderating growth. Given CoreWeave's value proposition is reliant on the continued expansion of the AI market to justify the specialized configurations it holds. Such a dynamic would be a meaningful change vs what we are seeing in the market today (AI demand growing rapidly and meaningfully exceeding demand) and is therefore not something we expect to ultimately occur.

**Exhibit 38: "AI downside" scenario implied supply sufficiency**



Source: 451 Research - part of S&P Global Market Intelligence, Goldman Sachs Global Investment Research

**Power: Current state of supply/demand and path to steady state:**

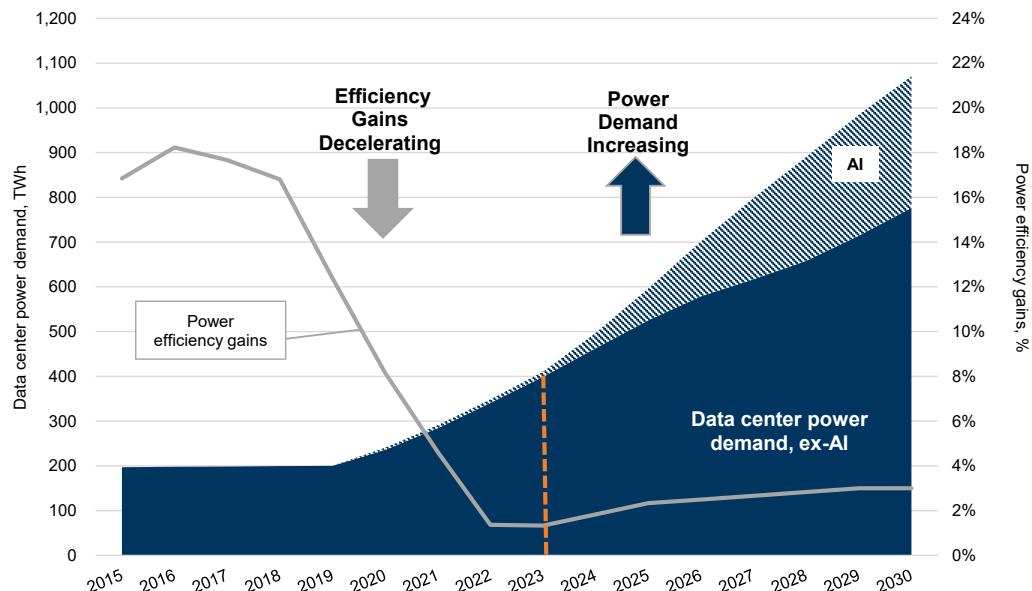
We collaborate with our Utility and Sustainability teams (led by Carly Davenport and Brian Singer, respectively) to flesh out where the market capacity stands today and how CoreWeave is positioned against that. **Overall, we see CoreWeave needing to demonstrate continued ability to source power and drive internal efficiencies to meet industry demand** ([Exhibit 39](#)).

**Elevated demand still driving shortage of supply...** The team continues to expect data demand — driven in part by AI and in part by traditional data center workloads – to catalyze generational growth in global power demand. With AI expected to make up roughly 30% of 2030 data center demand, our utilities team continues to expect a ~160% increase in global datacenter power demand by 2030 vs. 2023 levels.

- Overall data centers now contribute about ~100 bps of growth to our 2.5% US power demand CAGR through 2030, with about 60 bps coming from non-AI data center workloads. As such, if the team would adjust AI growth out of their model, they would still expect to see US power demand CAGR of about 2% through 2030, well above growth over the last two decades.
- Base case implies datacenter power demand moves from 1%-2% of overall global power demand to 3%-4% by 2030. In the US, the pace of mix increase is even greater, more than doubling by 2030 from 4% in 2023. If global datacenter growth in 2030 vs. 2023 levels were its own country, it would be a top 10 global power consumer.

**Exhibit 39: As efficiency gains have decelerated, data demand growth is driving a surge in datacenter power use, with AI acceleration expected to continue**

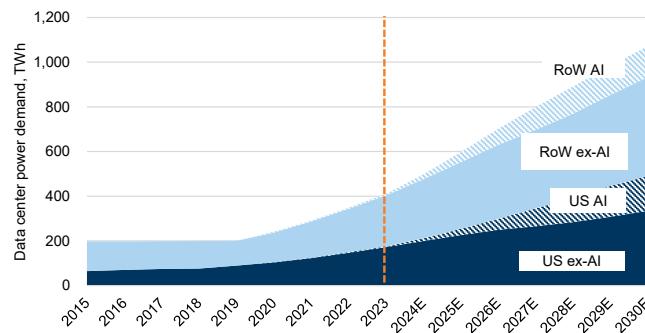
Datacenter electricity consumption, TWh (LHS) and power efficiency gains ex-AI, % (RHS)



Source: Masanet et al. (2020), IEA, Cisco, Goldman Sachs Global Investment Research

**Exhibit 40: After being flat for 2015-19, we have seen data center power demand accelerate in 2021-23 and expect a ~160% increase through the rest of the decade**

Global datacenter electricity consumption, TWh; includes AI and excludes cryptocurrency



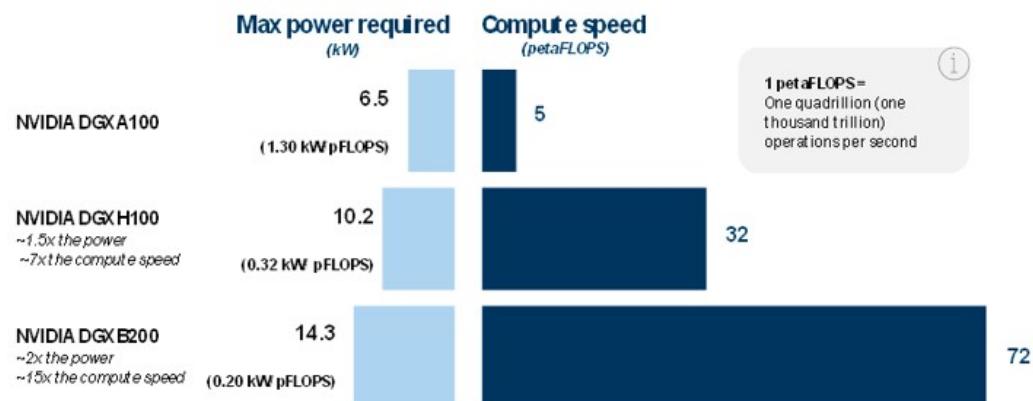
Source: Masnet et al. (2020), Cisco, IEA, Goldman Sachs Global Investment Research

**...Though expect power intensity to moderate.** Our teams note that the market has seen new AI innovations increase max power consumption per server but the increase in computing speed per server grew by an even greater level. This represents a meaningful reduction in power intensity. Similarly, DeepSeek raised the questions as to the pace of innovation of how much compute speed will be needed based on efficiency of AI solutions per unit of compute speed. Recent evolutions of Nvidia server system specifications are indicative of increasing max power per server but with lower power intensity relative to computing speed (for training) - as seen in [Exhibit 41](#).

- Given CoreWeave is bringing on many new clusters or next-gen Nvidia chips, the puts and takes of these dynamics is critical for CoreWeave to ensure it has a sufficient amount of power coming inline to support future growth.
- **CoreWeave's partnerships critical to meeting future demand.** This leads us to see CoreWeave's partnerships, such as that with Core Scientific, as a key pillar of the company's business model. If CoreWeave is unable to procure power or have rights to power coming online in the future, this can limit the workload capacity it can support and therefore the net new business it can generate.

**Exhibit 41: We have seen new AI innovations increase max power consumption per server but increase computing speed per server by an even greater level, representing a meaningful reduction in power intensity**

Recent evolution of NVIDIA server system specifications is indicative of increasing max power per server but with lower power intensity relative to computing speed (for training)



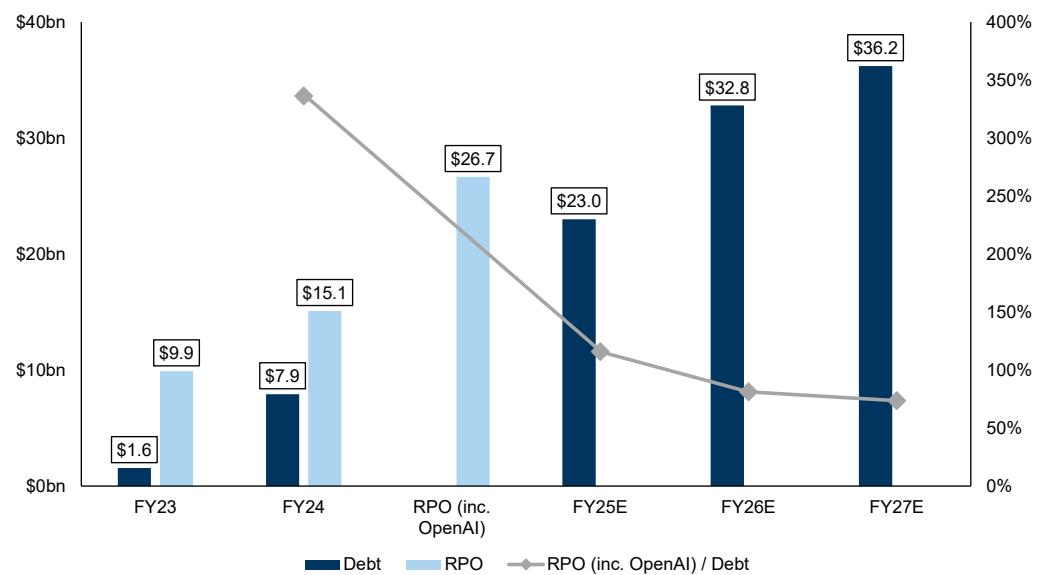
Source: Nvidia, Goldman Sachs Global Investment Research

## Debt Burden: Revenue-backed financing lowers risks associated with scaling debt needs

We see CoreWeave's current valuation reflecting an outsized focus on the magnitude of the company's debt load while under-appreciating its favorable contract dynamics. While debt will be the primary financing instrument to sustain CoreWeave's capital-intensive business, we index toward the terms of such capital and ultimately see the debt as manageable. Capital needs are largely an output of contracted deals and are amortized against the cash flows of customer contracts, lowering the risk of competing cash flow needs. As a function of CapEx spending – which is taken on only once a deal is signed – the scaling debt load is a reflection of growing demand (and RPO commitments). Strong payback periods on its CapEx investments (~2.5 years) further underpins our view.

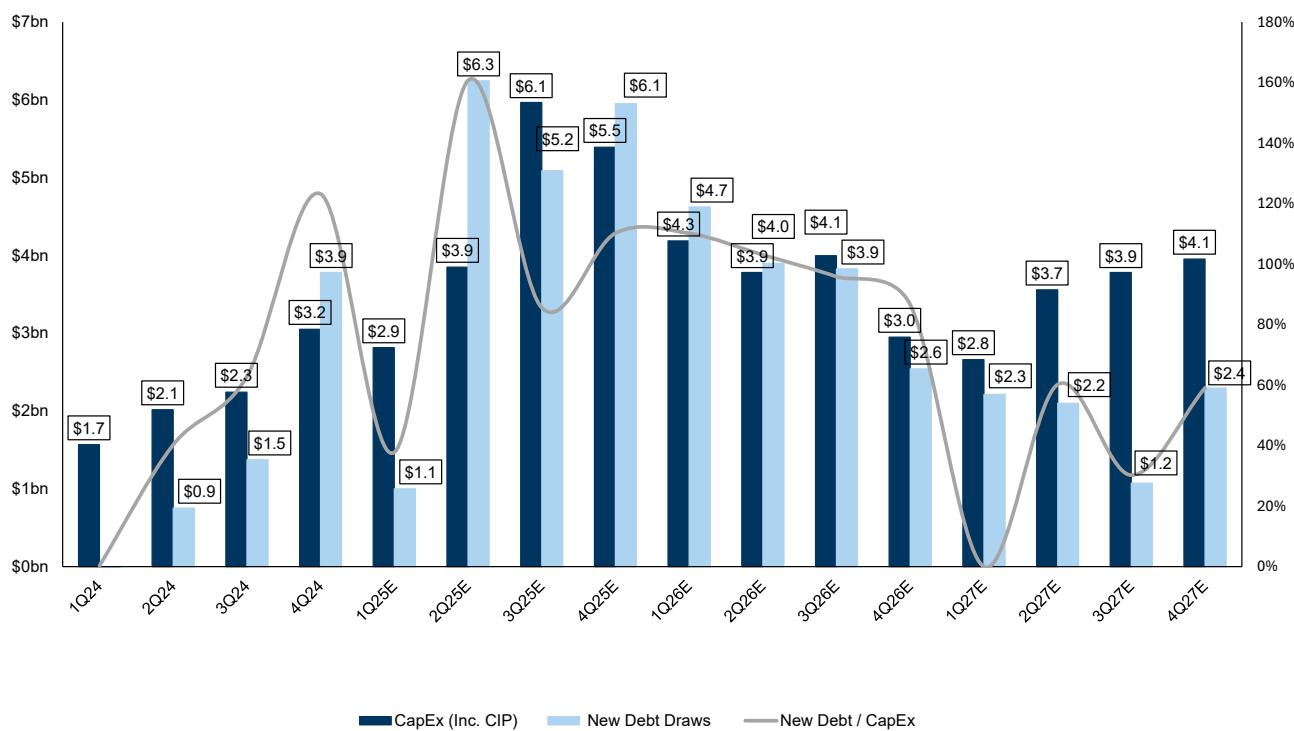
**Debt supported by strong contract dynamics:** As CoreWeave's future debt is largely backed by the cash flows of individual contracts, we are comfortable underwriting a significant ramping of debt levels. First, CoreWeave does not purchase its GPUs ahead of contract signing. Because the company does not purchase components without a committed contract, we see the risk of CoreWeave spending debt on unmonetizable assets are relatively low risk. Furthermore, with an average cash payback period of 2.5 years, individual contract dynamics are quite favorable to quick repayment of debt. CoreWeave's average contract length is ~4 years, which suggests the CoreWeave has a ~1.5 years of higher-flow through cash generation. This is supported by the company's favorable RPO to debt balance – nearly 2x covered in FY24 with \$15.1bn of RPO and \$7.9bn in debt. With the deployment of Blackwell commitments and the addition of the OpenAI contract in 2025, we expect the RPO-to-debt ratio to level at ~116% with \$26.7bn in RPO and \$23.0bn in debt ([Exhibit 42](#)). As CoreWeave continues to sign new contracts over the coming years, we expect the company's debt to RPO balance to remain favorable.

**Exhibit 42: CoreWeave has 116% RPO coverage on FY25 debt, 81% in FY26**  
 RPO (includes OpenAI) as a % of Debt



Source: Company data, Goldman Sachs Global Investment Research

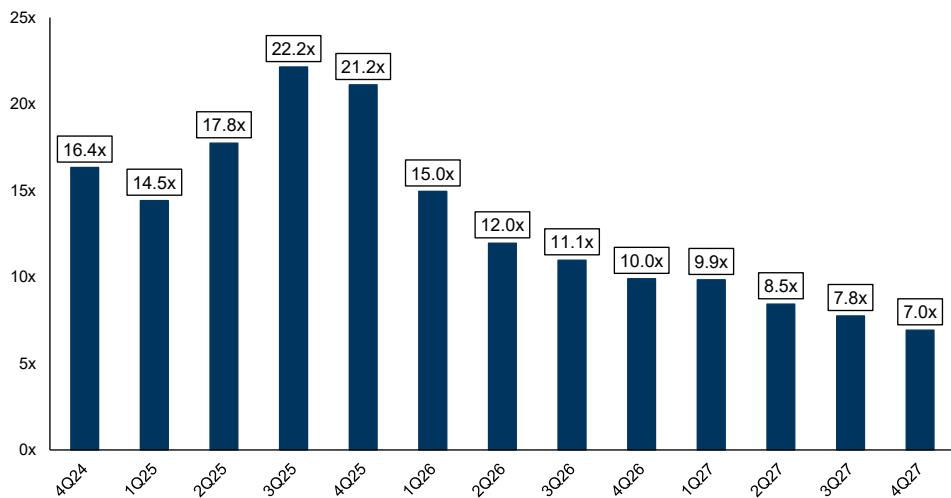
- **CapEx leads debt expansion:** We expect CoreWeave's incremental debt draws to an increasingly lower percentage of CapEx over time. As CoreWeave's business scales, we project that the company's scaling cash flows will more materially offset CoreWeave's future CapEx needs. As CoreWeave sees an average cash payback period of 2.5 years on their ~4 year average contract duration, we see this cash flow beginning to offset CoreWeave's CapEx demands as the company scales.

**Exhibit 43: Expect future CapEx to be less reliant on debt issuances as OCF scales**


Source: Company data, Goldman Sachs Global Investment Research

**Leverage expected to decline meaningfully:** As of FY24, the business was 22.3x levered on a Debt/Non-GAAP EBIT basis. We expect leverage to come down as earnings scale, forecasting 28.3x Debt/Non-GAAP EBIT leverage in FY25, 12.8x in FY26, and 8.4x in FY27. We use Debt/Non-GAAP EBIT as this best represents CoreWeave's earnings generating potential excluding the impact from financing. While debt is expected to increase to \$36.2bn in FY27, we see non-GAAP EBIT scaling significantly to \$4.3bn as CoreWeave's ROI on net assets improves and the company is able to drive better leverage across the P&L.

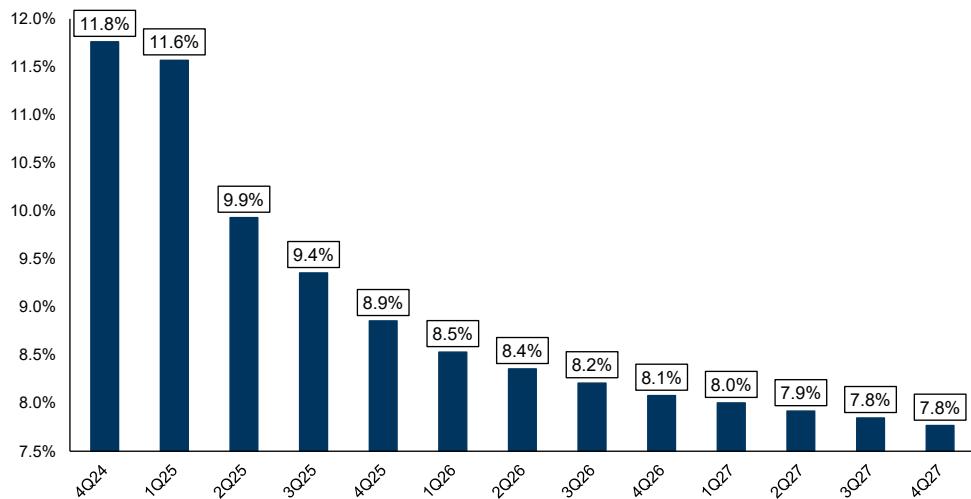
---

**Exhibit 44: Leverage expected to reduce from 16.4x today to 7.0x by 4Q27**


Source: Company data, Goldman Sachs Global Investment Research

**Cost of capital to come down from 11.8% to 7.8% by FY27:** As we expand upon in the following section, while debt expands meaningfully from \$7.9bn to \$36.2bn, we expect the cost of capital to be reduced meaningfully, alleviating some of the pressure of the debt burden. We see a lower cost of capital as a key pillar for the long-term positioning of the company as it can determine CoreWeave's ability to compete against hyperscalers in the longer-term. (more on this in next section)

---

**Exhibit 45: Expect cost of capital to come down from 11.8% today to 7.8% in 4Q27**


Source: Company data, Goldman Sachs Global Investment Research

**Breakdown of debt facilities:** We expect CoreWeave to continue to use a variety of

debt facilities to finance its purchases of networking and compute equipment. Though CoreWeave still has ~\$4bn of undrawn debt (as of December 2024), we expect CoreWeave to raise additional debt in the coming years as their cumulative CapEx spend reaches \$48bn from FY25-FY27. As outlined above, this debt will be raised as CoreWeave continues to purchase infrastructure after signing new contracts. In particular, alongside their existing facilities, we see CoreWeave using a mix of mostly asset-backed securities and parent-level debt (which will become more favorable should the company successfully diversify its customer base) as well as some OEM financing. Over time, we expect CoreWeave's average cost of capital to come down as the company continues to receive more favorable financing arrangements. This will happen as the business scales, CoreWeave receives benefits from the transparency of a public company, and the company builds a longer track-record of creditworthiness.

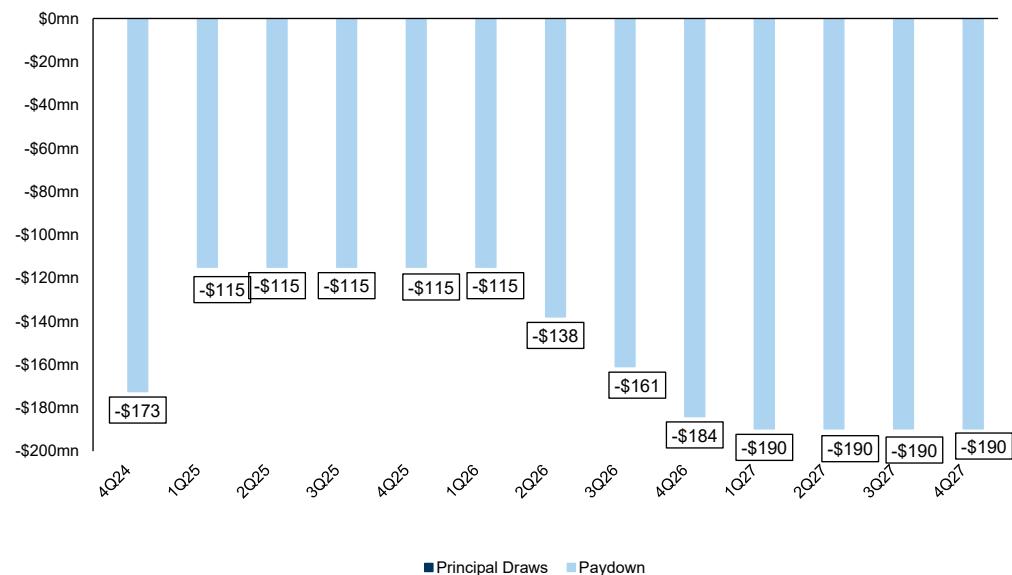
#### Exhibit 46: Breakdown of CoreWeave's current and expected future debt facilities

Debt Facilities				
Name	Size of Facility	Interest Rate	Date Signed	Maturity
DDTL 1.0	\$2.3bn	9.62% + SOFR	July, 2023	March 2028
DDTL 2.0	\$7.6bn	6.00% + SOFR	May, 2024	May 2029
OEM Vendor Financing	\$1.2bn (as of filing)	9-11%	February-December, 2024	--
Revolving Credit Facility	\$0.65bn	1.75% + SOFR	July, 2024	--
Term Loan	\$1bn	6.25% + SOFR	December 2024	--
Future Parent-Level or ABS Debt	--	3.75% + SOFR (GSe)	--	--

Source: Company data, Goldman Sachs Global Investment Research

- **Delayed Draw Term Loan Facility 1.0 (DDTL1):** The Delayed Draw Term Loan Facility 1.0 was entered into in July 2023, providing CoreWeave with up to \$2.3bn in principal, fully drawn as of December 2024. We expect CoreWeave to paydown this debt over time, beginning at 5% of the total balance, maturing in March 2028. Interest is set at 9.62% + SOFR.

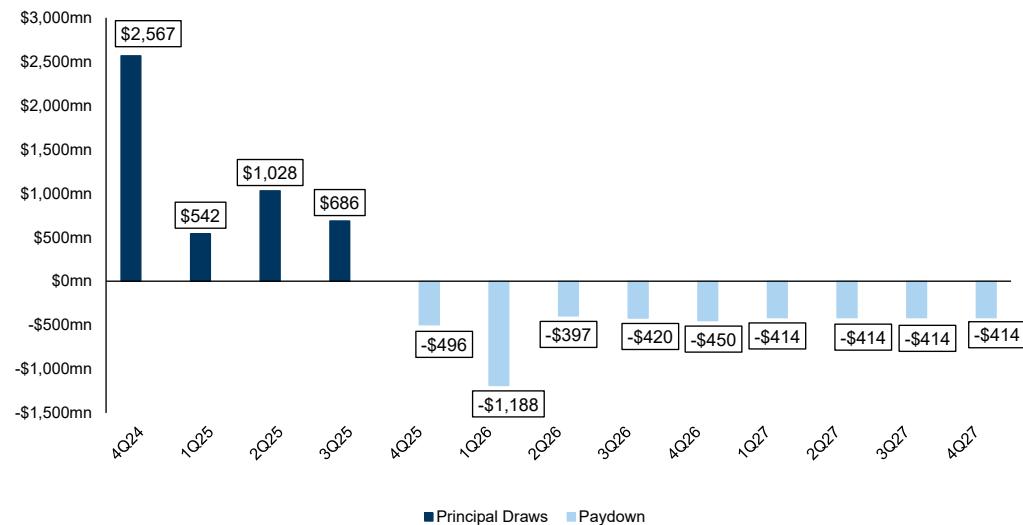
---

**Exhibit 47: DDTL 1.0 expected draws and paydowns each quarter**


Source: Company data, Goldman Sachs Global Investment Research

- **Delayed Draw Term Loan Facility 2.0 (DDTL2):** The Delayed Draw Term Loan Facility 2.0 was entered into in May 2024, providing CoreWeave with up to \$7.6bn in principal, \$3.8bn of which had been drawn by December 2024. We expect CoreWeave to continue to draw from this facility through F3Q25, with quarterly repayments beginning in F4Q25 and Maturity in May 2029. Interest is set at 6.0% + SOFR for specified investment-grade entities (the majority of the principal).

---

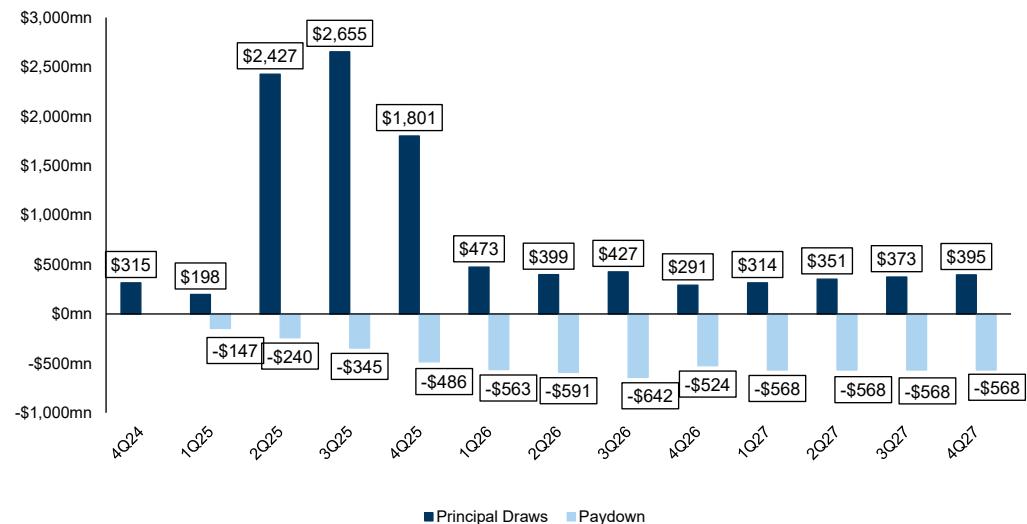
**Exhibit 48: DDTL 2.0 expected draws and paydowns each quarter**


Source: Company data, Goldman Sachs Global Investment Research

- **OEM Vendor Financing:** Between February and December 2024, CoreWeave entered into various financing agreements with Nvidia, with a notional balance of

\$1.2bn as of December 2024. We expect OEM Financing to play a more minor role in future financing than it has historically, especially should the company successfully diversify its customer base. We expect principal draws of \$7.0bn in FY25, \$1.6bn in FY26, and \$1.4bn in FY27. Interest on this debt was 9.2% in F4Q24, and we expect this interest rate to come down gradually over time.

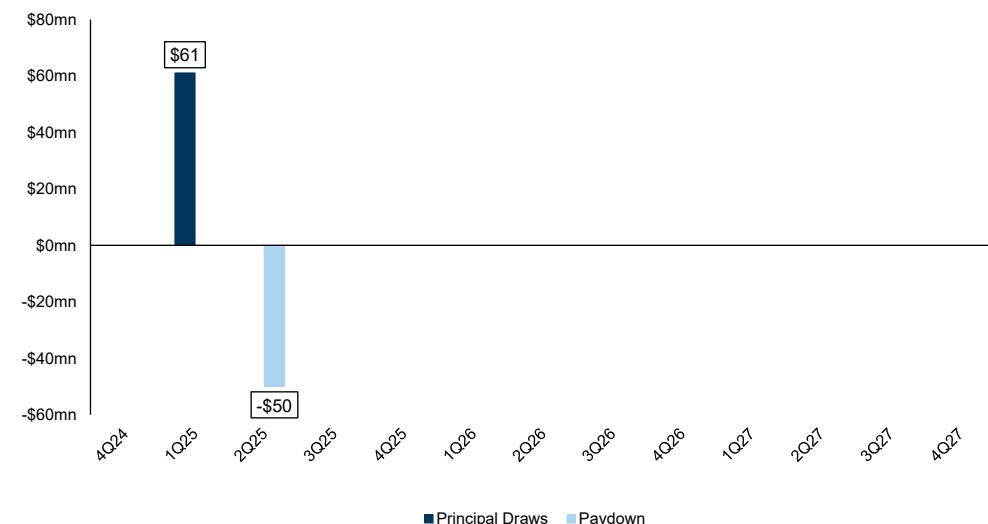
#### Exhibit 49: OEM expected draws and paydowns each quarter



Source: Company data, Goldman Sachs Global Investment Research

- **Revolving Credit Facility (RCF):** The Revolving Credit Facility was entered into in July 2024, providing CoreWeave with up to \$650mn in principal, not drawn as of December 2024. We expect CoreWeave to draw \$61mn of this debt in F1Q25 and pay \$50mn of this debt back in F2Q25. Interest is set at 1.75% + SOFR.

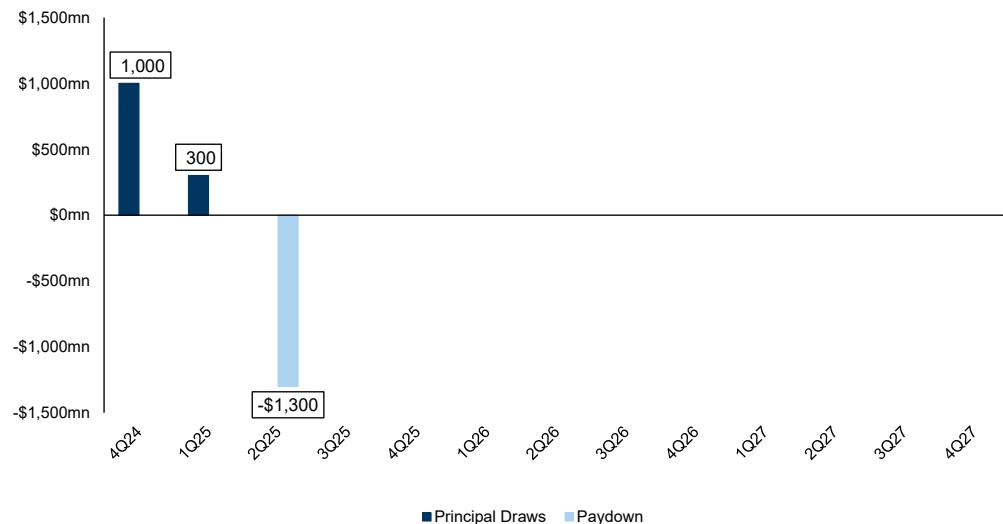
#### Exhibit 50: RCF expected draws and paydowns each quarter



Source: Company data, Goldman Sachs Global Investment Research

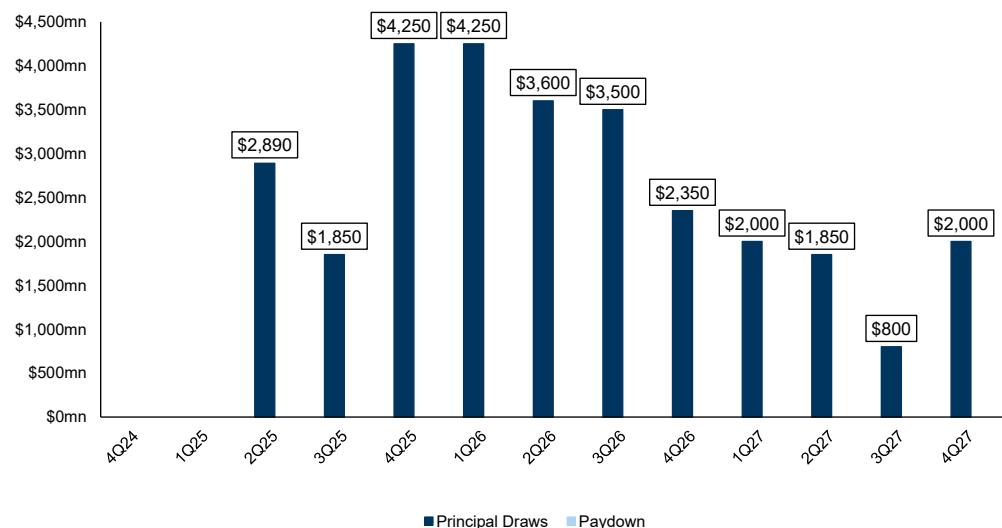
- **Term Loan (Bridge Loan):** The Term Loan was entered into in December 2024, providing CoreWeave with \$1bn in principal (which can be increased to \$1.5bn with the exercise of an uncommitted accordion feature). \$1bn was fully drawn as of December 2024. We expect a further \$300mn to be drawn in F1Q25, and the loan to be fully paid back in F2Q25 with IPO proceeds. Interest is set at 6.25% + SOFR.

**Exhibit 51: Term Loan expected draws and paydowns each quarter**

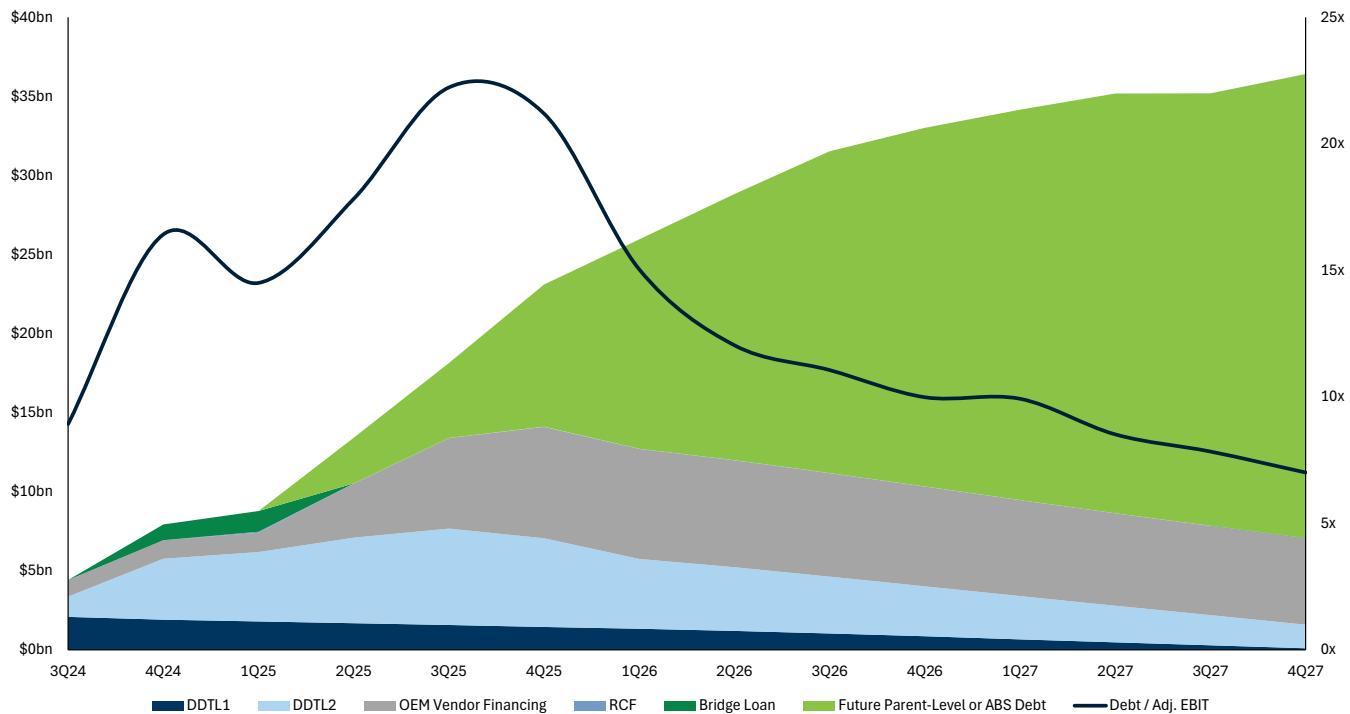


Source: Company data, Goldman Sachs Global Investment Research

- **Future Parent-Level or Asset-Backed Security (ABS) Debt:** Due to the capital demands of the business and limited free cash flow generation, we estimate CoreWeave will need to raise further debt beginning in F2Q25. We estimate the future Parent-Level or ABS debt needed by first calculating Cash Available for Sweep (Beginning Cash + Operating Cash Flow + Investment Cash Flow + Drawdown of Existing Debt Facilities + Issuance of Common Stock – Restricted Cash – Scheduled Debt Payments). We also assume CoreWeave operates with a minimum cash balance of \$2bn considering their prudent financing approach. Once Cash Available for Sweep becomes negative, CoreWeave will need to raise more debt to service its customers and meet its future commitments. Based on this, we layer in our assumptions for new debt raised, with an estimated interest rate for new debt of 3.75% + SOFR in 2025/26/27. We estimate additional debt needed scales from \$9bn in FY25 to \$23bn in FY26, and \$29bn in FY27.

**Exhibit 52: Future Debt expected draws and paydowns each quarter**


Source: Company data, Goldman Sachs Global Investment Research

**Exhibit 53: Total debt expected to scale as revenue and CapEx scale, though with Debt/EBIT leverage coming down meaningfully**


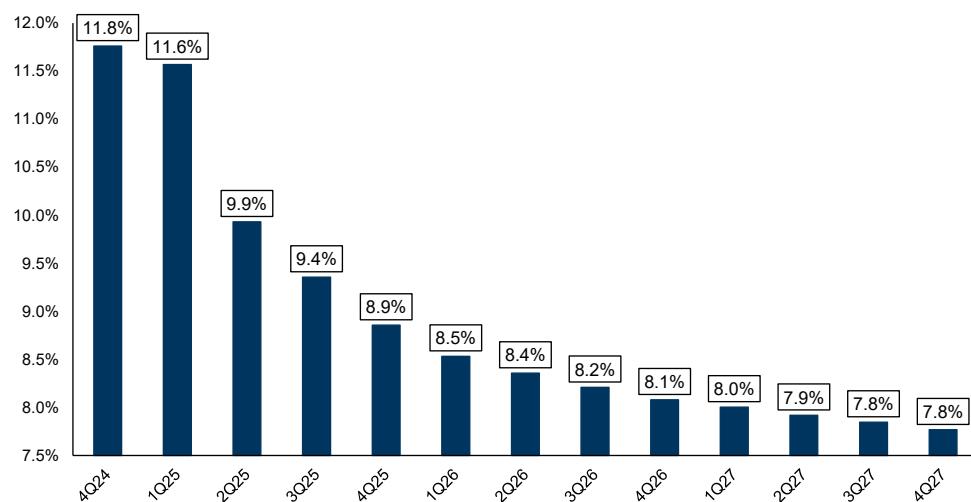
Source: Company data, Goldman Sachs Global Investment Research

## Debt Burden: Path to lower cost of capital

We believe CoreWeave lowering its cost of capital is necessary, with proof points needing to be built before investors are likely to more fully underwrite a low- to mid-single digit CAC. To compete with the hyperscalers in the longer term, it is necessary for CoreWeave to lower its cost of debt. The hyperscalers have access to significantly lower-interest rate debt (longer operating history, diversified revenue base, etc), with a weighted average cost of debt of 3.5% (vs. CoreWeave's 11.8% in FY24). Additionally, all have existing cashflow generating businesses that can support the build out of AI infrastructure without being too reliant on the debt market. As GPU supply constraints normalize, CoreWeave will need to lower the cost of their infrastructure in order to stay competitive on pricing. Our estimates assume the company's weighted average cost of capital comes down to 9.0% in FY25, 8.0% in FY26, and 7.8% in FY27 ([Exhibit 54](#)). CoreWeave's cost of capital may benefit from more favorable interest rate terms on new debt facilities with greater scale, a lengthening operating history, increased transparency and standardization of financial reporting post IPO, liquidity of the public markets, and customer diversification.

**Exhibit 54: Cost of debt expected to come down from 11.8% today to 7.8% by 4Q27**

Weighted Average Cost of CoreWeave's Debt



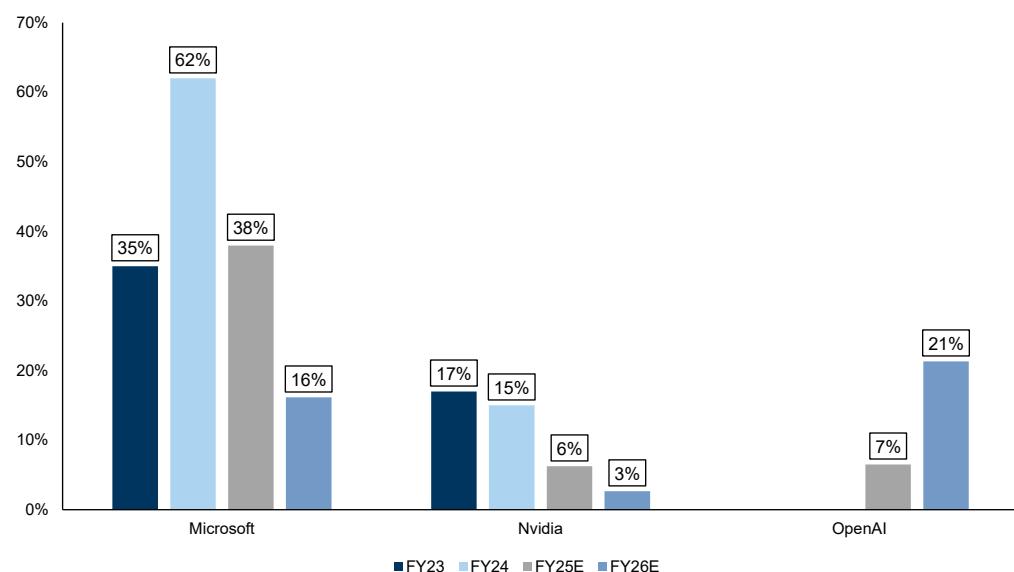
Source: Company data, Goldman Sachs Global Investment Research

- **Improving scale and lengthening operating history:** As CoreWeave continues to scale (GSe top line growth of >100% in 2025/26) and lengthens its operating history, these factors could help support the company's credit perception and allow them to raise future tranches of debt at a lower interest rate. With \$4bn+ in revenue in 2025, and expectations to exceed \$10bn/\$15bn in 25/26, CoreWeave is increasingly becoming a scaled player in the AI ecosystem. This can give the company the flexibility to raise larger principals of debt with at lower risk for creditors. In a similar vein, CoreWeave is building out its credit history, which over time will give creditors

further confidence underwriting their cash flows. We note that CoreWeave was only first launched in 2017, with revenues scaling meaningfully over the last two years. Therefore, the company is still building its track record of delivering against its contracts.

- **IPO increases visibility and liquidity:** CoreWeave's IPO should provide transparency on CoreWeave's business, giving creditors more standardized insight into the company's underlying fundamentals and market reception. With audited financial statements, analyst coverage, and an increasingly scaled internal financial organization, creditors should be able to establish incrementally more confidence in the durability of the business. This could potentially expand the types of creditors the company will appeal to, which may lead to more diversification of financing parties vs. the concentration in a handful of lenders seen pre-IPO. Lastly, the liquidity of an IPO should also give lenders more avenues to hedge risk, and at a lower cost to CoreWeave, which previously provided concessions such as dividends, stock across various class structures, preferred converts, etc.
- **Customer diversification:** Should CoreWeave successfully diversify its customer base, the company may be able to pool customer contracts together to dilute the credit risk of any one contract. We draw a loose similarity to the construct to a mortgage-backed security, in which the cashflows of a group of mortgages are pooled together and borrowed against. The risk of any given contract is diversified across the group of customers and creditors can have more confidence in the broader risk profile of the group of contracts. While we expect CoreWeave's customer base to diversify over time (expanded upon in prior sections), we believe the customer base will remain relatively concentrated between a few large customers for the foreseeable future ([Exhibit 55](#)). For this reason, management has put an emphasis on landing customers with high-credit worthiness already (MSFT, NVDA), before looking to expand to smaller, and in some cases, newer cohorts of companies.

**Exhibit 55: Expect majority of CoreWeave's revenue to come from a few large customers for the foreseeable future**



FY25/FY26 Microsoft assumes revenue grows to \$1.8bn/year, the midpoint between MSFT contribution in FY24 and OpenAI's yearly revenue contribution. FY25/FY26 Nvidia assumes revenue contribution is flat. FY25/FY26 OpenAI revenue is GSe.

Source: Company data, Goldman Sachs Global Investment Research

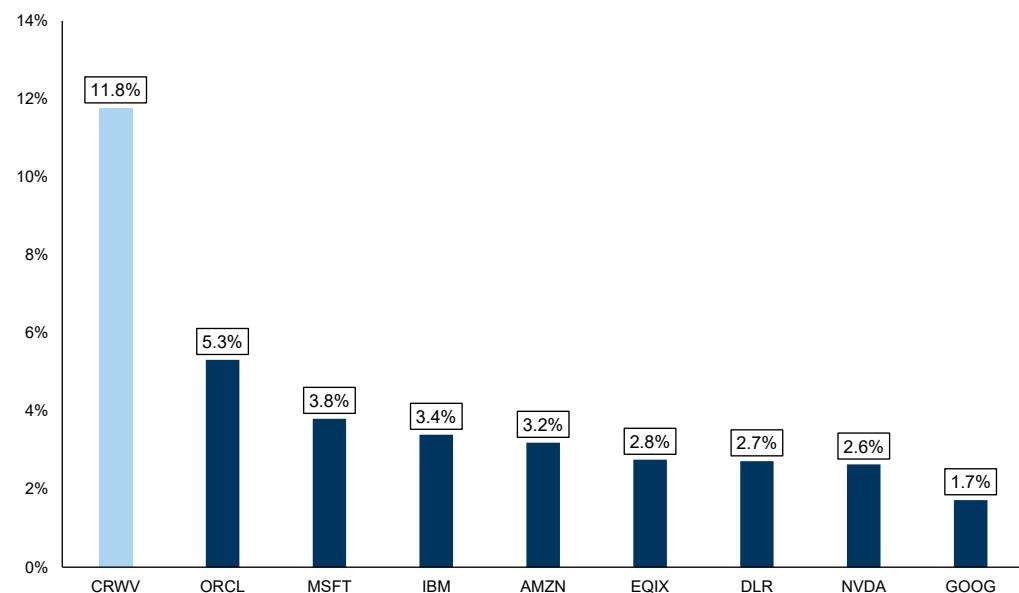
■ **Risks that can elongate trajectory toward lower Cost of Capital:** Though we are comfortable underwriting an improving cost of capital, CoreWeave still has a long path before it reaches a comparable weighted average cost of capital to peers. CoreWeave's weighted average cost of debt was 11.8% at the end of 2024, vs. an average of 3.2% for peers ([Exhibit 56](#)). The following benefits these peers hold underpins their lower cost of financing and more stable debt costs (via fixed rates): longer operating history and track record, broadly diversified customer base, and highly profitable adjacent businesses (enterprise software, search, and e-commerce) that can support the capital intensity of their cloud investments. CoreWeave does not benefit from similar dynamics. In addition to the evolution of the dynamics laid out in the prior paragraph, we offer additional risks that can elongate the trajectory on which CoreWeave will be able to achieve a low- to mid-single digit cost of debt:

- **Success of GTM:** CoreWeave's go-to-market motion is still very nascent. With non-GAAP S&M spend accounting for less than 1% of revenue, CoreWeave has been able to benefit from demand conversion due to the supply constraints in market. The pace in which the company will be able to scale its more directed sales motion will likely determine when they will see a more diversified customer base that can adapt the terms available to them for financing. Today, since CoreWeave's customer base is much more concentrated, there is a higher weighting on any one customer not fulfilling its commitments on CoreWeave's overall financing options.
- **Tariffs:** Since CRWV is heavily reliant on procuring GPUs and data centers, the input costs that go into building training and inference systems for its customers face evolving uncertainty due to tariffs. Their ability to pass along

higher COGS and tech and Infrastructure costs to their end customers will be a significant source of uncertainty. While the larger hyperscalers face this uncertainty as well, they have much more diversified customer bases, better ability to cushion spikes in component costs and more avenues to mitigate partially if not fully through pricing their more elaborate cloud offerings more appropriately.

If maintained, the roughly 8pp delta between CoreWeave and the hyperscalers can make it difficult for CoreWeave to maintain its positioning and compete with the hyperscalers in a market with high entry costs.

**Exhibit 56: CoreWeave's cost of capital is significantly higher than its competitors**  
Weighted Average Cost of Debt



Source: Company data, Goldman Sachs Global Investment Research

- **Earnings show high sensitivity to cost of capital:** With interest expense accounting for 59% of non-GAAP EBIT in FY27, any shift in CoreWeave's cost of capital can have a meaningful impact on the earnings profile of the company. In order to gauge this impact, we present an illustrative sensitivity analysis below. The analysis aims to show how the company's leverage profile impacts the earnings power of the company. At 3.65% SOFR and 3.75% cost of capital (our modeled assumptions for FY27), earnings are expected to be \$2.38 or roughly 66% lower than if the company would be able to generate the same margin profile with no leverage. As we evaluate different scenarios, we see that any ~100bps change to either cost of capital or SOFR would have a 40-50 cent impact to FY27 EPS.

**Exhibit 57: Estimate a 1pp change in FY27 SOFR or debt cost would result in a \$0.40 - \$0.50 impact to FY27 EPS**

Interest Expense Impact to EPS with Different SOFR Estimates and Cost of Future Debt

SOFR (FY27)	Future Debt Cost (+ SOFR)				
	1.75%	2.75%	3.75%	4.75%	5.75%
1.65%	-\$2.77	-\$3.21	-\$3.65	-\$4.10	-\$4.54
2.65%	-\$3.26	-\$3.70	-\$4.14	-\$4.59	-\$5.03
3.65%	-\$3.75	-\$4.19	-\$4.64	-\$5.08	-\$5.52
4.65%	-\$4.24	-\$4.68	-\$5.13	-\$5.57	-\$6.01
5.65%	-\$4.73	-\$5.18	-\$5.62	-\$6.06	-\$6.50

Source: Company data, Goldman Sachs Global Investment Research

# Valuation

## Peer Group

CoreWeave is a singularity as an AI cloud hyperscaler, operating at scale with estimated 141% revenue growth and 18% non-GAAP operating margins this year. As such, we look at a blend of the following categories to help us determine the right relative valuation for the company:

- **General-purpose cloud providers (Microsoft, Alphabet, Google, Oracle, IBM, Digital Ocean):** The hyperscalers operate AI clouds like CoreWeave, but they are also at-scale businesses with other major business lines (IBM falls in a similar vein). Digital Ocean is the only pure-play publicly traded cloud provider, but its customers are significantly smaller in scale (their largest cohort is \$8k+ customers). If CoreWeave is able to execute on its financial model, successfully navigate the shift from training to inference, and demonstrate capital deployment efficiency, we believe CoreWeave can be valued similarly to this group of companies (average of 16x EV/EBIT).
- **High-growth hardware (Nvidia):** Nvidia operates in a similar high-growth, asset-heavy business with significant exposure to AI-related revenue. We note that Nvidia offers component hardware vs. CoreWeave's compute and infrastructure software services.
- **Value-add Hardware redistributors (CDW, Arrow Electronics, Avnet):** The value-add hardware redistributors add services on top of tech hardware distribution, are heavily levered, and low-margin. In a scenario where AI compute is commoditized, CoreWeave could see their margins compress while maintaining a similar leverage profile.

We also expect investors to include other Gen-AI native peers as they come to market. We exclude Nebius (NBIS) from our valuation methodology given the limited consensus estimates available.

In aggregate, we believe these represent a fair peer group that can help investors triangulate CoreWeave's fair value across the key factors likely to be top of mind for investors: 1) high-growth 2) asset-heavy business model, with high CapEx, depreciation and re-investment costs and 3) leveraged capital structure with a more commoditized product.

**Exhibit 58: CoreWeave comps table**

	<u>Sales Growth</u>			<u>EBIT Margin</u>			<u>EV / EBIT</u>		
	CY25	CY26	CY27	CY25	CY26	CY27	CY25	CY26	CY27
AMZN	9%	10%	11%	11%	13%	14%	23.1x	19.3x	15.3x
ARW	-1%	7%	13%	3%	4%	4%	9.2x	6.9x	5.9x
AVT	-1%	7%		3%	4%		9.3x	7.5x	7.1x
CDW	2%	5%	4%	9%	9%	9%	13.1x	12.3x	11.8x
DOCN	13%	14%	16%	26%	29%	27%	15.6x	12.7x	11.4x
GOOG	10%	10%	11%	33%	33%	34%	14.1x	12.7x	11.1x
IBM	4%	3%	5%	19%	19%	20%	21.7x	20.5x	18.9x
MSFT	12%	13%	17%	45%	46%	45%	20.7x	18.0x	15.7x
NVDA	53%	22%	18%	64%	67%	64%	18.5x	14.6x	12.9x
ORCL	10%	16%	22%	43%	43%	46%	16.7x	14.5x	10.9x
<b>Average</b>	<b>11%</b>	<b>11%</b>	<b>13%</b>	<b>26%</b>	<b>27%</b>	<b>29%</b>	<b>16.2x</b>	<b>13.9x</b>	<b>12.1x</b>
<b>CoreWeave</b>	<b>141%</b>	<b>135%</b>	<b>43%</b>	<b>11%</b>	<b>20%</b>	<b>25%</b>	<b>51.3x</b>	<b>12.3x</b>	<b>6.8x</b>

Comps as of 4/21/2025

Source: Company data, Goldman Sachs Global Investment Research

**Valuation Methodology**

We leverage an EV/EBIT valuation methodology to arrive at our valuation for CoreWeave. Focusing on EBIT allows investors to account for CoreWeave's asset-heavy business model and account for the impact of depreciation.

This also reflects investors' wariness of the long-term growth durability of the company given it does not account for the high-growth trajectory of CoreWeave over the next three years.

Our approach balances various scenarios around CoreWeave's future market positioning as the peer group has a blend of software solutions (a long-term growth investment for CoreWeave that represent the company's long-term potential) and more value-add solutions (which would represent a downside scenario where AI compute becomes more commoditized). We note, however, that our comp group is 70% weighted towards high growth/cloud-focused tech companies and only 30% weighted towards the value-add hardware resellers.

**This leads us to land at a PT of \$54.** We set a range of EV/EBIT (GAAP) multiples based on where the comparable companies in CoreWeave's peer group are trading on Consensus' Q5 – Q8 (second next-twelve-months). We took an average multiple (13.9x) of this group and applied that to CoreWeave's 2027 EBIT (\$3.9bn) to arrive at a 2026 enterprise value of \$53.9bn. We subtract 1Q26 net debt of \$23.0bn to arrive at a 12-month out equity value of \$30.9bn and a value per share of \$54.

**Exhibit 59: Use an EV / EBIT multiple of 13.9x to reach a PT of \$54**  
 \$mn except per share data

<b>EV / EBIT</b>	
FY25 EBIT	510
FY26 EBIT	2,127
FY27 EBIT	3,872
FY26 (STM) Comps EV / EBIT	<b>13.9x</b>
EV (NTM Multiple*STM EBIT)	53,815
Implied CRWV FY26 EV / EBIT	25.3x
Net Debt (1Q26)	22,967
<b>NTM Equity Value</b>	<b>30,848</b>

<b>Implied PT</b>	<b>54.00</b>
Current Price	<b>35.42</b>
<b>Implied upside</b>	<b>52%</b>

Source: Company data, Goldman Sachs Global Investment Research

As CoreWeave scales and progresses toward net income, investors may pivot to P/E to best account for interest expense and the leverage of the business. We currently expect CoreWeave to generate non-GAAP net income for the first time in 2027.

## Key Risks

- **Durability of Gen-AI as the next step function in technology.** CoreWeave's current offerings and investments are allocated to driving innovation around how high-powered, GPU-intensive Gen-AI workloads are serviced. Any changes in this demand are likely to highly affect CoreWeave's long-term growth. In such a scenario, the company's components may be repurposed to other end-market use cases that may need similar components. Nevertheless, the company expects to diversify and expand its customer base into industries such as pharmaceuticals and financial services over time.
- **Increased competition may risk diminishing value differentiation.** As additional Gen-AI capacity comes online via other cloud providers' investment in first-party data centers (expected in 2H), CoreWeave faces a risk of increased competition. As more capacity comes online across a number of providers, the market is likely to expect more demand to be met (driving the industry toward supply/demand toward equilibrium), reducing pricing power. CoreWeave will need to lean into its competitive differentiation (accelerated time-to-market, superior performance, specialized-infrastructure, etc.) to show its distinguished value proposition and continued right to win.
- **Tariffs may increase input costs, supply chain timelines.** Since CRWV is heavily reliant on procuring GPUs and data centers, the input costs that go into building training and inference systems for its customers face evolving uncertainty due to tariffs. Their ability to pass along higher COGS and tech and Infrastructure costs to their end customers will be a significant source of uncertainty. While the larger hyperscalers face this uncertainty as well, they have much more diversified customer bases, are better able to cushion spikes in component costs and have more avenues to mitigate partially if not fully through pricing their more elaborate cloud offerings more appropriately.
- **Continued balance across all stakeholders. CoreWeave's financial structure is tightly dependent on an interwoven network of stakeholders.** Any potential disruption to any one part of the business can have outsized impact across all elements of the organization. CoreWeave currently engages with a number of financing parties, partners, suppliers, and customers and adequately balances the needs of each one. Since the company operates in an asset-heavy business, CoreWeave requires numerous financing partners and suppliers. If supply chain headwinds, interest rate variability or customer demand shifts, it can have significant impact to the ongoing operations of CoreWeave's business. CoreWeave is currently managing these risks well by: 1) having a large mix of committed contract revenue (>95%) 2) high concentration of large established customers (with Microsoft accounting for over 62% of revenue in CY24), 3) having tight supply chain relationships and 4) procuring components on a just-in-time basis. Should any of these pillars face challenges (i.e. supply chain delays) or the profile of CoreWeave's customer base evolves (as we move through the Gen-AI adoption cycle and company diversifies its customer base), the company's credit risk may increase. This

can then impact future financing needs, which CoreWeave needs to meet customer commitments (with future debt needs reflected in our estimates). Therefore, maintaining a balance across these stakeholders is imperative for CoreWeave to meet current expectations.

- **Constrained procurement of additional compute and data center equipment amid limited quantity, high competition.** CoreWeave risks not being able to meet additional customer needs and therefore drive incremental growth should it not successfully procure the necessary components to service the project. The company's close supply chain relationships will continue to be important as constraints around key components are unrelenting. In particular, as three suppliers account for 46%, 16% and 14% of purchases in FY24, impacts to any of those relationships would have a material effect of CoreWeave's operations. The rapid pace of innovation and desired adoption of Gen-AI are outpacing the constrained production capabilities of building out GPUs, bringing power online and building out data centers, the main components of AI-based data centers. CoreWeave acquires the needed GPUs on a just-in-time basis based on customer contracts. Its close partnership with Nvidia is a key pillar of this framework. For power and data center space, CoreWeave deploys at-risk capital (not backed by matching a customer contract with spend) to get ahead of demand. CoreWeave current has enough capacity to meet existing commitments but should bottlenecks or competition increase, there is a risk to future growth. As of December 2024, CoreWeave had 1.3 GW of contracted power with only 360MW of it active. Lastly, due to more supply is coming online from hyperscalers' prior investments over the last two years, increased visibility into accessibility of these parameters may be needed for CoreWeave to have line of sight to its future capacity and increase confidence around their ability to meet customers' needs, which can increase customer confidence around choosing CoreWeave.
- **Pace of innovation on silicone layer may compress useful life of computing equipment assets.** CoreWeave's computing equipment is depreciated over a six-year period. As it stands, the company's average contract duration is ~4 years, with a payback period of 2.5 years on its GPUs. This allows the company to drive additional revenue on assets written off the balance sheet. Should the pace of innovation on the chip layer accelerate meaningfully, CoreWeave may need to adjust this accounting assumption, which would impact the level of depreciation recognized and compress the average contract duration per chip generation. Investors are likely to gain visibility into this dynamic as a growing mix of GPUs pass the 2.5 payback period and 4-year contract duration time frames over the next few years. For reference, CoreWeave's fleet of GPUs has grown to over 250K by year-end 2024, from 53K in 2023 and 17K GPUs in 2022.
- **Rapid growth may put a strain on company's finance function.** CoreWeave's exponential growth – up 121x from 2022-2024 – may have put a strain on the company's finance function. CoreWeave's non-GAAP G&A expense has only increased ~19x over the same period suggesting that the support function around procurement and financing may not have kept up with the company's needs. CoreWeave has addressed initiatives to improve internal controls throughout 2025

and 2026. In addition to expanding accounting, finance, and operations functions to create and execute business processes, CoreWeave is investing in building out IT General Controls and increased technology deployment across the company.

- **High exposure to variable rate debt vehicles increases exposure to volatile macro environment, risk of elevated cost of capital + interest expense.** The volatility in current macro conditions may present a headwind for CoreWeave should interest rates step above current expectations. The company's high leverage profile, with \$8bn in debt, is highly exposed to market conditions considering the floating interest rate of all credit facilities. The most common benchmark across CoreWeave's various debt obligations is the Secured Overnight Financing Rate (SOFR). If the cost of debt increases substantially, this may limit CoreWeave's ability to invest in revenue-generating areas and impact financial metrics such as EPS. Interest Expense accounted for ~5% of the company's total principal in 2024. This can also impact the availability and cost of future capital (currently expected to be ~3.75% + SOFR), and therefore the trajectory of improvement in the company's average cost of capital. We currently expect this to decline from 11.8% in CY24 to 9.0%, 8.1% and 7.8% in CY25/26/27, respectively.
- **Business needs may continue to drive high rate of cash burn, push out FCF profitability.** Given the asset-heavy nature of CoreWeave's business, the company's investment in fixed assets has a direct correlation with the future revenue growth of the business. For this reason, our estimates do not expect material improvement in the business' free cash flow through FY26. Estimates assume FCF remains in the negative \$3-16bn range through FY27. As long as business prospects continue to require CapEx investments, the achievability of FCF profitability may continue to be pushed out. CoreWeave's nascent software solutions, which are currently not being monetized, are a long-term lever that may offset these FCF dynamics. Software can be scaled without incremental CapEx investments, making it a growth avenue that can improve unit economics and drive FCF improvement should management flow the high-margin revenue through to FCF. Largely, investors will need to be comfortable with the potential for the business' current cash burn rate to persist until clarity is provided around either CoreWeave's ability to reduce its dependency on net asset growth to drive revenue, or management reduces its reinvestment ratio.

# Disclosure Appendix

## Reg AC

We, Kash Rangan, Gili Naftalovich, Henry Dane, Matthew Martino, Selina Zhang and Nishad Patwardhan, hereby certify that all of the views expressed in this report accurately reflect our personal views about the subject company or companies and its or their securities. We also certify that no part of our compensation was, is or will be, directly or indirectly, related to the specific recommendations or views expressed in this report.

Unless otherwise stated, the individuals listed on the cover page of this report are analysts in Goldman Sachs' Global Investment Research division.

## GS Factor Profile

The Goldman Sachs Factor Profile provides investment context for a stock by comparing key attributes to the market (i.e. our universe of rated stocks) and its sector peers. The four key attributes depicted are: Growth, Financial Returns, Multiple (e.g. valuation) and Integrated (a composite of Growth, Financial Returns and Multiple). Growth, Financial Returns and Multiple are calculated by using normalized ranks for specific metrics for each stock. The normalized ranks for the metrics are then averaged and converted into percentiles for the relevant attribute. The precise calculation of each metric may vary depending on the fiscal year, industry and region, but the standard approach is as follows:

**Growth** is based on a stock's forward-looking sales growth, EBITDA growth and EPS growth (for financial stocks, only EPS and sales growth), with a higher percentile indicating a higher growth company. **Financial Returns** is based on a stock's forward-looking ROE, ROCE and CROCI (for financial stocks, only ROE), with a higher percentile indicating a company with higher financial returns. **Multiple** is based on a stock's forward-looking P/E, P/B, price/dividend (P/D), EV/EBITDA, EV/FCF and EV/Debt Adjusted Cash Flow (DACP) (for financial stocks, only P/E, P/B and P/D), with a higher percentile indicating a stock trading at a higher multiple. The **Integrated** percentile is calculated as the average of the Growth percentile, Financial Returns percentile and (100% - Multiple percentile).

Financial Returns and Multiple use the Goldman Sachs analyst forecasts at the fiscal year-end at least three quarters in the future. Growth uses inputs for the fiscal year at least seven quarters in the future compared with the year at least three quarters in the future (on a per-share basis for all metrics).

For a more detailed description of how we calculate the GS Factor Profile, please contact your GS representative.

## M&A Rank

Across our global coverage, we examine stocks using an M&A framework, considering both qualitative factors and quantitative factors (which may vary across sectors and regions) to incorporate the potential that certain companies could be acquired. We then assign a M&A rank as a means of scoring companies under our rated coverage from 1 to 3, with 1 representing high (30%-50%) probability of the company becoming an acquisition target, 2 representing medium (15%-30%) probability and 3 representing low (0%-15%) probability. For companies ranked 1 or 2, in line with our standard departmental guidelines we incorporate an M&A component into our target price. M&A rank of 3 is considered immaterial and therefore does not factor into our price target, and may or may not be discussed in research.

## Quantum

Quantum is Goldman Sachs' proprietary database providing access to detailed financial statement histories, forecasts and ratios. It can be used for in-depth analysis of a single company, or to make comparisons between companies in different sectors and markets.

## Disclosures

### Logo disclosure

Third party brands used in this report are the property of their respective owners, and are used here for informational purposes only. The use of such brands should not be viewed as an endorsement, affiliation or sponsorship by or for Goldman Sachs or any of its products/services.

**The rating(s) for CoreWeave Inc. is/are relative to the other companies in its/their coverage universe:** Adobe Inc., Atlassian Corp., Autodesk Inc., Bentley Systems Inc., Confluent Inc., CoreWeave Inc., Couchbase Inc., Datadog Inc., Definitive Healthcare Corp., Dropbox Inc., Dynatrace Inc., Elastic NV, GitLab Inc., Informatica Inc., Intuit Inc., Microsoft Corp., Monday.Com Ltd., MongoDB Inc., Oracle Corp., PagerDuty, Procure Technologies Inc., RingCentral, Rubrik Inc., Salesforce Inc., Samsara Inc., ServiceNow Inc., ServiceTitan Inc., Snowflake Inc., Twilio, Unity Software Inc., Weave Communications Inc., Workday Inc., Zoom Video Communications Inc., ZoomInfo Technologies Inc.

### Company-specific regulatory disclosures

The following disclosures relate to relationships between The Goldman Sachs Group, Inc. (with its affiliates, "Goldman Sachs") and companies covered by Goldman Sachs Global Investment Research and referred to in this research.

Goldman Sachs has received compensation for investment banking services in the past 12 months: CoreWeave Inc. (\$35.42)

Goldman Sachs expects to receive or intends to seek compensation for investment banking services in the next 3 months: CoreWeave Inc. (\$35.42)

Goldman Sachs had an investment banking services client relationship during the past 12 months with: CoreWeave Inc. (\$35.42)

Goldman Sachs has managed or co-managed a public or Rule 144A offering in the past 12 months: CoreWeave Inc. (\$35.42)

Goldman Sachs makes a market in the securities or derivatives thereof: CoreWeave Inc. (\$35.42)

## Distribution of ratings/investment banking relationships

Goldman Sachs Investment Research global Equity coverage universe

	Rating Distribution			Investment Banking Relationships		
	Buy	Hold	Sell	Buy	Hold	Sell
Global	49%	34%	17%	63%	57%	42%

As of April 1, 2025, Goldman Sachs Global Investment Research had investment ratings on 3,016 equity securities. Goldman Sachs assigns stocks as Buys and Sells on various regional Investment Lists; stocks not so assigned are deemed Neutral. Such assignments equate to Buy, Hold and Sell for the purposes of the above disclosure required by the FINRA Rules. See 'Ratings, Coverage universe and related definitions' below. The Investment

Banking Relationships chart reflects the percentage of subject companies within each rating category for whom Goldman Sachs has provided investment banking services within the previous twelve months.

## Regulatory disclosures

### Disclosures required by United States laws and regulations

See company-specific regulatory disclosures above for any of the following disclosures required as to companies referred to in this report: manager or co-manager in a pending transaction; 1% or other ownership; compensation for certain services; types of client relationships; managed/co-managed public offerings in prior periods; directorships; for equity securities, market making and/or specialist role. Goldman Sachs trades or may trade as a principal in debt securities (or in related derivatives) of issuers discussed in this report.

The following are additional required disclosures: **Ownership and material conflicts of interest:** Goldman Sachs policy prohibits its analysts, professionals reporting to analysts and members of their households from owning securities of any company in the analyst's area of coverage.

**Analyst compensation:** Analysts are paid in part based on the profitability of Goldman Sachs, which includes investment banking revenues. **Analyst as officer or director:** Goldman Sachs policy generally prohibits its analysts, persons reporting to analysts or members of their households from serving as an officer, director or advisor of any company in the analyst's area of coverage. **Non-U.S. Analysts:** Non-U.S. analysts may not be associated persons of Goldman Sachs & Co. LLC and therefore may not be subject to FINRA Rule 2241 or FINRA Rule 2242 restrictions on communications with subject company, public appearances and trading securities held by the analysts.

**Distribution of ratings:** See the distribution of ratings disclosure above. **Price chart:** See the price chart, with changes of ratings and price targets in prior periods, above, or, if electronic format or if with respect to multiple companies which are the subject of this report, on the Goldman Sachs website at <https://www.gs.com/research/hedge.html>.

### Additional disclosures required under the laws and regulations of jurisdictions other than the United States

The following disclosures are those required by the jurisdiction indicated, except to the extent already made above pursuant to United States laws and regulations. **Australia:** Goldman Sachs Australia Pty Ltd and its affiliates are not authorised deposit-taking institutions (as that term is defined in the Banking Act 1959 (Cth)) in Australia and do not provide banking services, nor carry on a banking business, in Australia. This research, and any access to it, is intended only for "wholesale clients" within the meaning of the Australian Corporations Act, unless otherwise agreed by Goldman Sachs. In producing research reports, members of Global Investment Research of Goldman Sachs Australia may attend site visits and other meetings hosted by the companies and other entities which are the subject of its research reports. In some instances the costs of such site visits or meetings may be met in part or in whole by the issuers concerned if Goldman Sachs Australia considers it is appropriate and reasonable in the specific circumstances relating to the site visit or meeting. To the extent that the contents of this document contains any financial product advice, it is general advice only and has been prepared by Goldman Sachs without taking into account a client's objectives, financial situation or needs. A client should, before acting on any such advice, consider the appropriateness of the advice having regard to the client's own objectives, financial situation and needs. A copy of certain Goldman Sachs Australia and New Zealand disclosure of interests and a copy of Goldman Sachs' Australian Sell-Side Research Independence Policy Statement are available at: <https://www.goldmansachs.com/disclosures/australia-new-zealand/index.html>. **Brazil:** Disclosure information in relation to CVM Resolution n. 20 is available at <https://www.gs.com/worldwide/brazil/area/gir/index.html>. Where applicable, the Brazil-registered analyst primarily responsible for the content of this research report, as defined in Article 20 of CVM Resolution n. 20, is the first author named at the beginning of this report, unless indicated otherwise at the end of the text. **Canada:** This information is being provided to you for information purposes only and is not, and under no circumstances should be construed as, an advertisement, offering or solicitation by Goldman Sachs & Co. LLC for purchasers of securities in Canada to trade in any Canadian security. Goldman Sachs & Co. LLC is not registered as a dealer in any jurisdiction in Canada under applicable Canadian securities laws and generally is not permitted to trade in Canadian securities and may be prohibited from selling certain securities and products in certain jurisdictions in Canada. If you wish to trade in any Canadian securities or other products in Canada please contact Goldman Sachs Canada Inc., an affiliate of The Goldman Sachs Group Inc., or another registered Canadian dealer. **Hong Kong:** Further information on the securities of covered companies referred to in this research may be obtained on request from Goldman Sachs (Asia) L.L.C. **India:** Further information on the subject company or companies referred to in this research may be obtained from Goldman Sachs (India) Securities Private Limited, Research Analyst - SEBI Registration Number INH000001493, 951-A, Rational House, Appasaheb Marathe Marg, Prabhadevi, Mumbai 400 025, India, Corporate Identity Number U74140MH2006FTC160634, Phone +91 22 6616 9000, Fax +91 22 6616 9001. Goldman Sachs may beneficially own 1% or more of the securities (as such term is defined in clause 2 (h) the Indian Securities Contracts (Regulation) Act, 1956) of the subject company or companies referred to in this research report. Investment in securities market are subject to market risks. Read all the related documents carefully before investing. Registration granted by SEBI and certification from NISM in no way guarantee performance of the intermediary or provide any assurance of returns to investors. Goldman Sachs (India) Securities Private Limited compliance officer and investor grievance contact details can be found at this link: <https://publishing.gs.com/disclosures/hedge.html#/general/equity>. **Japan:** See below. **Korea:** This research, and any access to it, is intended only for "professional investors" within the meaning of the Financial Services and Capital Markets Act, unless otherwise agreed by Goldman Sachs. Further information on the subject company or companies referred to in this research may be obtained from Goldman Sachs (Asia) L.L.C., Seoul Branch. **New Zealand:** Goldman Sachs New Zealand Limited and its affiliates are neither "registered banks" nor "deposit takers" (as defined in the Reserve Bank of New Zealand Act 1989) in New Zealand. This research, and any access to it, is intended for "wholesale clients" (as defined in the Financial Advisers Act 2008) unless otherwise agreed by Goldman Sachs. A copy of certain Goldman Sachs Australia and New Zealand disclosure of interests is available at: <https://www.goldmansachs.com/disclosures/australia-new-zealand/index.html>. **Russia:** Research reports distributed in the Russian Federation are not advertising as defined in the Russian legislation, but are information and analysis not having product promotion as their main purpose and do not provide appraisal within the meaning of the Russian legislation on appraisal activity. Research reports do not constitute a personalized investment recommendation as defined in Russian laws and regulations, are not addressed to a specific client, and are prepared without analyzing the financial circumstances, investment profiles or risk profiles of clients. Goldman Sachs assumes no responsibility for any investment decisions that may be taken by a client or any other person based on this research report. **Singapore:** Goldman Sachs (Singapore) Pte. (Company Number: 198602165W), which is regulated by the Monetary Authority of Singapore, accepts legal responsibility for this research, and should be contacted with respect to any matters arising from, or in connection with, this research. **Taiwan:** This material is for reference only and must not be reprinted without permission. Investors should carefully consider their own investment risk. Investment results are the responsibility of the individual investor. **United Kingdom:** Persons who would be categorized as retail clients in the United Kingdom, as such term is defined in the rules of the Financial Conduct Authority, should read this research in conjunction with prior Goldman Sachs research on the covered companies referred to herein and should refer to the risk warnings that have been sent to them by Goldman Sachs International. A copy of these risks warnings, and a glossary of certain financial terms used in this report, are available from Goldman Sachs International on request.

**European Union and United Kingdom:** Disclosure information in relation to Article 6 (2) of the European Commission Delegated Regulation (EU) (2016/958) supplementing Regulation (EU) No 596/2014 of the European Parliament and of the Council (including as that Delegated Regulation is implemented into United Kingdom domestic law and regulation following the United Kingdom's departure from the European Union and the European Economic Area) with regard to regulatory technical standards for the technical arrangements for objective presentation of investment recommendations or other information recommending or suggesting an investment strategy and for disclosure of particular interests or indications of conflicts of interest is available at <https://www.gs.com/disclosures/europeanpolicy.html> which states the European Policy for Managing Conflicts of Interest in Connection with Investment Research.

**Japan:** Goldman Sachs Japan Co., Ltd. is a Financial Instrument Dealer registered with the Kanto Financial Bureau under registration number Kinsho

69, and a member of Japan Securities Dealers Association, Financial Futures Association of Japan Type II Financial Instruments Firms Association, The Investment Trusts Association, Japan, and Japan Investment Advisers Association. Sales and purchase of equities are subject to commission pre-determined with clients plus consumption tax. See company-specific disclosures as to any applicable disclosures required by Japanese stock exchanges, the Japanese Securities Dealers Association or the Japanese Securities Finance Company.

## Ratings, coverage universe and related definitions

**Buy (B), Neutral (N), Sell (S)** Analysts recommend stocks as Buys or Sells for inclusion on various regional Investment Lists. Being assigned a Buy or Sell on an Investment List is determined by a stock's total return potential relative to its coverage universe. Any stock not assigned as a Buy or a Sell on an Investment List with an active rating (i.e., a stock that is not Rating Suspended, Not Rated, Early-Stage Biotech, Coverage Suspended or Not Covered), is deemed Neutral. Each region manages Regional Conviction Lists, which are selected from Buy rated stocks on the respective region's Investment Lists and represent investment recommendations focused on the size of the total return potential and/or the likelihood of the realization of the return across their respective areas of coverage. The addition or removal of stocks from such Conviction Lists are managed by the Investment Review Committee or other designated committee in each respective region and do not represent a change in the analysts' investment rating for such stocks.

**Total return potential** represents the upside or downside differential between the current share price and the price target, including all paid or anticipated dividends, expected during the time horizon associated with the price target. Price targets are required for all covered stocks. The total return potential, price target and associated time horizon are stated in each report adding or reiterating an Investment List membership.

**Coverage Universe:** A list of all stocks in each coverage universe is available by primary analyst, stock and coverage universe at <https://www.gs.com/research/hedge.html>.

**Not Rated (NR).** The investment rating, target price and earnings estimates (where relevant) are removed pursuant to Goldman Sachs policy when Goldman Sachs is acting in an advisory capacity in a merger or in a strategic transaction involving this company, when there are legal, regulatory or policy constraints due to Goldman Sachs' involvement in a transaction, and in certain other circumstances. **Early-Stage Biotech (ES).** An investment rating and a target price are not assigned pursuant to Goldman Sachs policy when this company neither has a drug, treatment or medical device that has passed a Phase II clinical trial nor a license to distribute a post-Phase II drug, treatment or medical device. **Rating Suspended (RS).** Goldman Sachs Research has suspended the investment rating and price target for this stock, because there is not a sufficient fundamental basis for determining an investment rating or target price. The previous investment rating and target price, if any, are no longer in effect for this stock and should not be relied upon. **Coverage Suspended (CS).** Goldman Sachs has suspended coverage of this company. **Not Covered (NC).** Goldman Sachs does not cover this company.

## Global product; distributing entities

Goldman Sachs Global Investment Research produces and distributes research products for clients of Goldman Sachs on a global basis. Analysts based in Goldman Sachs offices around the world produce research on industries and companies, and research on macroeconomics, currencies, commodities and portfolio strategy. This research is disseminated in Australia by Goldman Sachs Australia Pty Ltd (ABN 21 006 797 897); in Brazil by Goldman Sachs do Brasil Corretora de Títulos e Valores Mobiliários S.A.; Public Communication Channel Goldman Sachs Brazil: 0800 727 5764 and / or contatogoldmanbrasil@gc.com. Available Weekdays (except holidays), from 9am to 6pm. Canal de Comunicação com o Públco Goldman Sachs Brasil: 0800 727 5764 e/ou contatogoldmanbrasil@gc.com. Horário de funcionamento: segunda-feira à sexta-feira (exceto feriados), das 9h às 18h; in Canada by Goldman Sachs & Co. LLC; in Hong Kong by Goldman Sachs (Asia) L.L.C.; in India by Goldman Sachs (India) Securities Private Ltd.; in Japan by Goldman Sachs Japan Co., Ltd.; in the Republic of Korea by Goldman Sachs (Asia) L.L.C., Seoul Branch; in New Zealand by Goldman Sachs New Zealand Limited; in Russia by OOO Goldman Sachs; in Singapore by Goldman Sachs (Singapore) Pte. (Company Number: 198602165W); and in the United States of America by Goldman Sachs & Co. LLC. Goldman Sachs International has approved this research in connection with its distribution in the United Kingdom.

Goldman Sachs International ("GSI"), authorised by the Prudential Regulation Authority ("PRA") and regulated by the Financial Conduct Authority ("FCA") and the PRA, has approved this research in connection with its distribution in the United Kingdom.

**European Economic Area:** GSI, authorised by the PRA and regulated by the FCA and the PRA, disseminates research in the following jurisdictions within the European Economic Area: the Grand Duchy of Luxembourg, Italy, the Kingdom of Belgium, the Kingdom of Denmark, the Kingdom of Norway, the Republic of Finland and the Republic of Ireland; GSI - Succursale de Paris (Paris branch) which is authorised by the French Autorité de contrôle prudentiel et de resolution ("ACPR") and regulated by the Autorité de contrôle prudentiel et de resolution and the Autorité des marchés financiers ("AMF") disseminates research in France; GSI - Sucursal en España (Madrid branch) authorized in Spain by the Comisión Nacional del Mercado de Valores disseminates research in the Kingdom of Spain; GSI - Sweden Bankfilial (Stockholm branch) is authorized by the SFSA as a "third country branch" in accordance with Chapter 4, Section 4 of the Swedish Securities and Market Act (Sw. lag (2007:528) om värdepappersmarknaden) disseminates research in the Kingdom of Sweden; Goldman Sachs Bank Europe SE ("GSBE") is a credit institution incorporated in Germany and, within the Single Supervisory Mechanism, subject to direct prudential supervision by the European Central Bank and in other respects supervised by German Federal Financial Supervisory Authority (Bundesanstalt für Finanzdienstleistungsaufsicht, BaFin) and Deutsche Bundesbank and disseminates research in the Federal Republic of Germany and those jurisdictions within the European Economic Area where GSI is not authorised to disseminate research and additionally, GSBE, Copenhagen Branch filial af GSBE, Tyskland, supervised by the Danish Financial Authority disseminates research in the Kingdom of Denmark; GSBE - Sucursal en España (Madrid branch) subject (to a limited extent) to local supervision by the Bank of Spain disseminates research in the Kingdom of Spain; GSBE - Succursale Italia (Milan branch) to the relevant applicable extent, subject to local supervision by the Bank of Italy (Banca d'Italia) and the Italian Companies and Exchange Commission (Commissione Nazionale per le Società e la Borsa "Consob") disseminates research in Italy; GSBE - Succursale de Paris (Paris branch), supervised by the AMF and by the ACPR disseminates research in France; and GSBE - Sweden Bankfilial (Stockholm branch), to a limited extent, subject to local supervision by the Swedish Financial Supervisory Authority (Finansinspektionen) disseminates research in the Kingdom of Sweden.

## General disclosures

This research is for our clients only. Other than disclosures relating to Goldman Sachs, this research is based on current public information that we consider reliable, but we do not represent it is accurate or complete, and it should not be relied on as such. The information, opinions, estimates and forecasts contained herein are as of the date hereof and are subject to change without prior notification. We seek to update our research as appropriate, but various regulations may prevent us from doing so. Other than certain industry reports published on a periodic basis, the large majority of reports are published at irregular intervals as appropriate in the analyst's judgment.

Goldman Sachs conducts a global full-service, integrated investment banking, investment management, and brokerage business. We have investment banking and other business relationships with a substantial percentage of the companies covered by Global Investment Research. Goldman Sachs & Co. LLC, the United States broker dealer, is a member of SIPC (<https://www.sipc.org>).

Our salespeople, traders, and other professionals may provide oral or written market commentary or trading strategies to our clients and principal trading desks that reflect opinions that are contrary to the opinions expressed in this research. Our asset management area, principal trading desks and investing businesses may make investment decisions that are inconsistent with the recommendations or views expressed in this research.

The analysts named in this report may have from time to time discussed with our clients, including Goldman Sachs salespersons and traders, or may discuss in this report, trading strategies that reference catalysts or events that may have a near-term impact on the market price of the equity securities discussed in this report, which impact may be directionally counter to the analyst's published price target expectations for such stocks. Any such trading strategies are distinct from and do not affect the analyst's fundamental equity rating for such stocks, which rating reflects a stock's return potential relative to its coverage universe as described herein.

We and our affiliates, officers, directors, and employees will from time to time have long or short positions in, act as principal in, and buy or sell, the securities or derivatives, if any, referred to in this research, unless otherwise prohibited by regulation or Goldman Sachs policy.

The views attributed to third party presenters at Goldman Sachs arranged conferences, including individuals from other parts of Goldman Sachs, do not necessarily reflect those of Global Investment Research and are not an official view of Goldman Sachs.

Any third party referenced herein, including any salespeople, traders and other professionals or members of their household, may have positions in the products mentioned that are inconsistent with the views expressed by analysts named in this report.

This research is not an offer to sell or the solicitation of an offer to buy any security in any jurisdiction where such an offer or solicitation would be illegal. It does not constitute a personal recommendation or take into account the particular investment objectives, financial situations, or needs of individual clients. Clients should consider whether any advice or recommendation in this research is suitable for their particular circumstances and, if appropriate, seek professional advice, including tax advice. The price and value of investments referred to in this research and the income from them may fluctuate. Past performance is not a guide to future performance, future returns are not guaranteed, and a loss of original capital may occur. Fluctuations in exchange rates could have adverse effects on the value or price of, or income derived from, certain investments.

Certain transactions, including those involving futures, options, and other derivatives, give rise to substantial risk and are not suitable for all investors. Investors should review current options and futures disclosure documents which are available from Goldman Sachs sales representatives or at <https://www.theocc.com/about/publications/character-risks.jsp> and <https://www.fiadocumentation.org/fia/regulatory-disclosures/fia-uniform-futures-and-options-on-futures-risk-disclosures-booklet-pdf-version-2018>. Transaction costs may be significant in option strategies calling for multiple purchase and sales of options such as spreads. Supporting documentation will be supplied upon request.

**Differing Levels of Service provided by Global Investment Research:** The level and types of services provided to you by Goldman Sachs Global Investment Research may vary as compared to that provided to internal and other external clients of GS, depending on various factors including your individual preferences as to the frequency and manner of receiving communication, your risk profile and investment focus and perspective (e.g., marketwide, sector specific, long term, short term), the size and scope of your overall client relationship with GS, and legal and regulatory constraints. As an example, certain clients may request to receive notifications when research on specific securities is published, and certain clients may request that specific data underlying analysts' fundamental analysis available on our internal client websites be delivered to them electronically through data feeds or otherwise. No change to an analyst's fundamental research views (e.g., ratings, price targets, or material changes to earnings estimates for equity securities), will be communicated to any client prior to inclusion of such information in a research report broadly disseminated through electronic publication to our internal client websites or through other means, as necessary, to all clients who are entitled to receive such reports.

All research reports are disseminated and available to all clients simultaneously through electronic publication to our internal client websites. Not all research content is redistributed to our clients or available to third-party aggregators, nor is Goldman Sachs responsible for the redistribution of our research by third party aggregators. For research, models or other data related to one or more securities, markets or asset classes (including related services) that may be available to you, please contact your GS representative or go to <https://research.gs.com>.

Disclosure information is also available at <https://www.gs.com/research/hedge.html> or from Research Compliance, 200 West Street, New York, NY 10282.

**© 2025 Goldman Sachs.**

You are permitted to store, display, analyze, modify, reformat, and print the information made available to you via this service only for your own use. You may not resell or reverse engineer this information to calculate or develop any index for disclosure and/or marketing or create any other derivative works or commercial product(s), data or offering(s) without the express written consent of Goldman Sachs. You are not permitted to publish, transmit, or otherwise reproduce this information, in whole or in part, in any format to any third party without the express written consent of Goldman Sachs. This foregoing restriction includes, without limitation, using, extracting, downloading or retrieving this information, in whole or in part, to train or finetune a machine learning or artificial intelligence system, or to provide or reproduce this information, in whole or in part, as a prompt or input to any such system.